

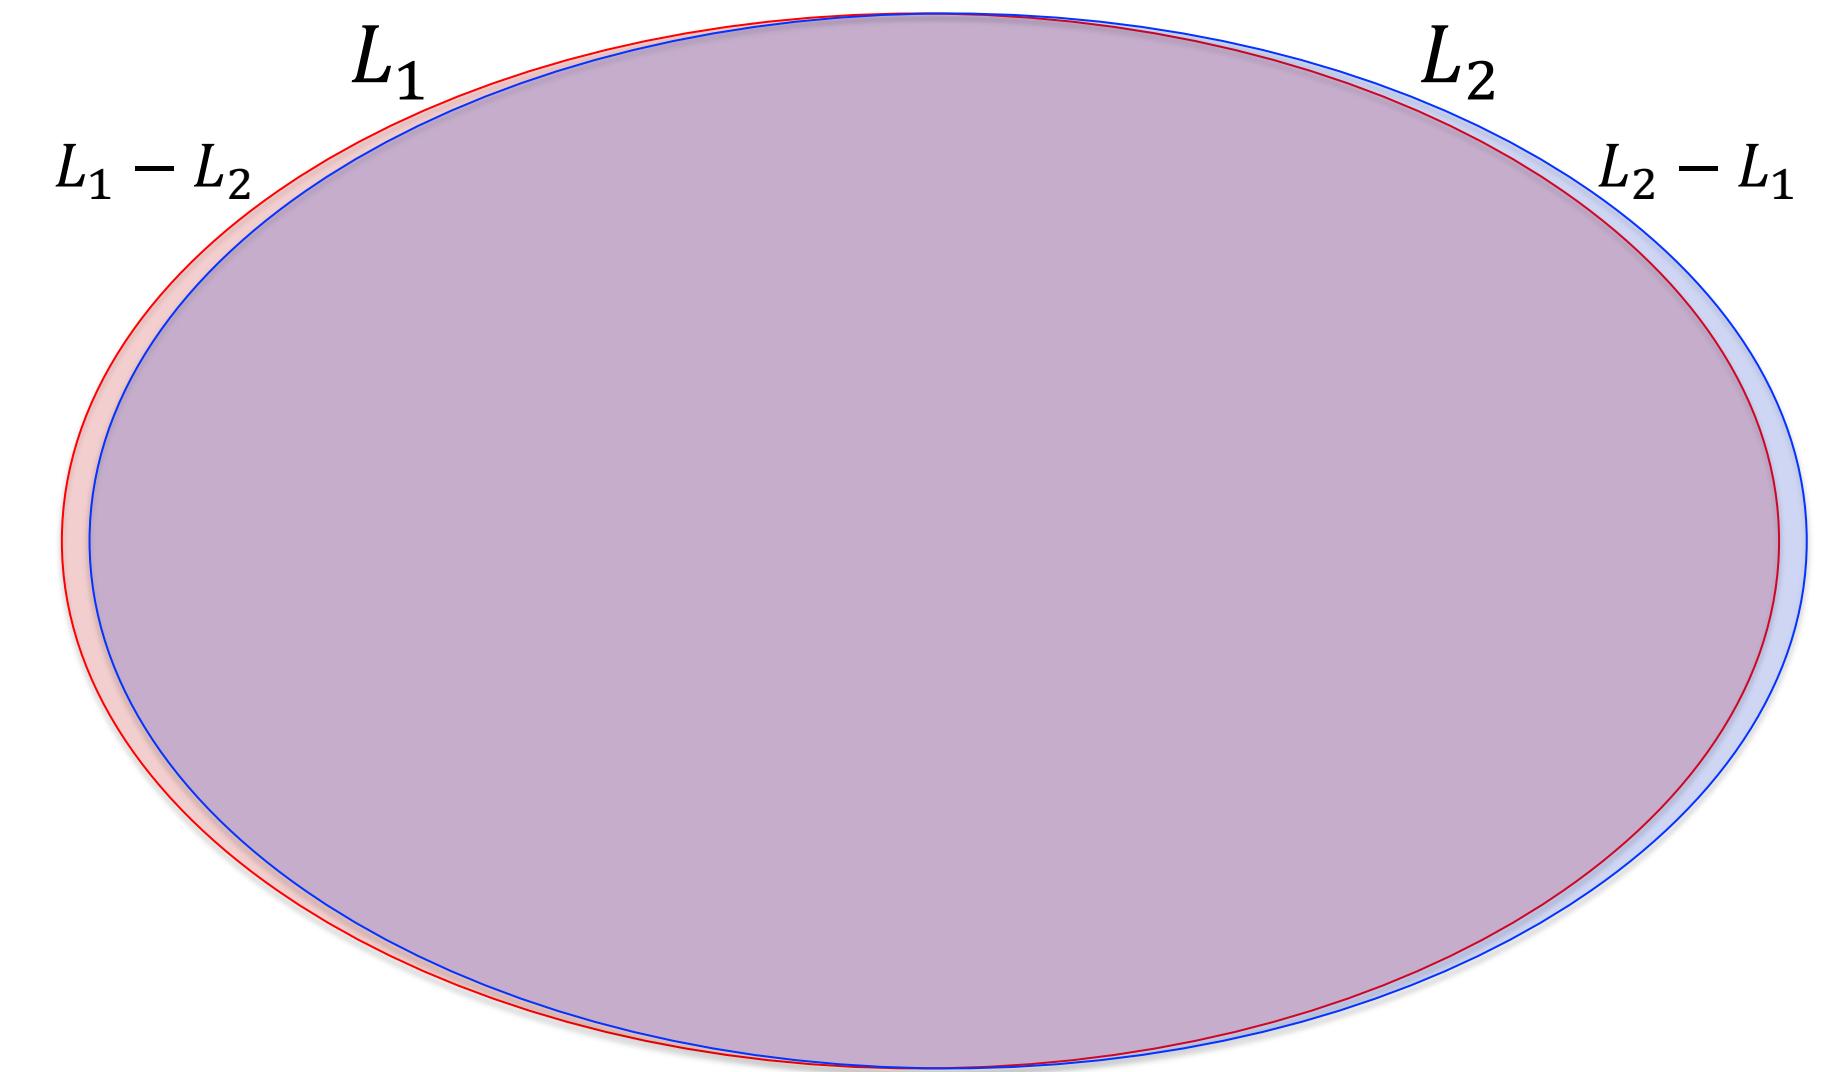
Invertible Bloom Lookup Tables

Lucas Kellar - 12/04/2024

The Set Reconciliation Problem

- Extract differences between very large sets with high overlap
- Datasets that hold small amount of data at extraction time, but large amounts before then

Set Reconciliation

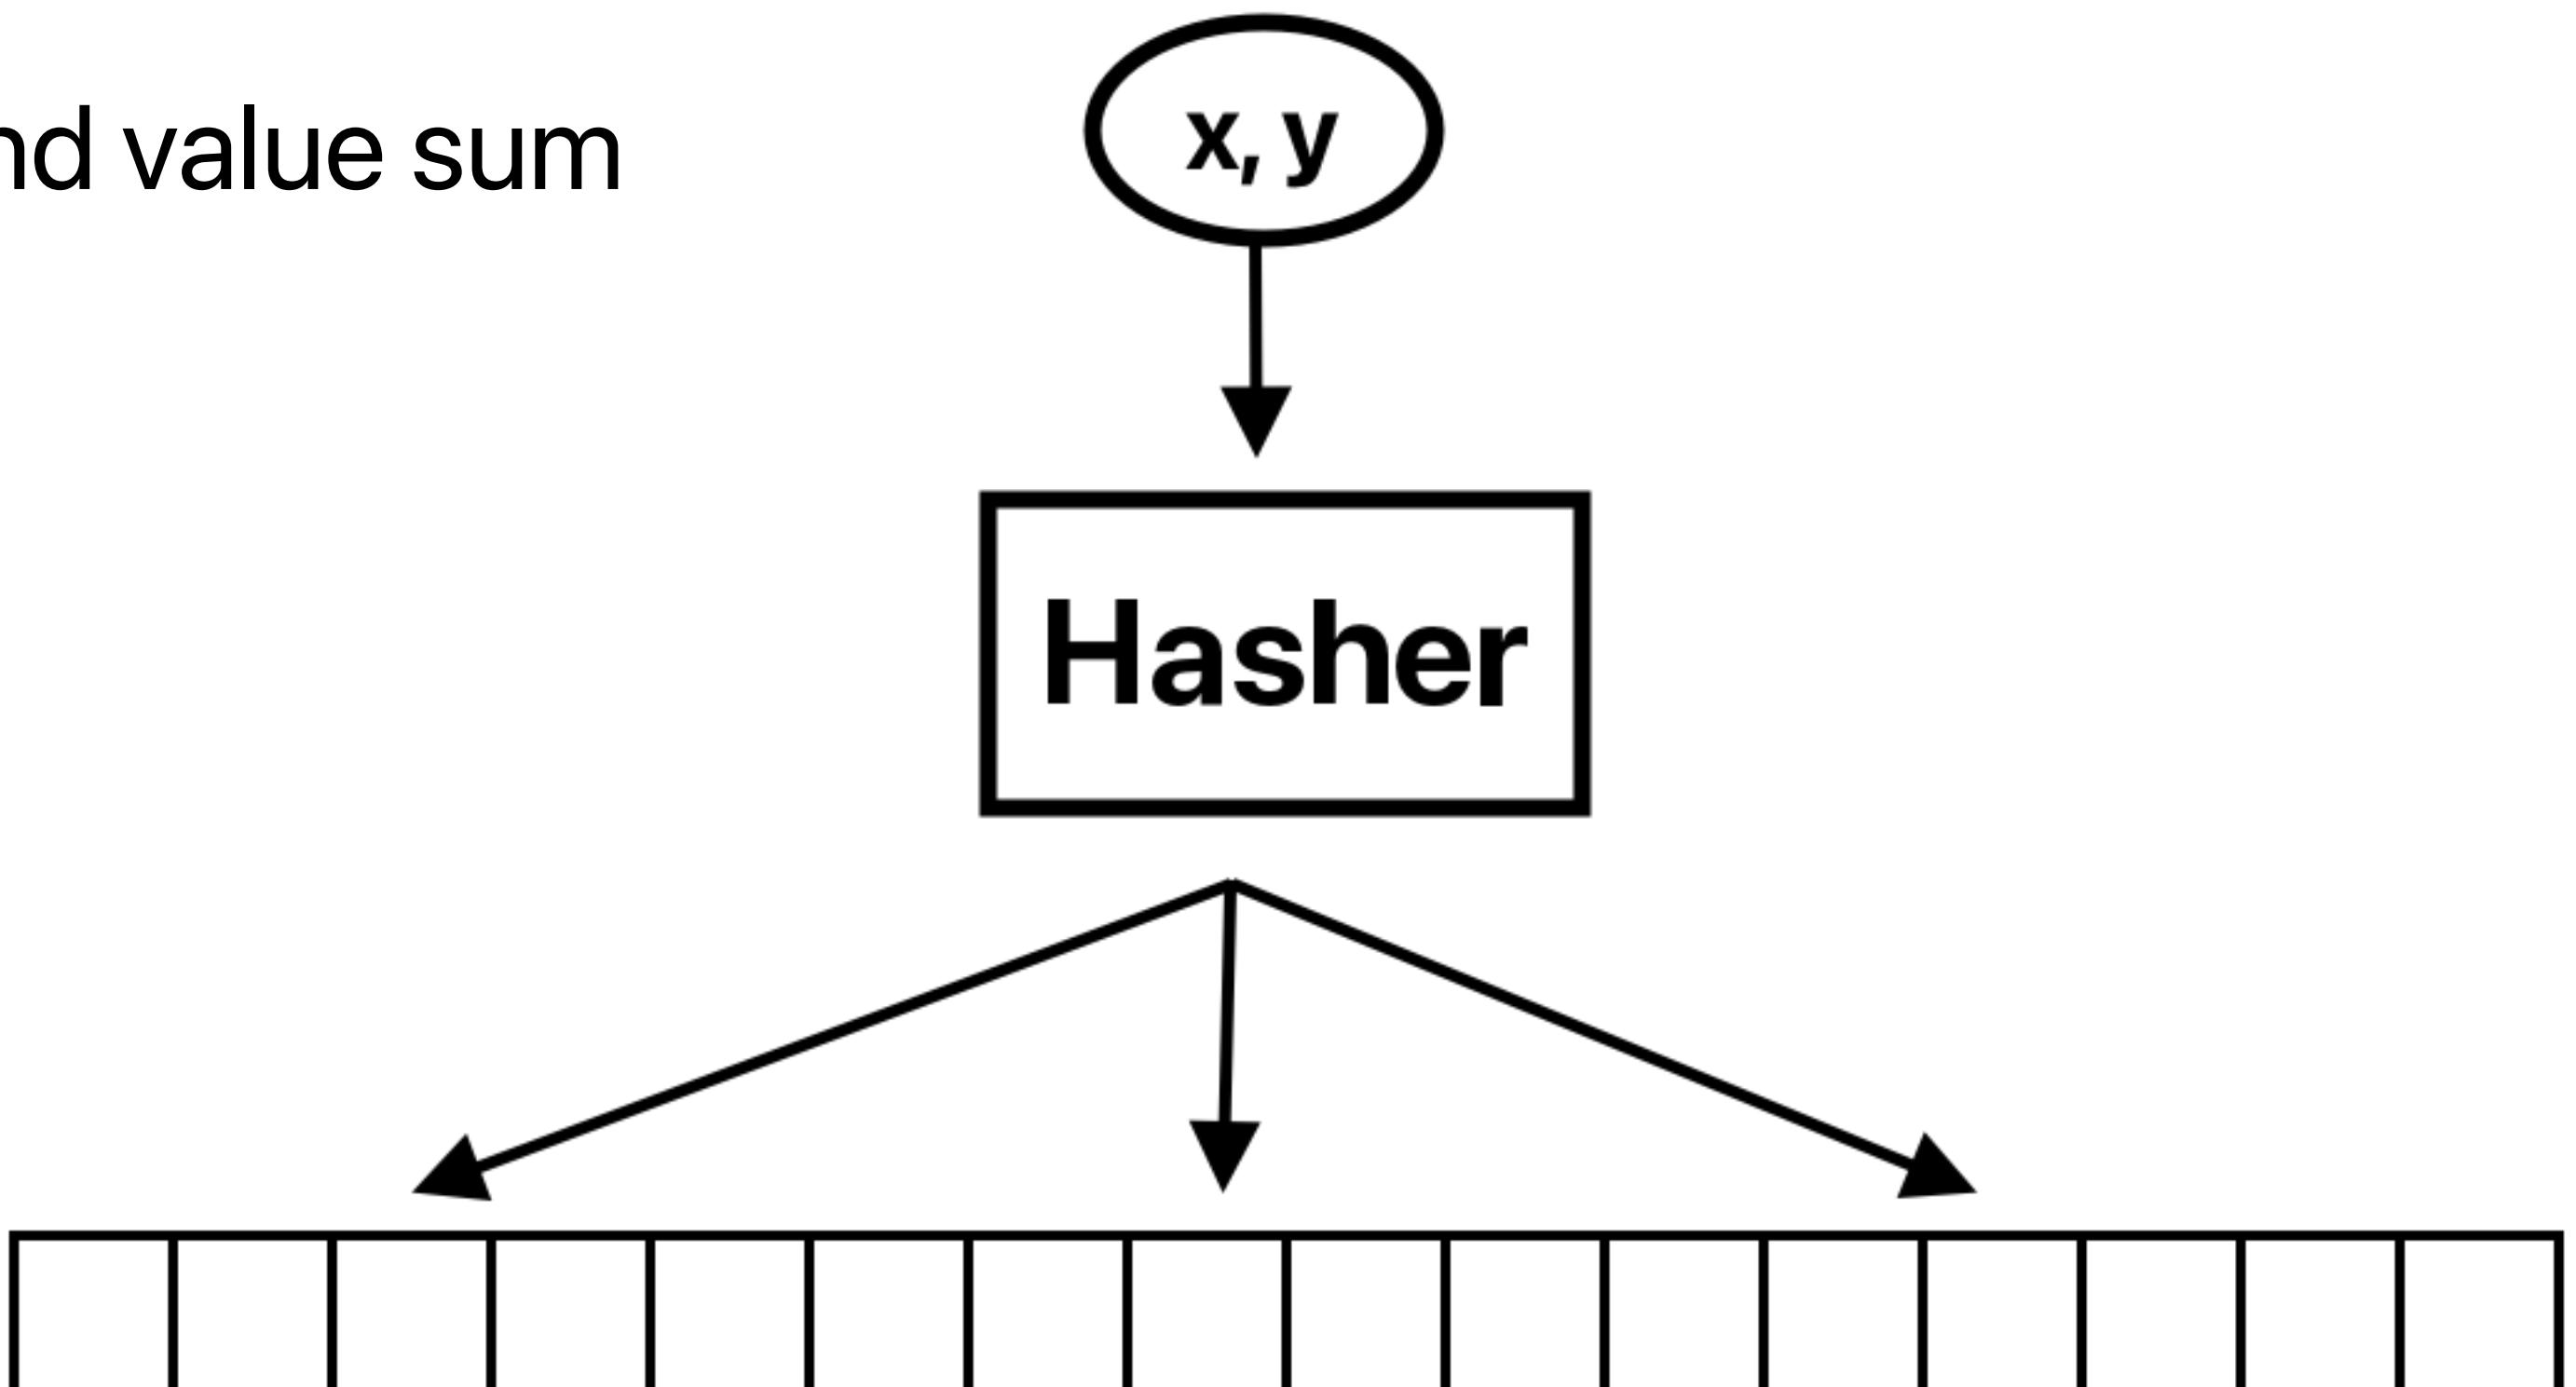


Implementation

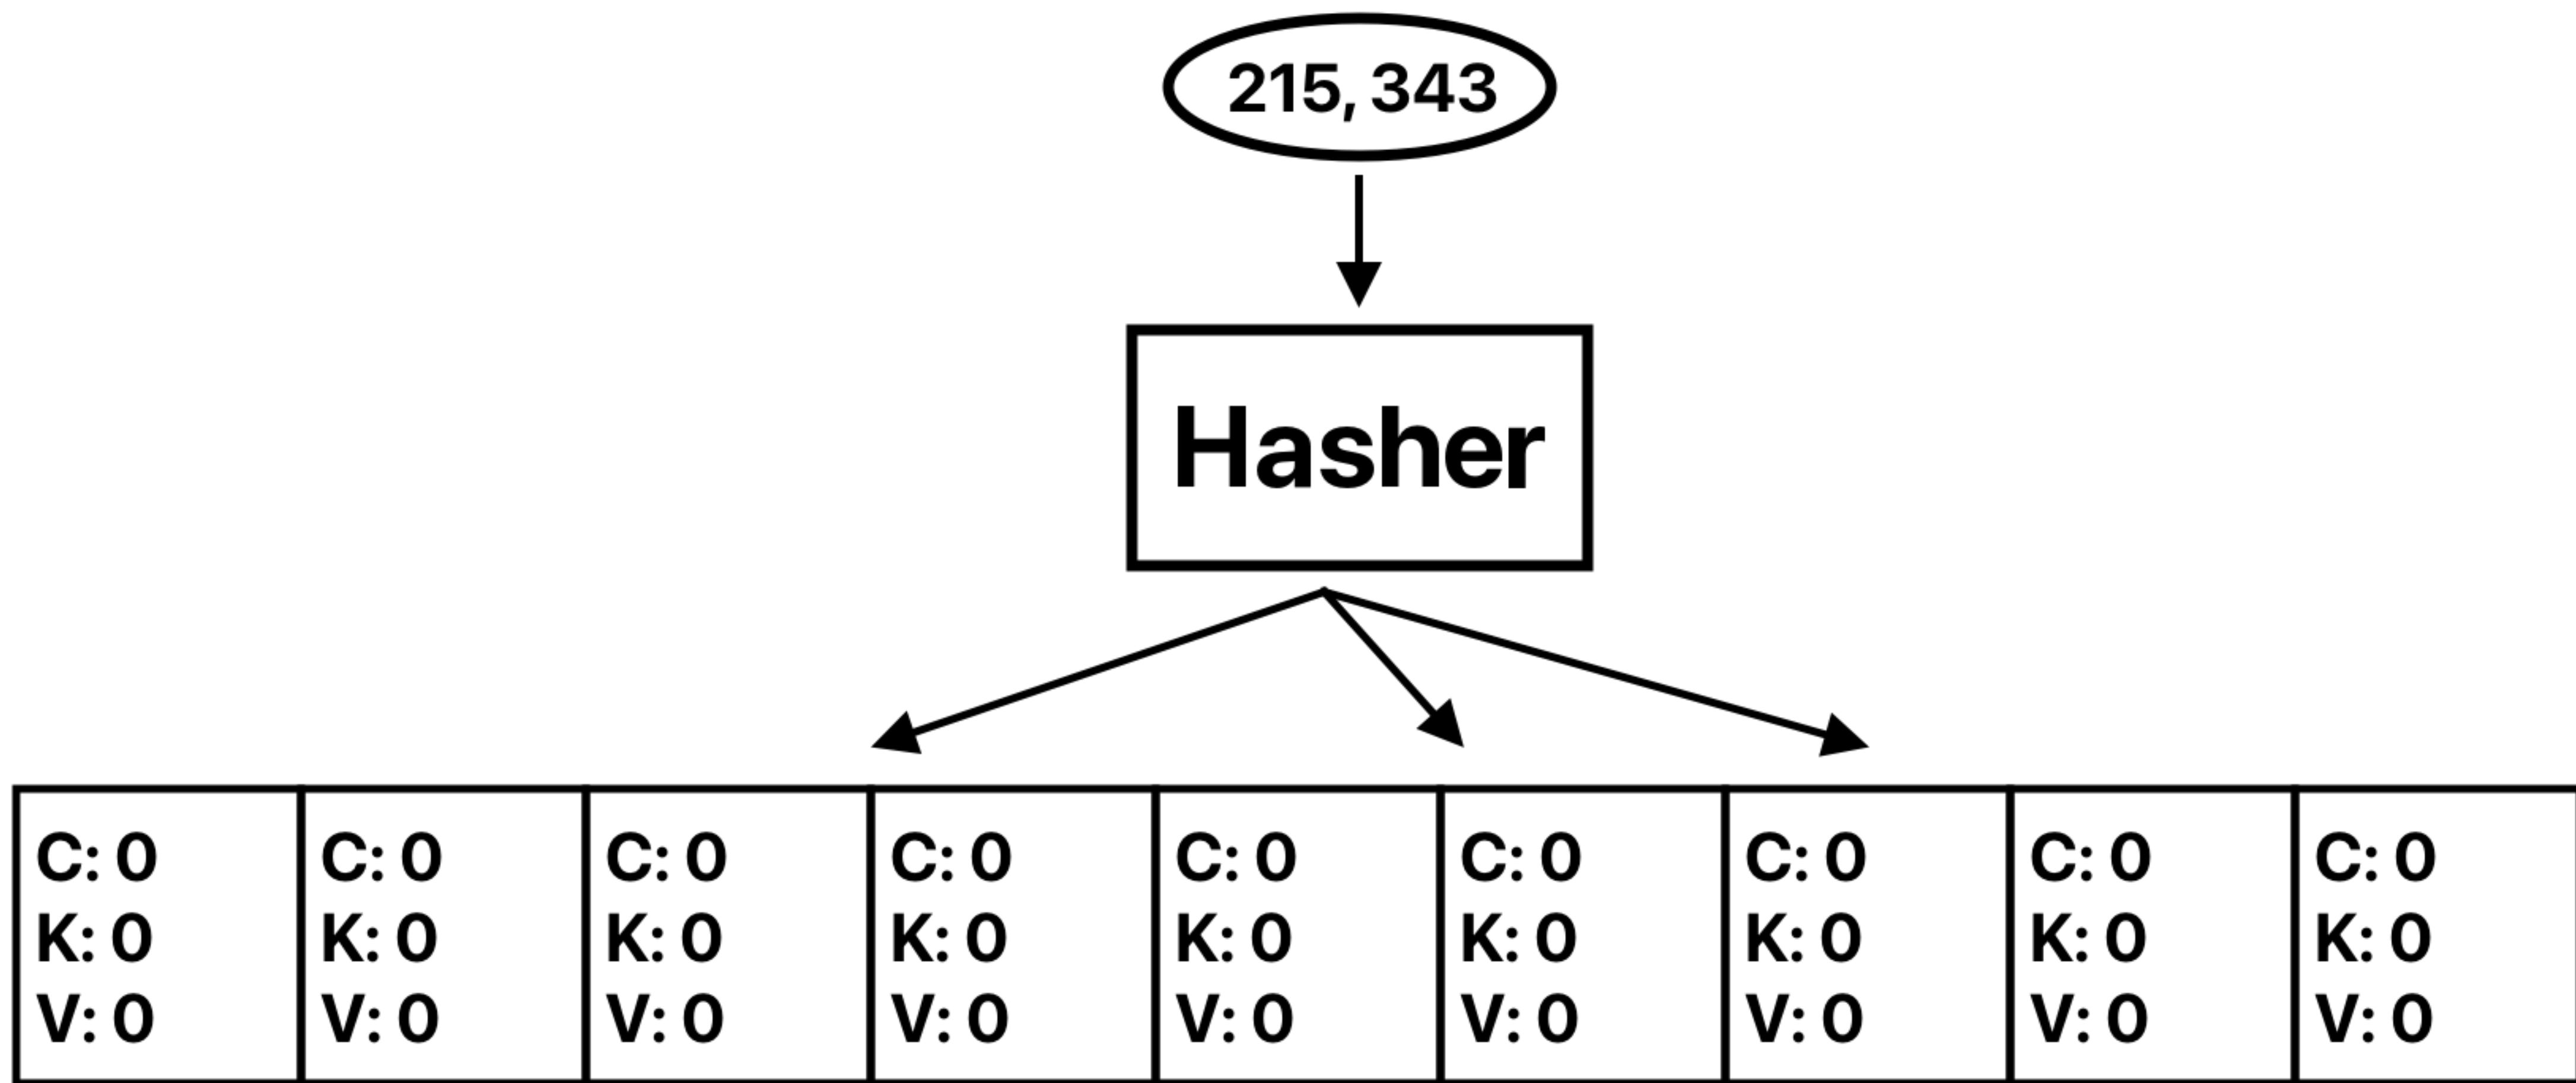
- Original IBLT (2011)
- HPW IBLT (2023)
- Stacked IBLT (2024)

Original IBLT

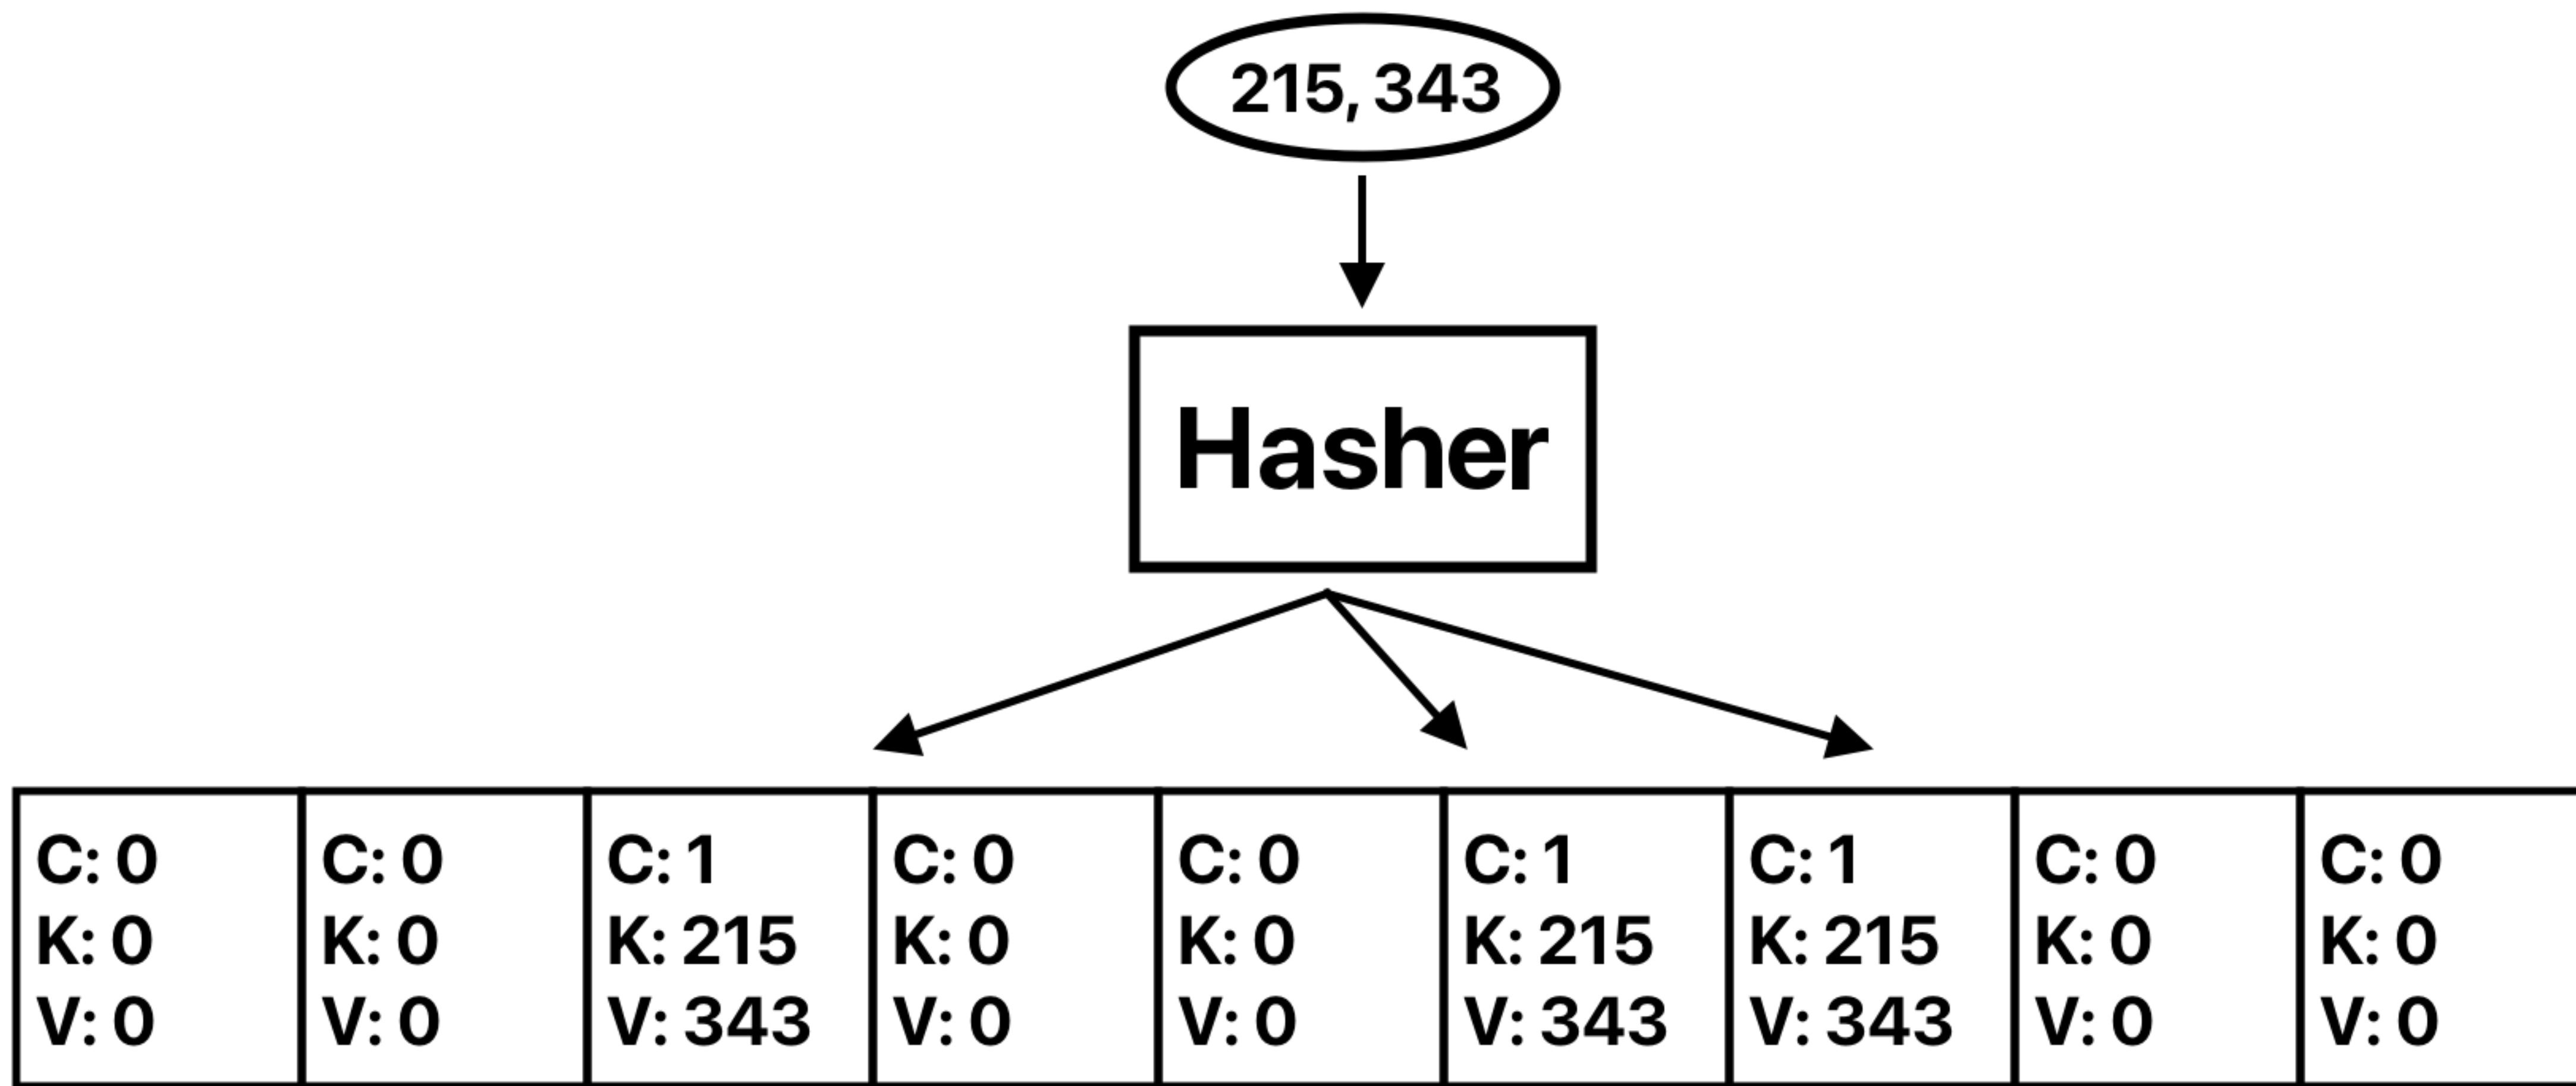
- Key/Value Store
- Each cell stores count, key sum, and value sum
- Operations:
 - Insert(x, y)
 - Delete(x, y)
 - Get(x)
 - ListEntries()



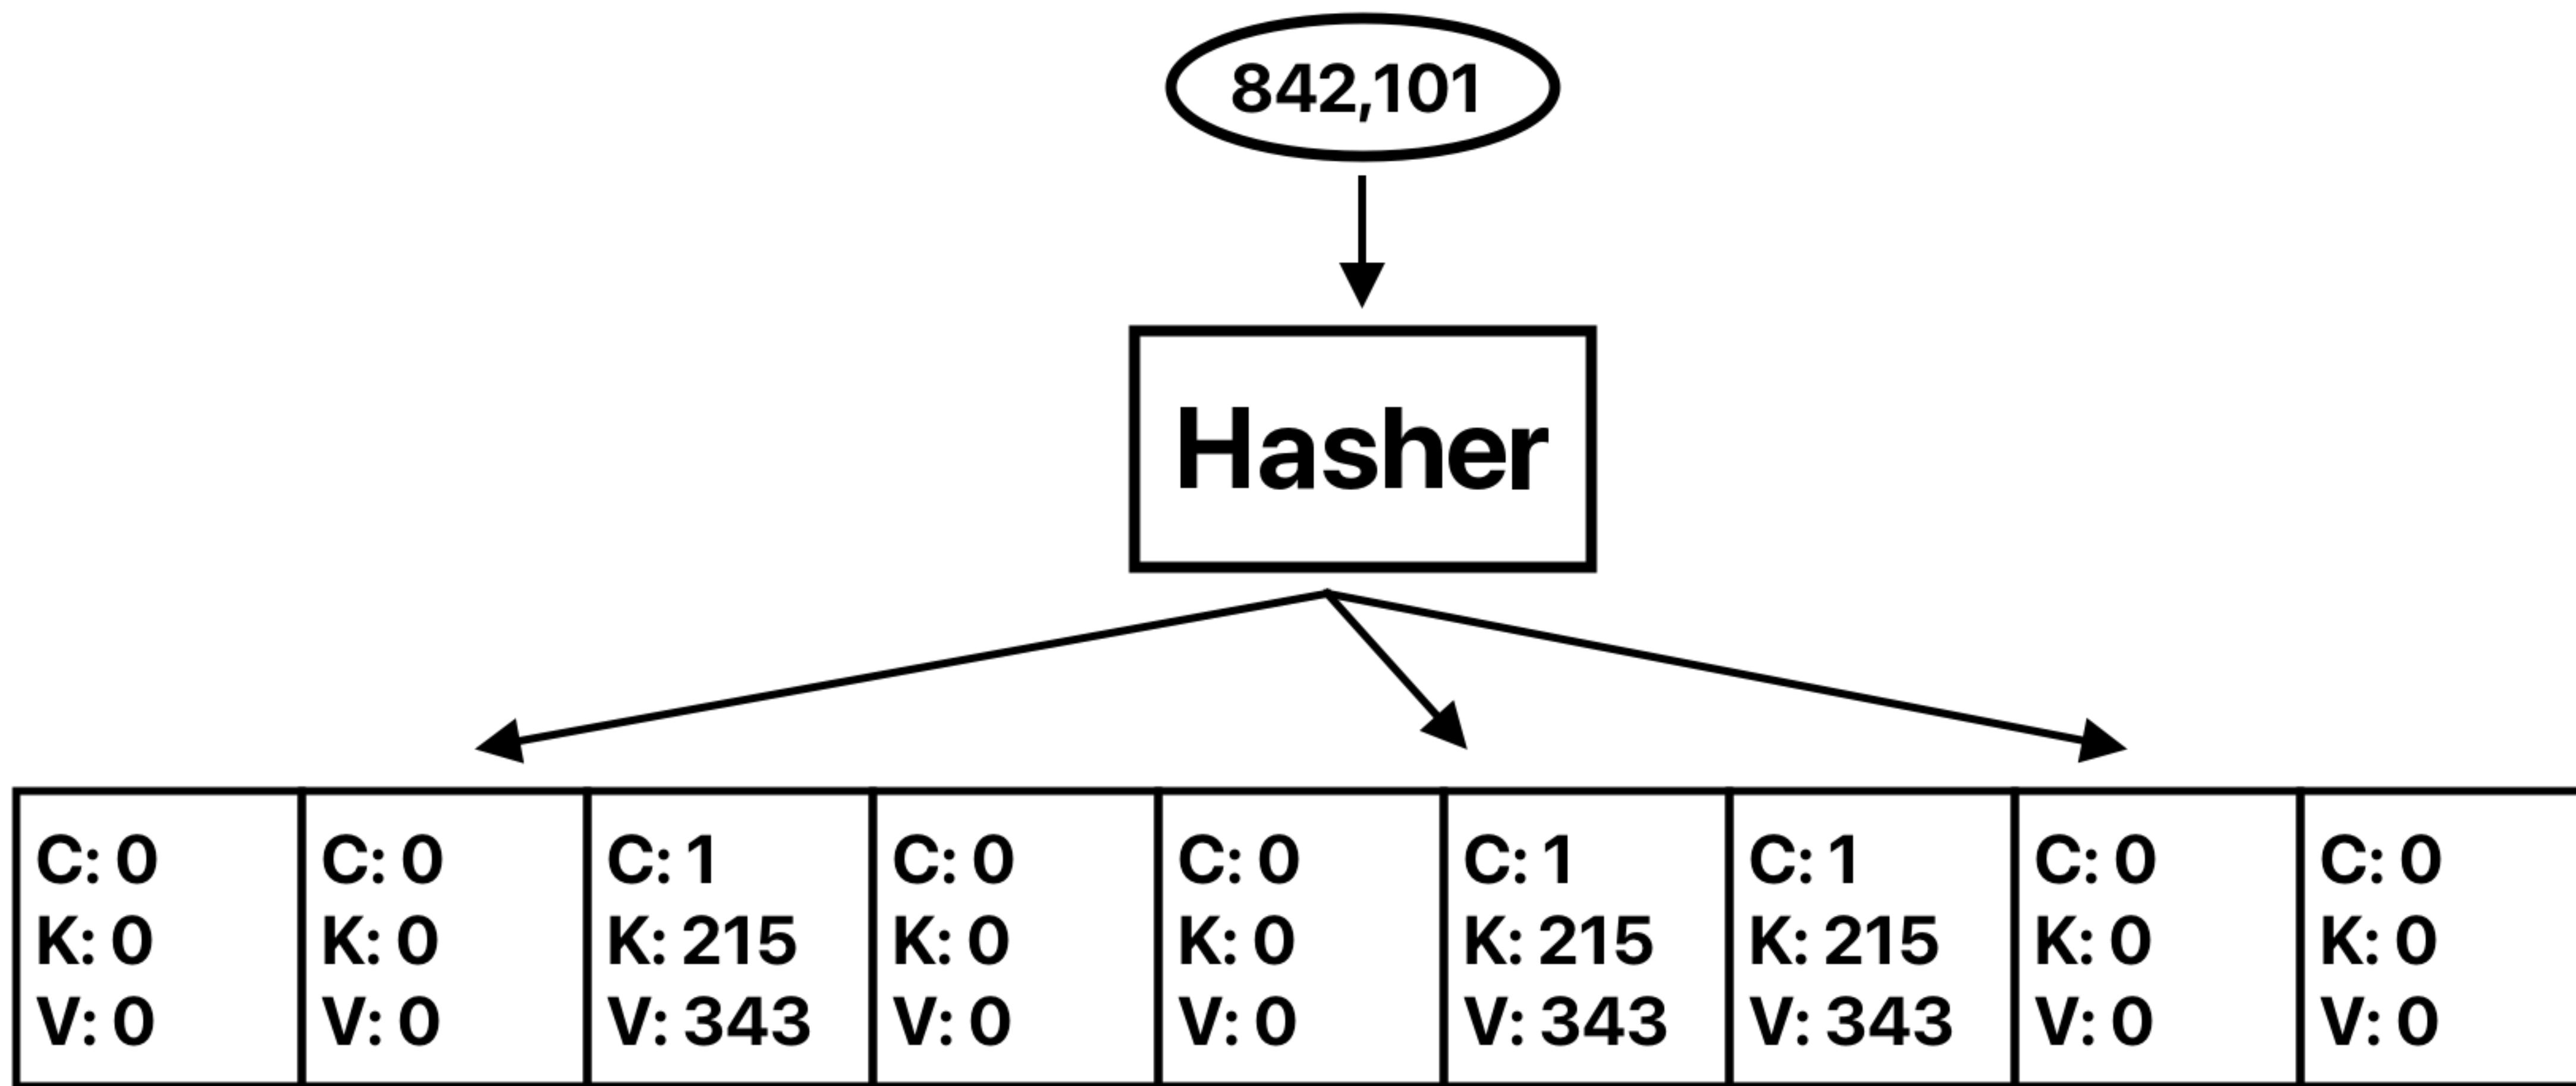
Original IBLT - Insert(x, y)



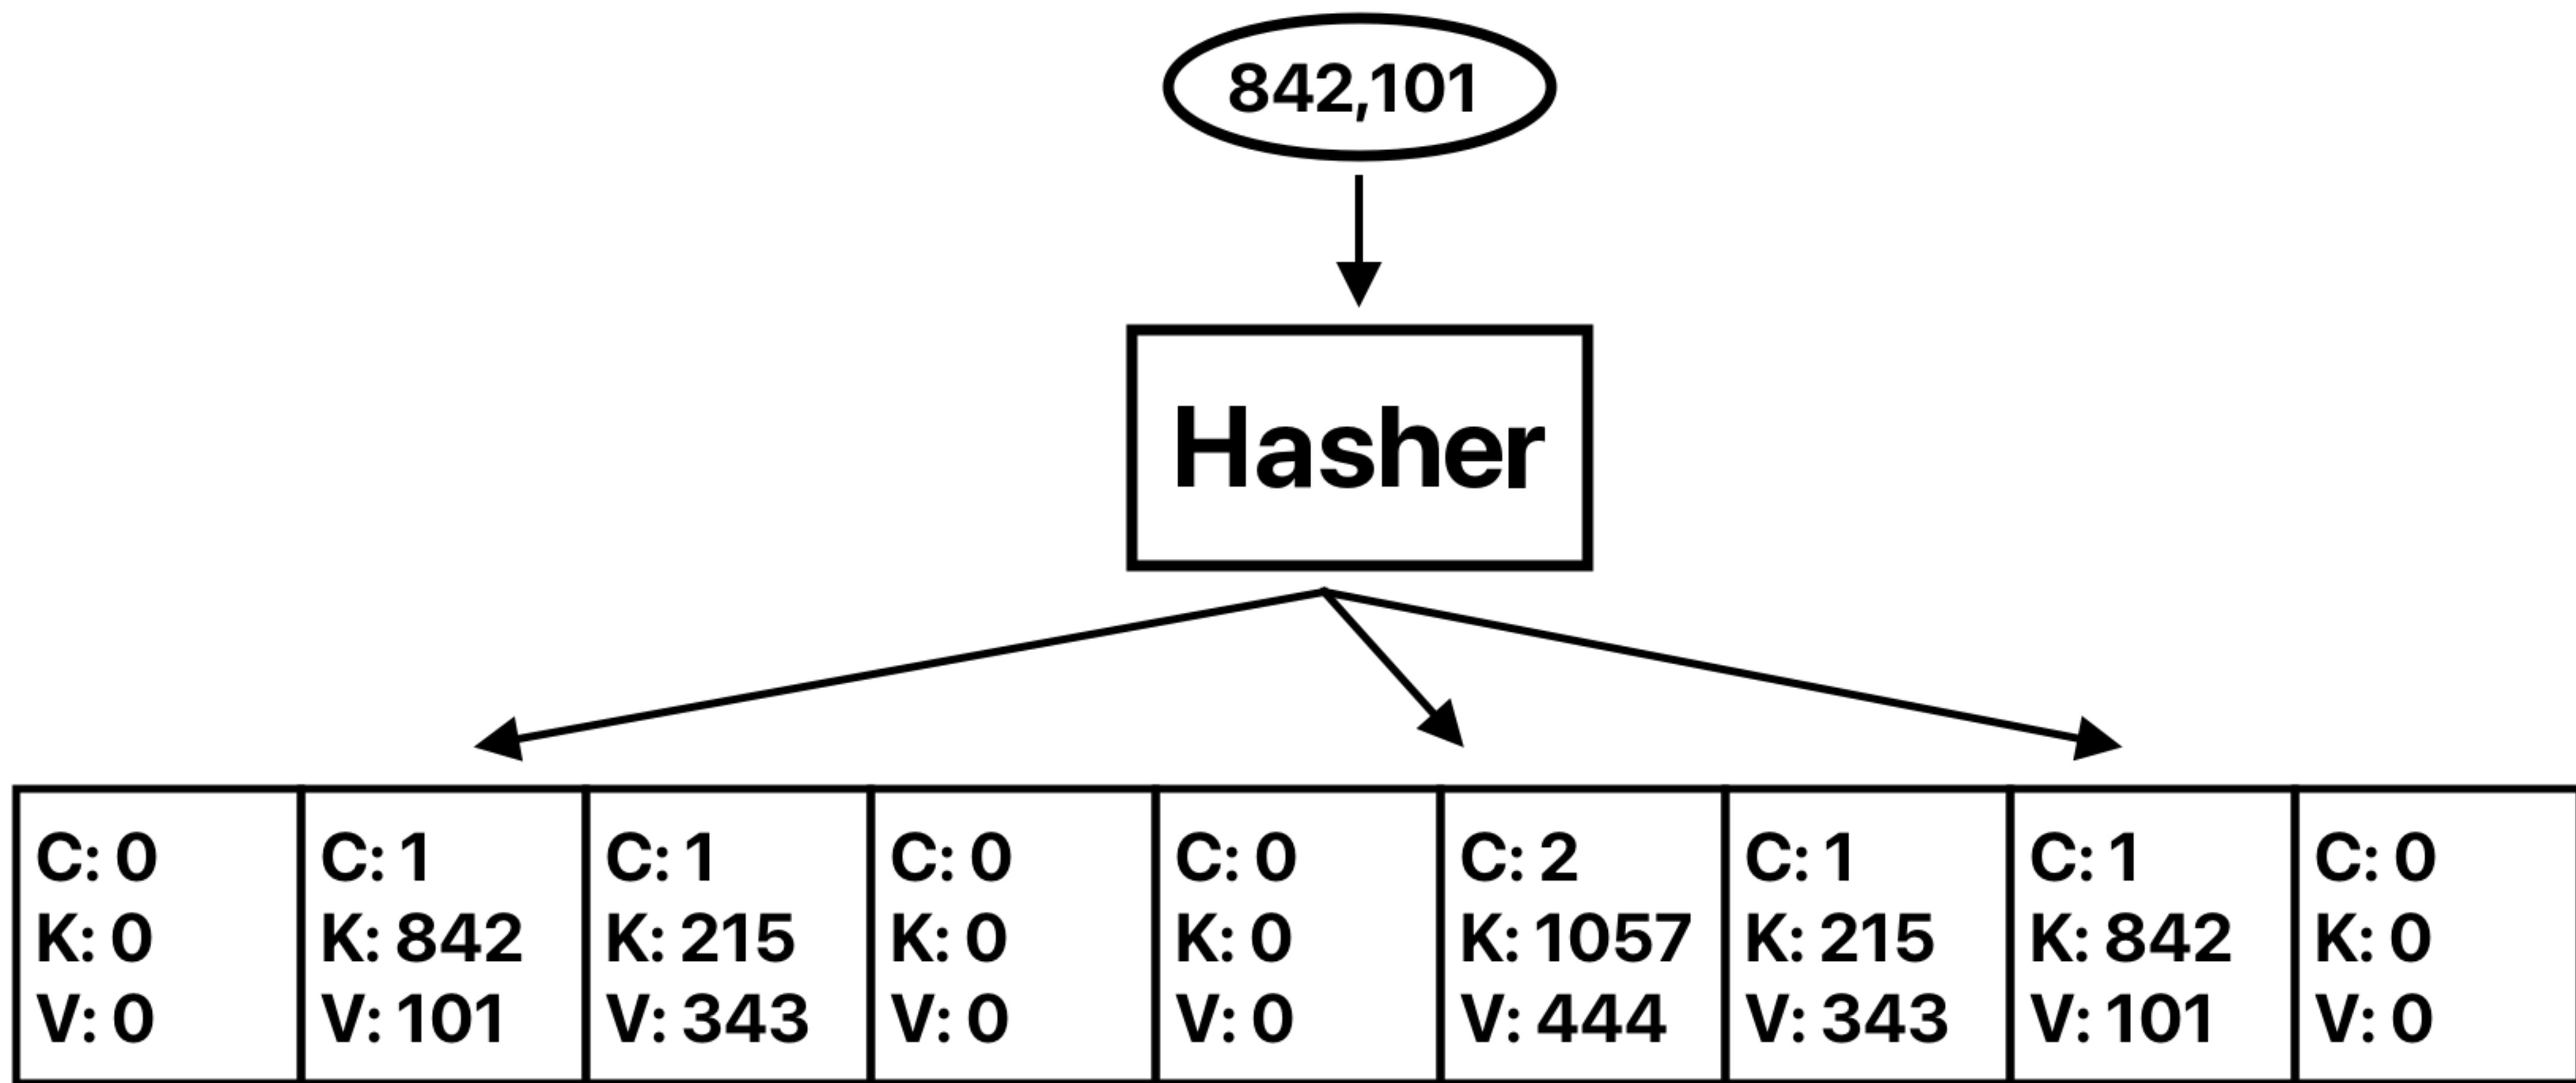
Original IBLT - Insert(x, y)



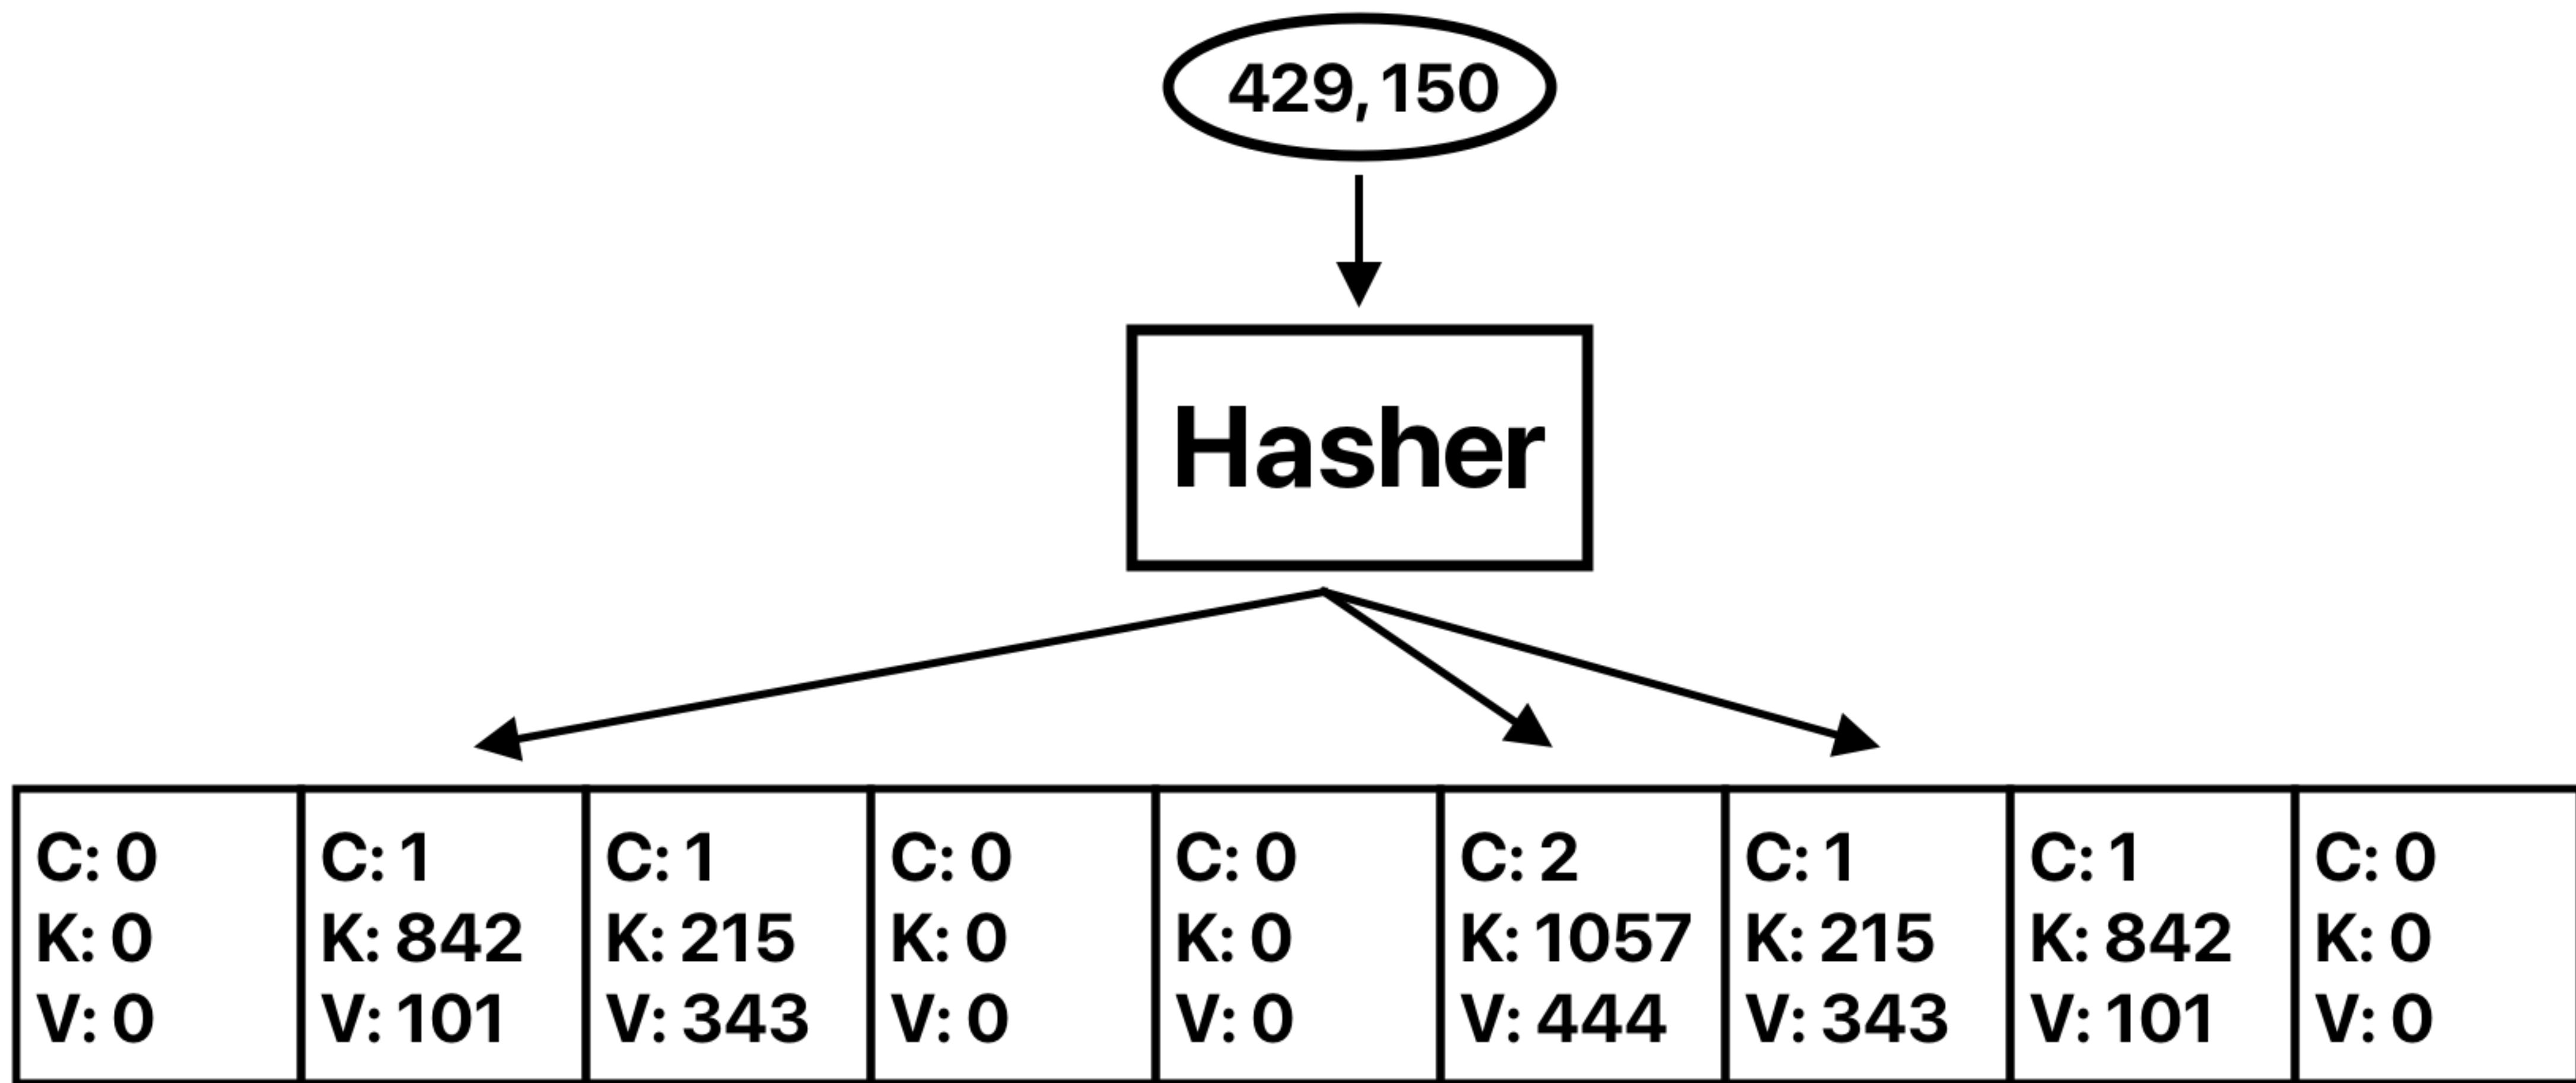
Original IBLT - Insert(x, y)



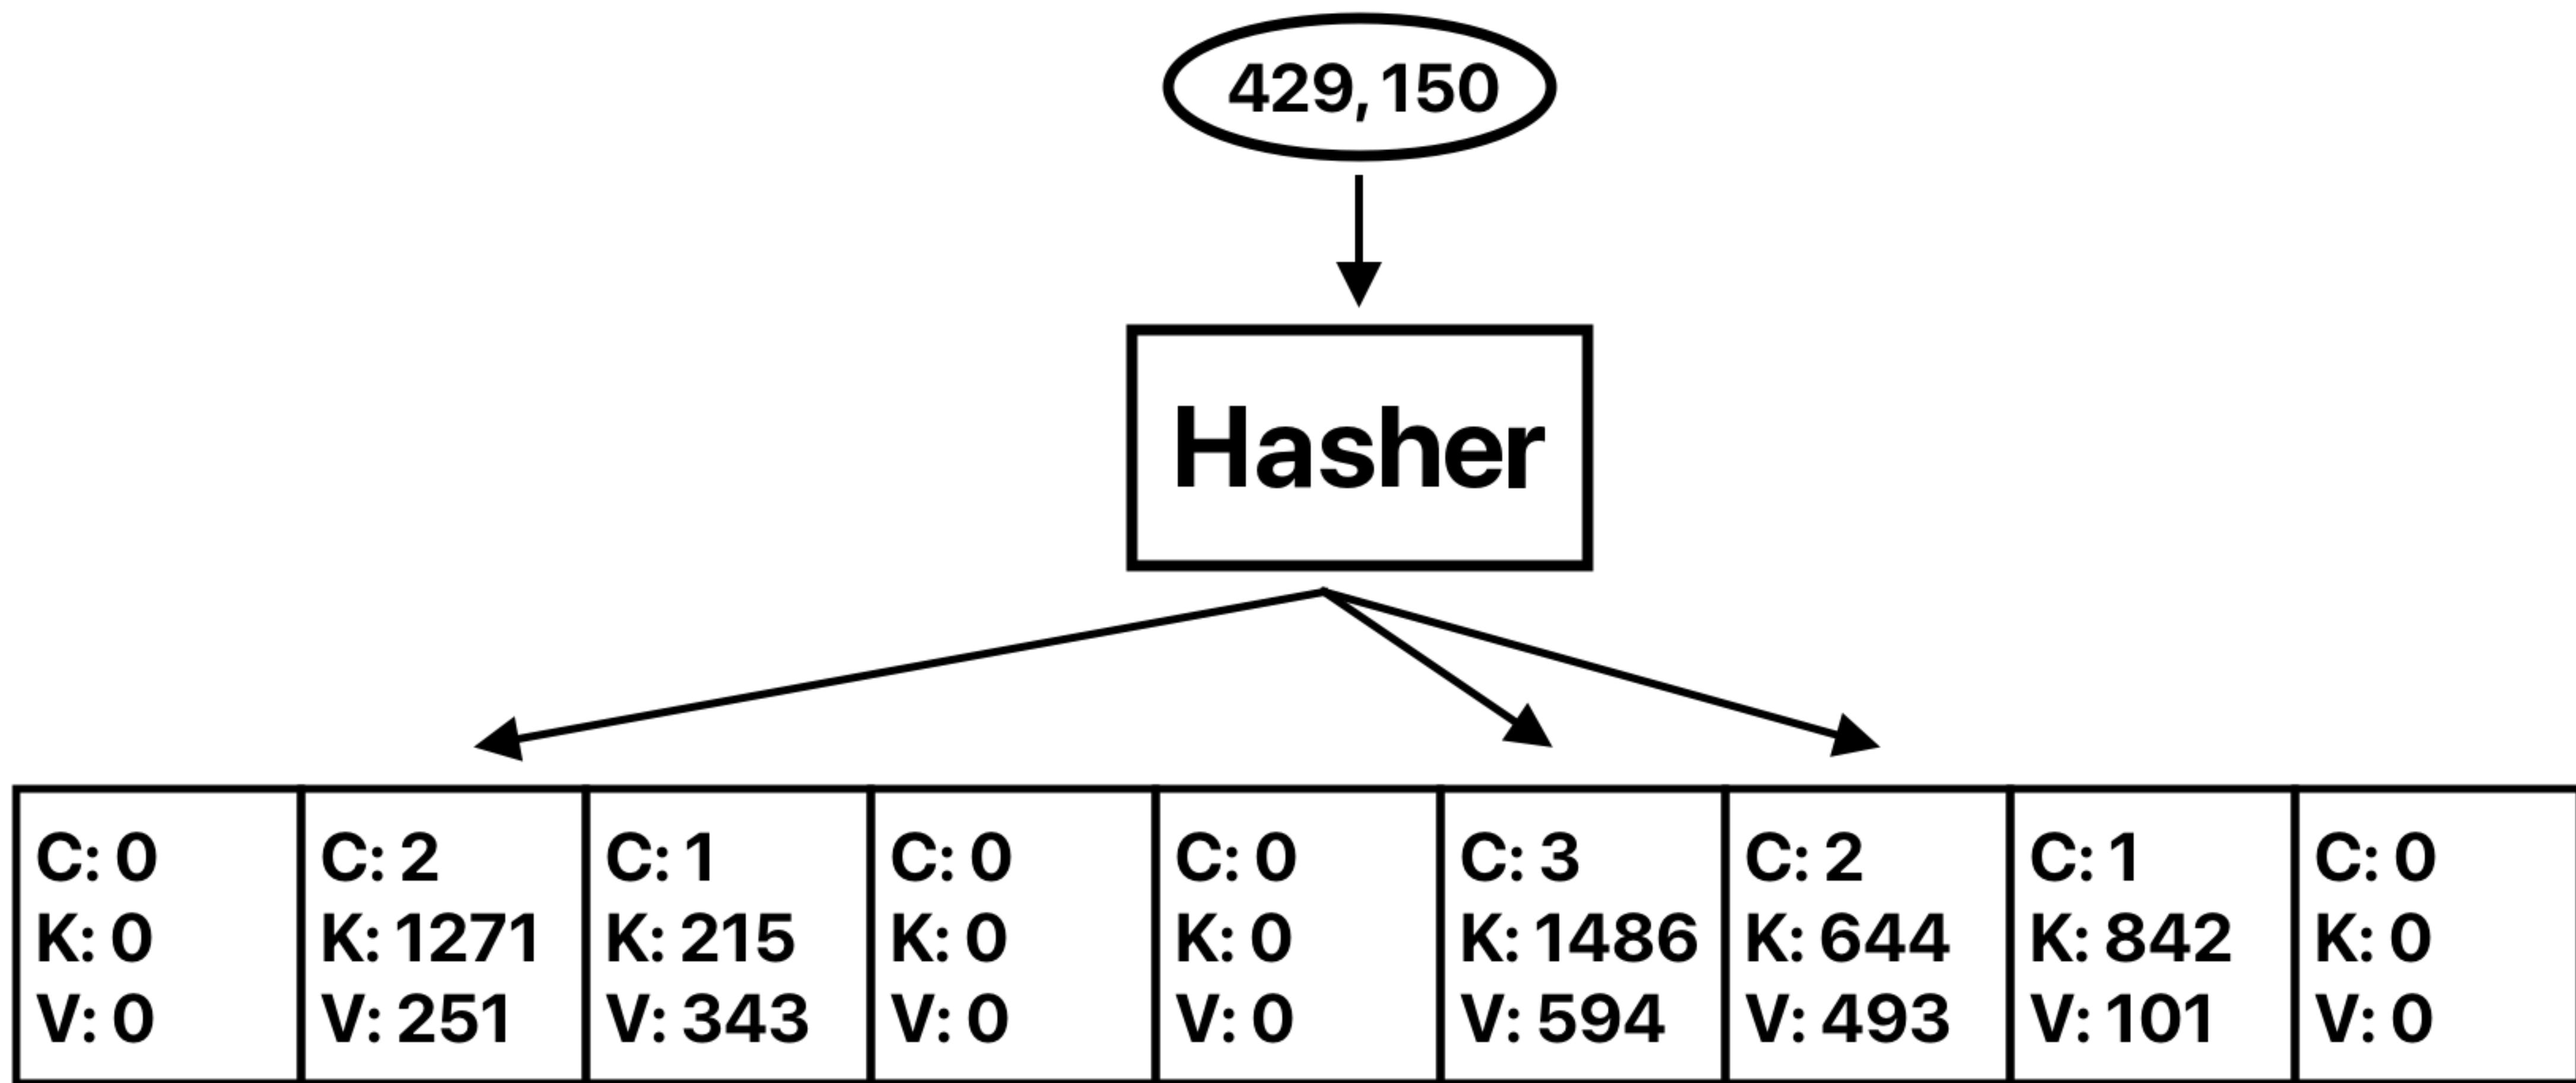
Original IBLT - Insert(x, y)



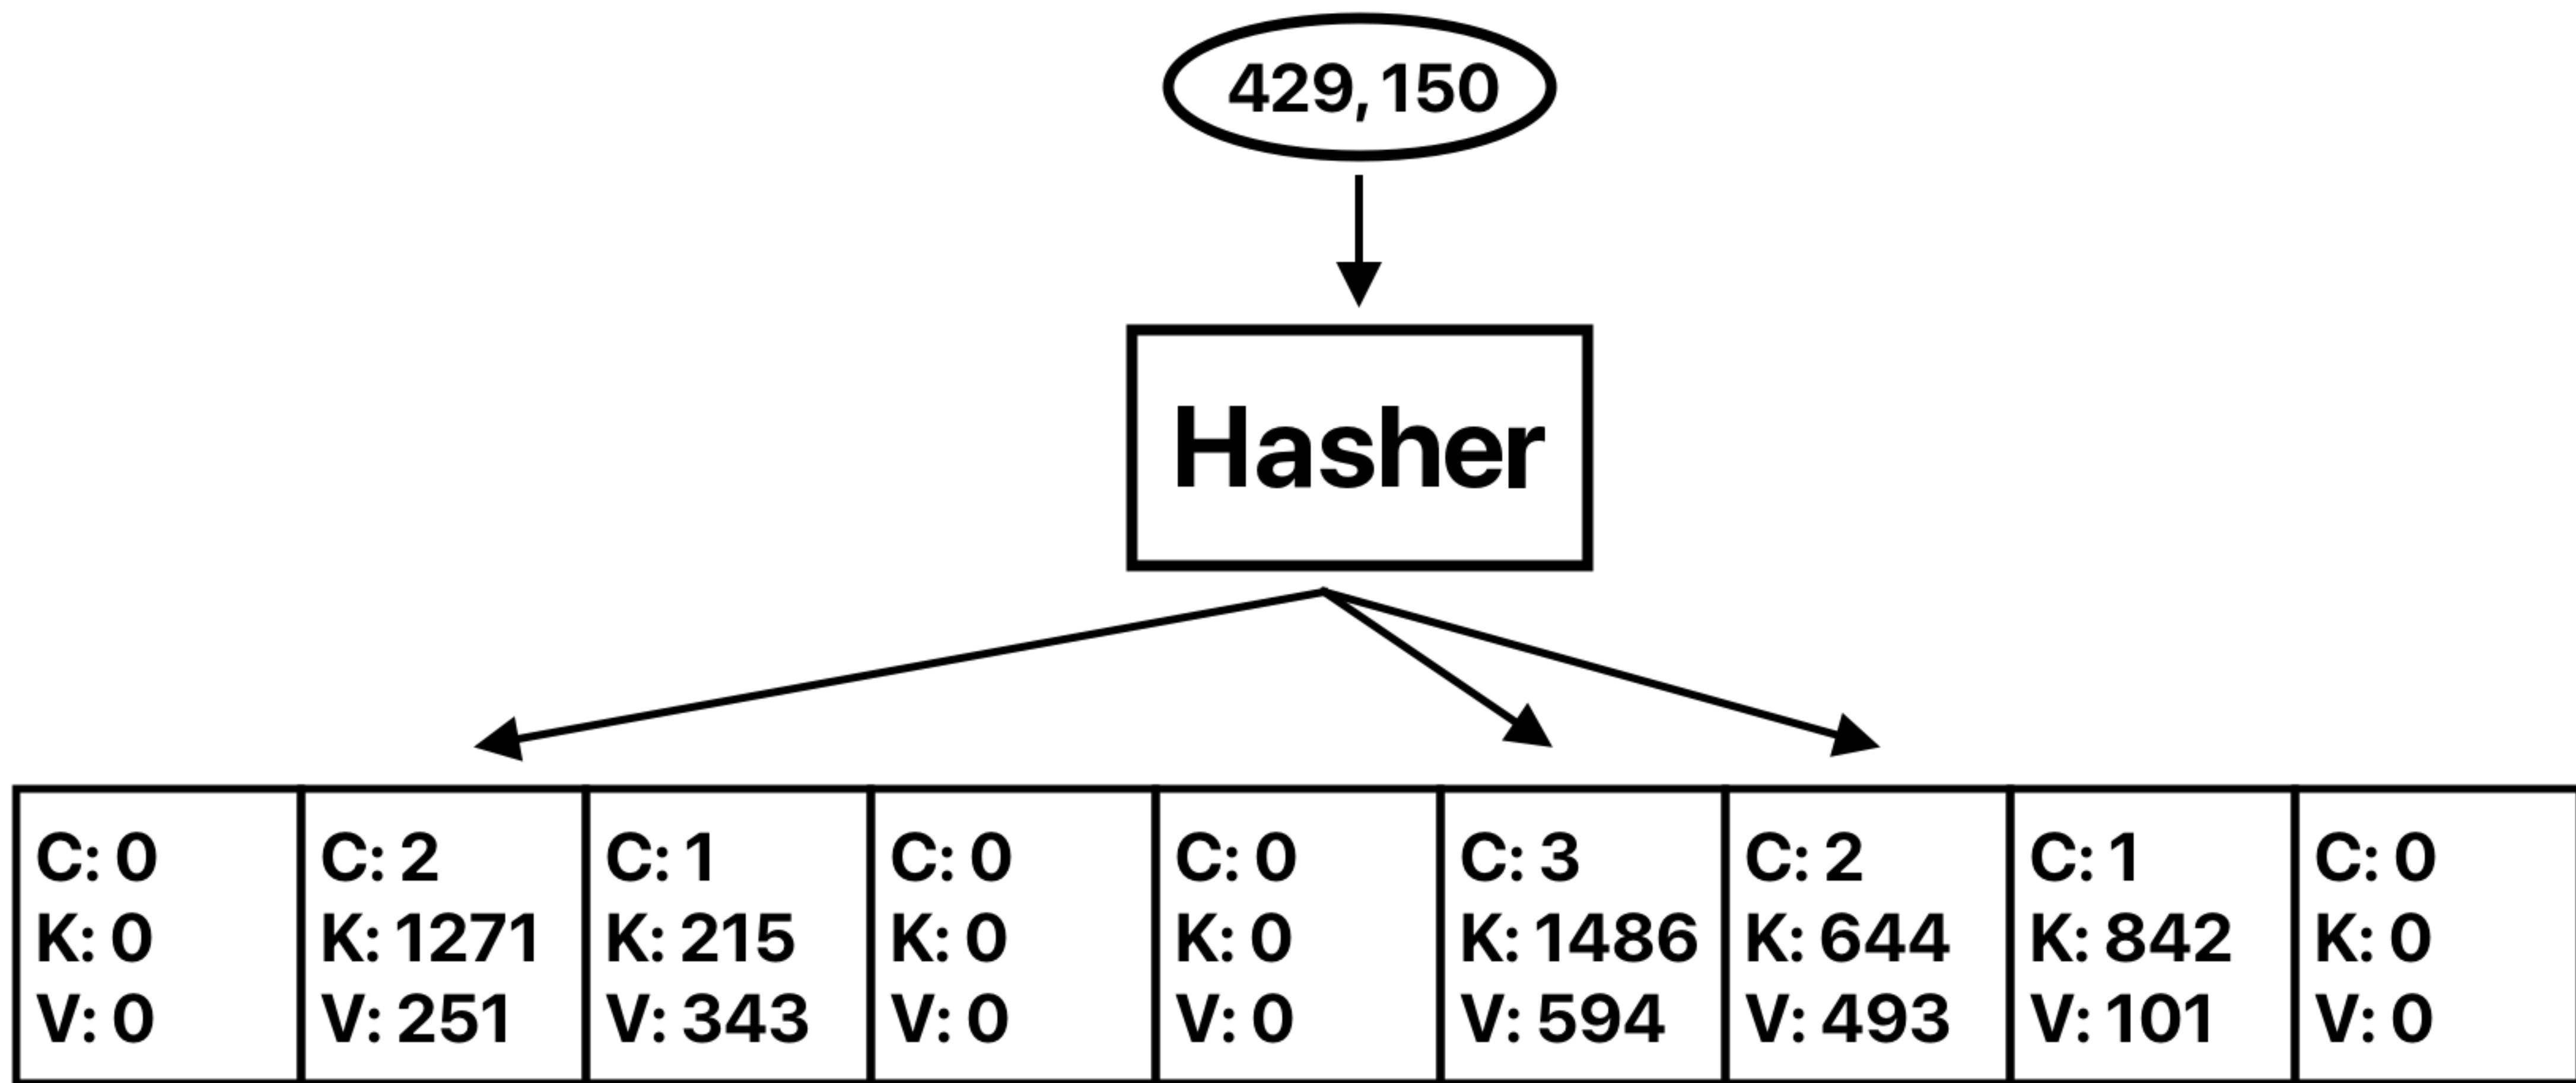
Original IBLT - Insert(x, y)



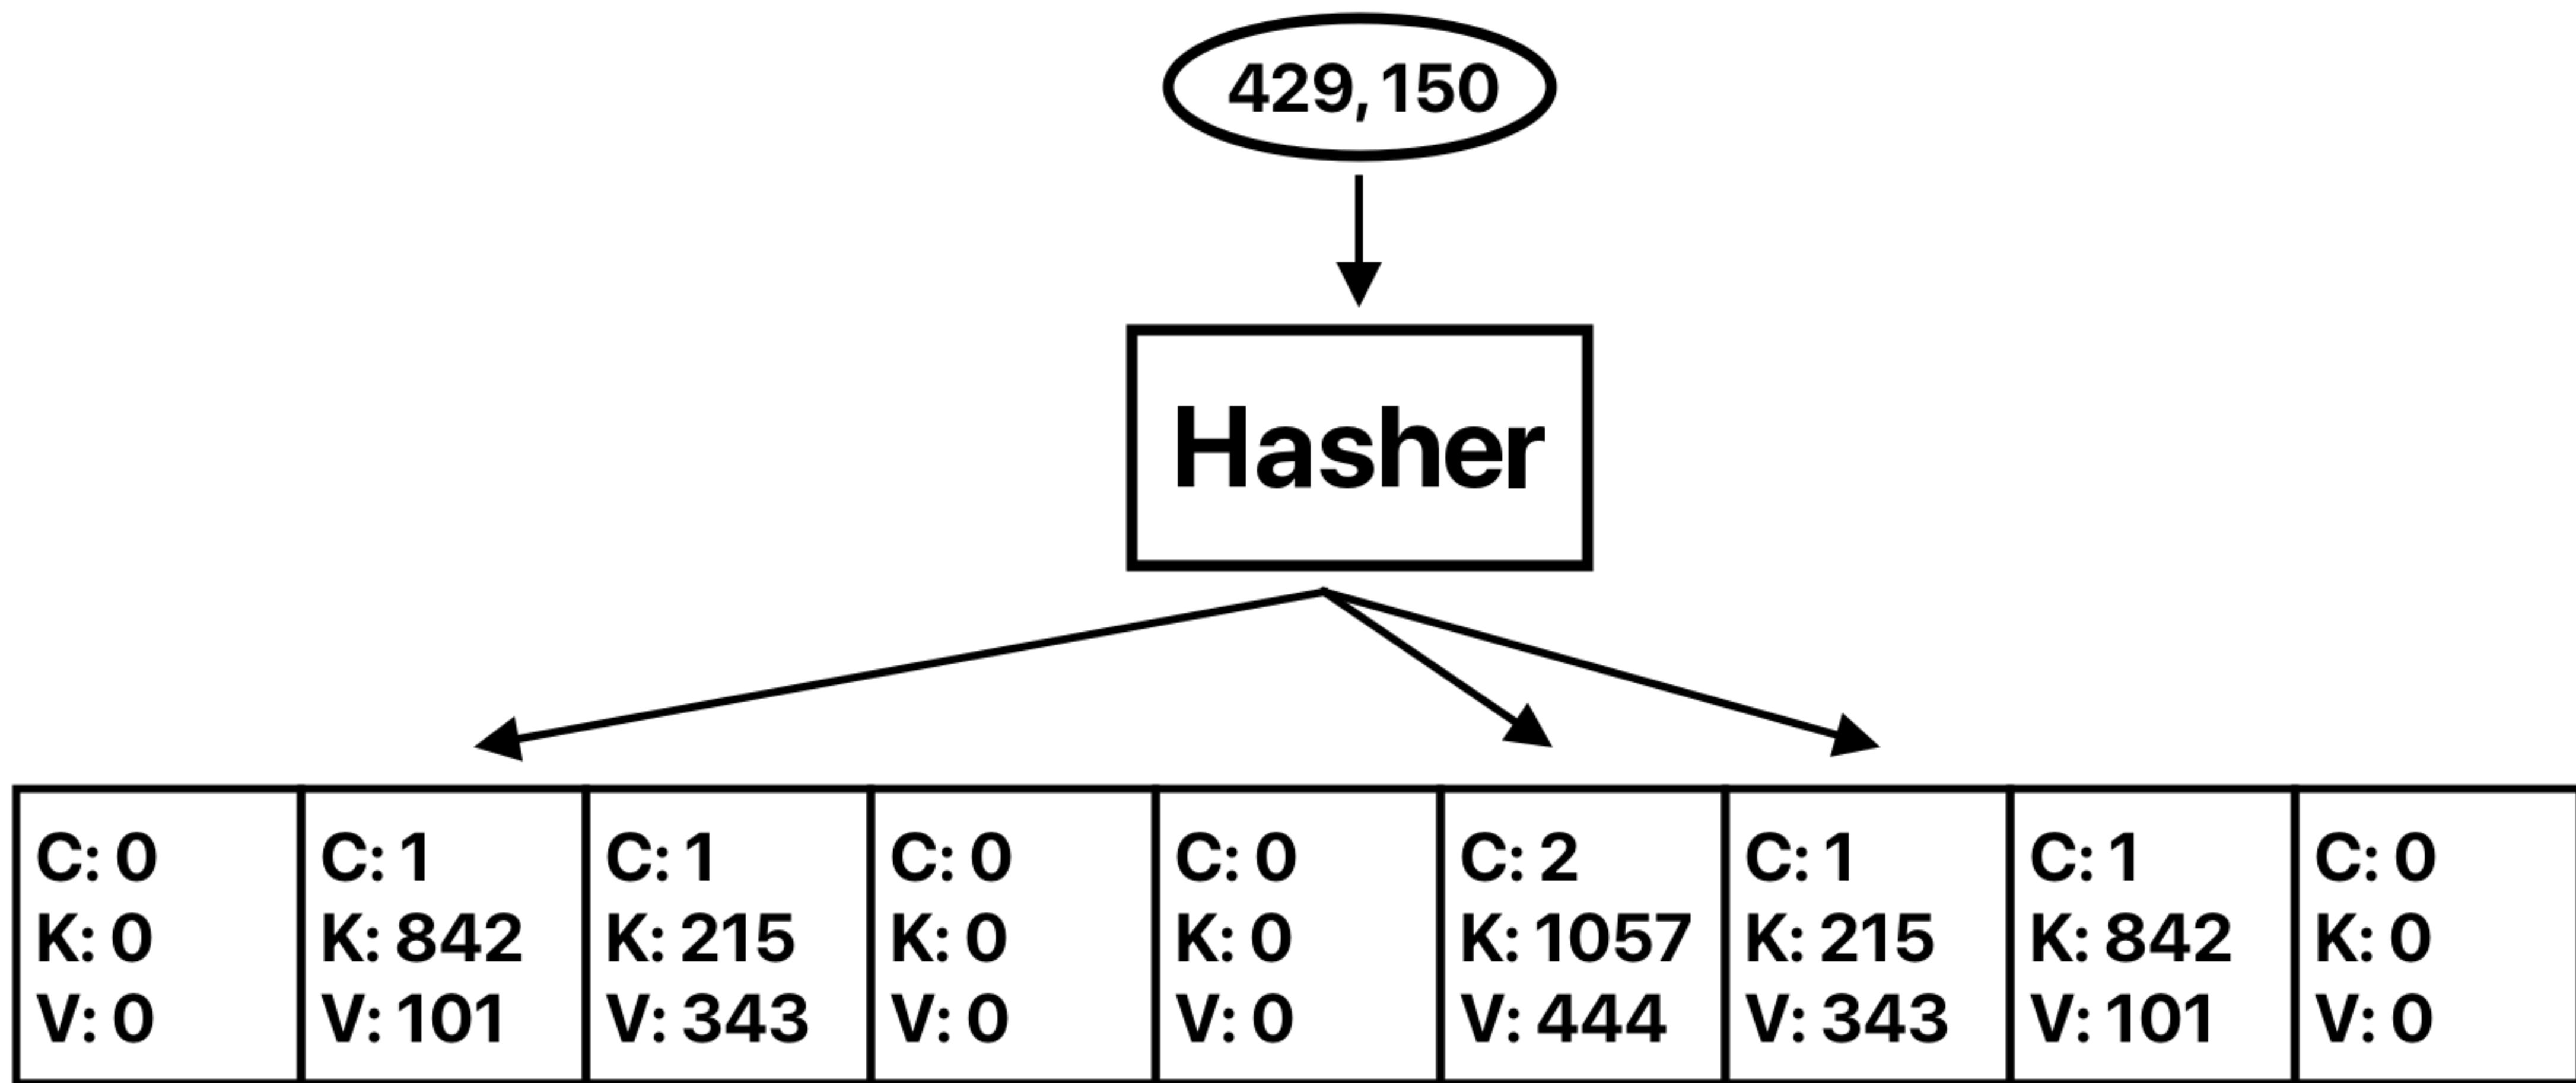
Original IBLT - Insert(x, y)



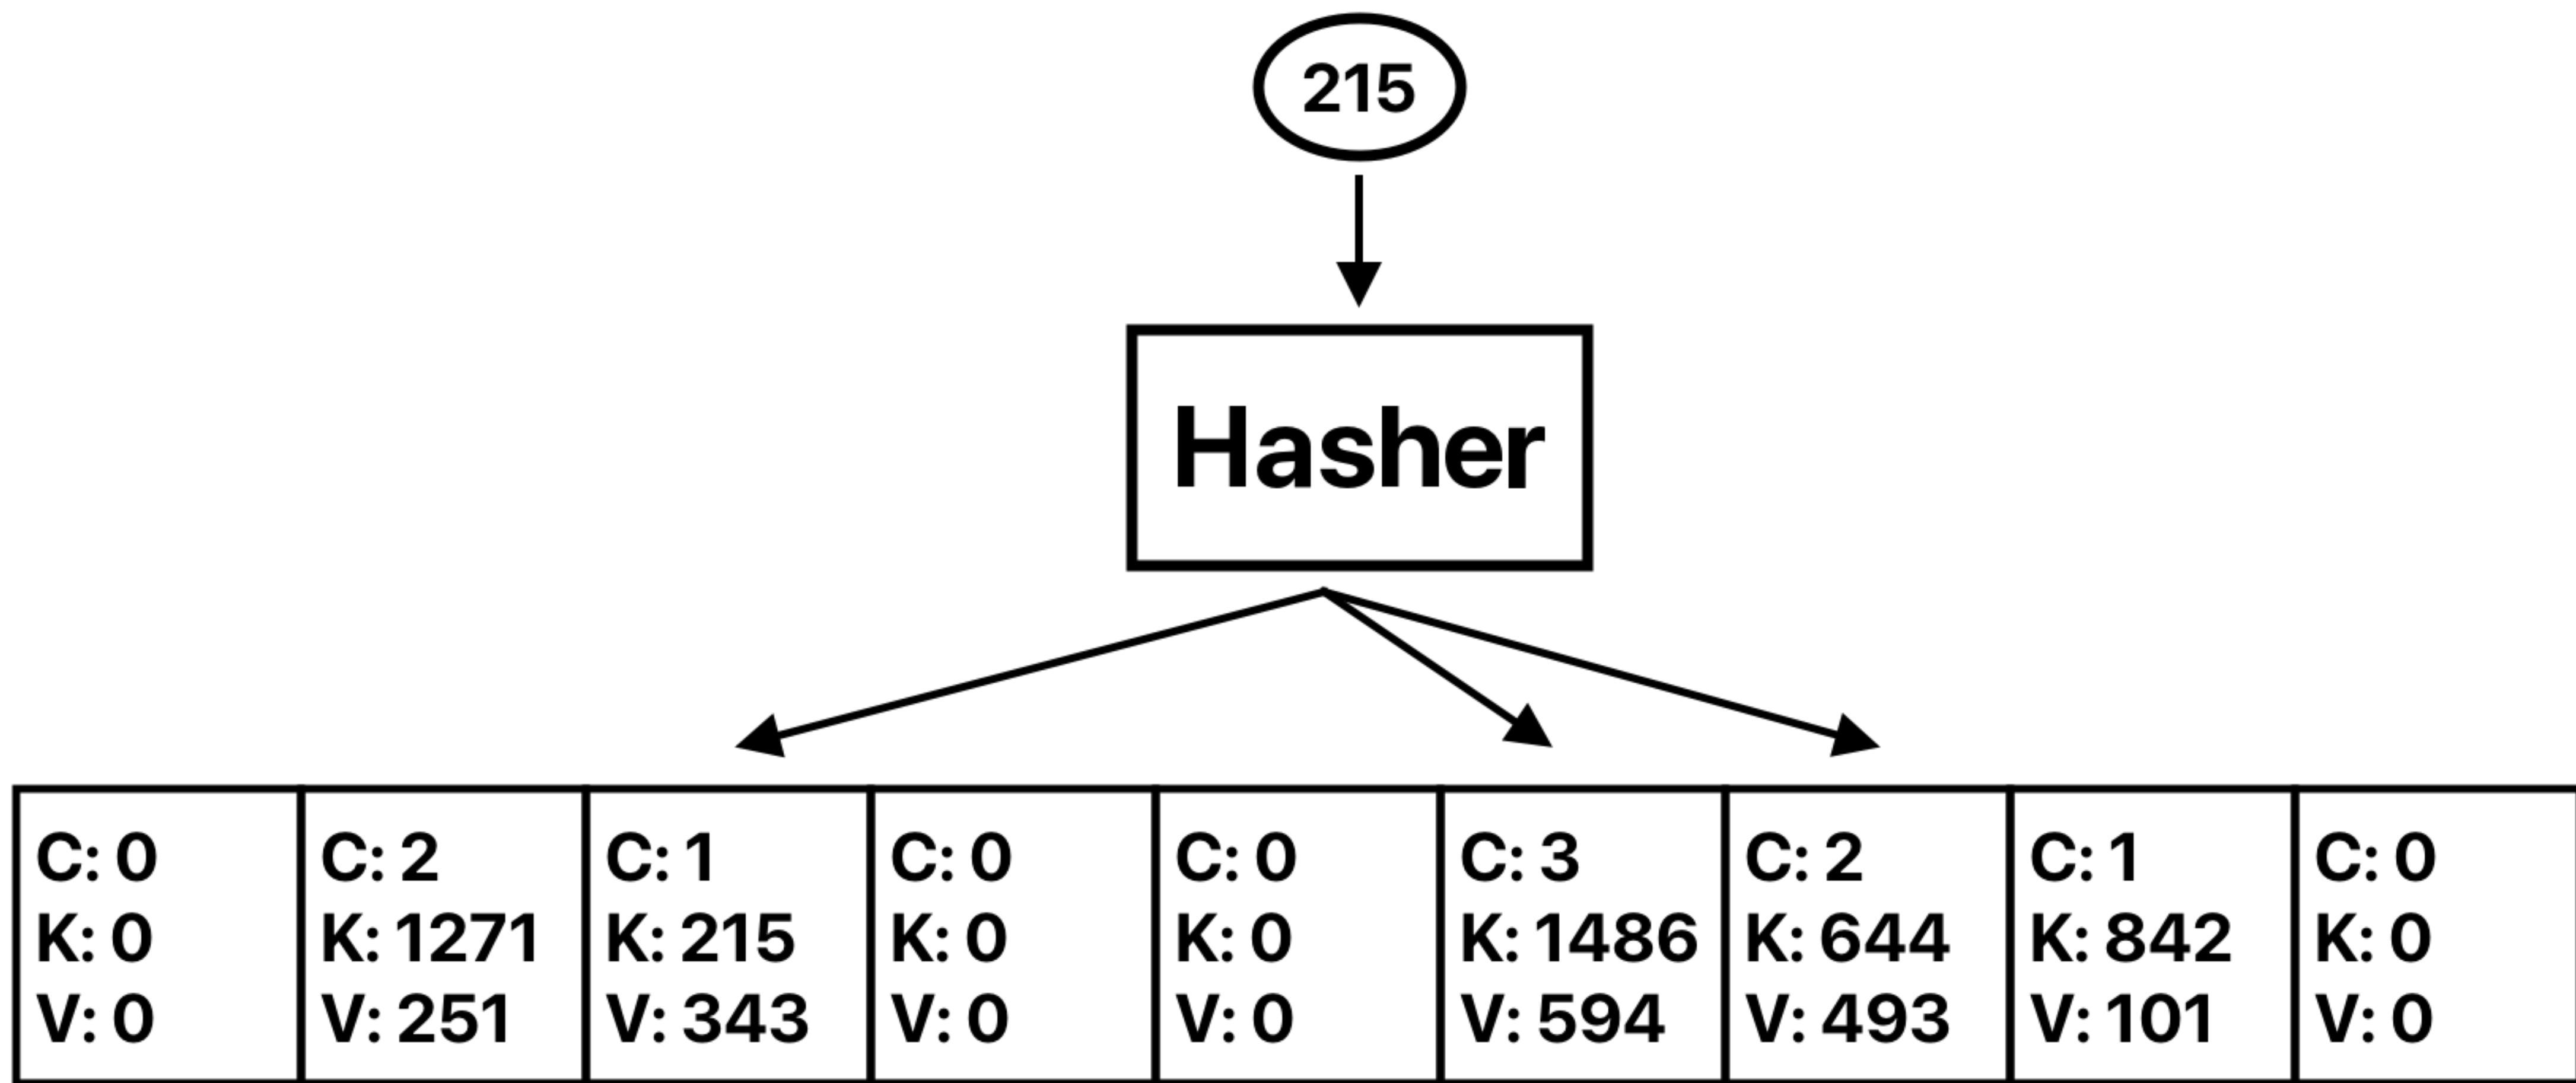
Original IBLT - Delete(x, y)



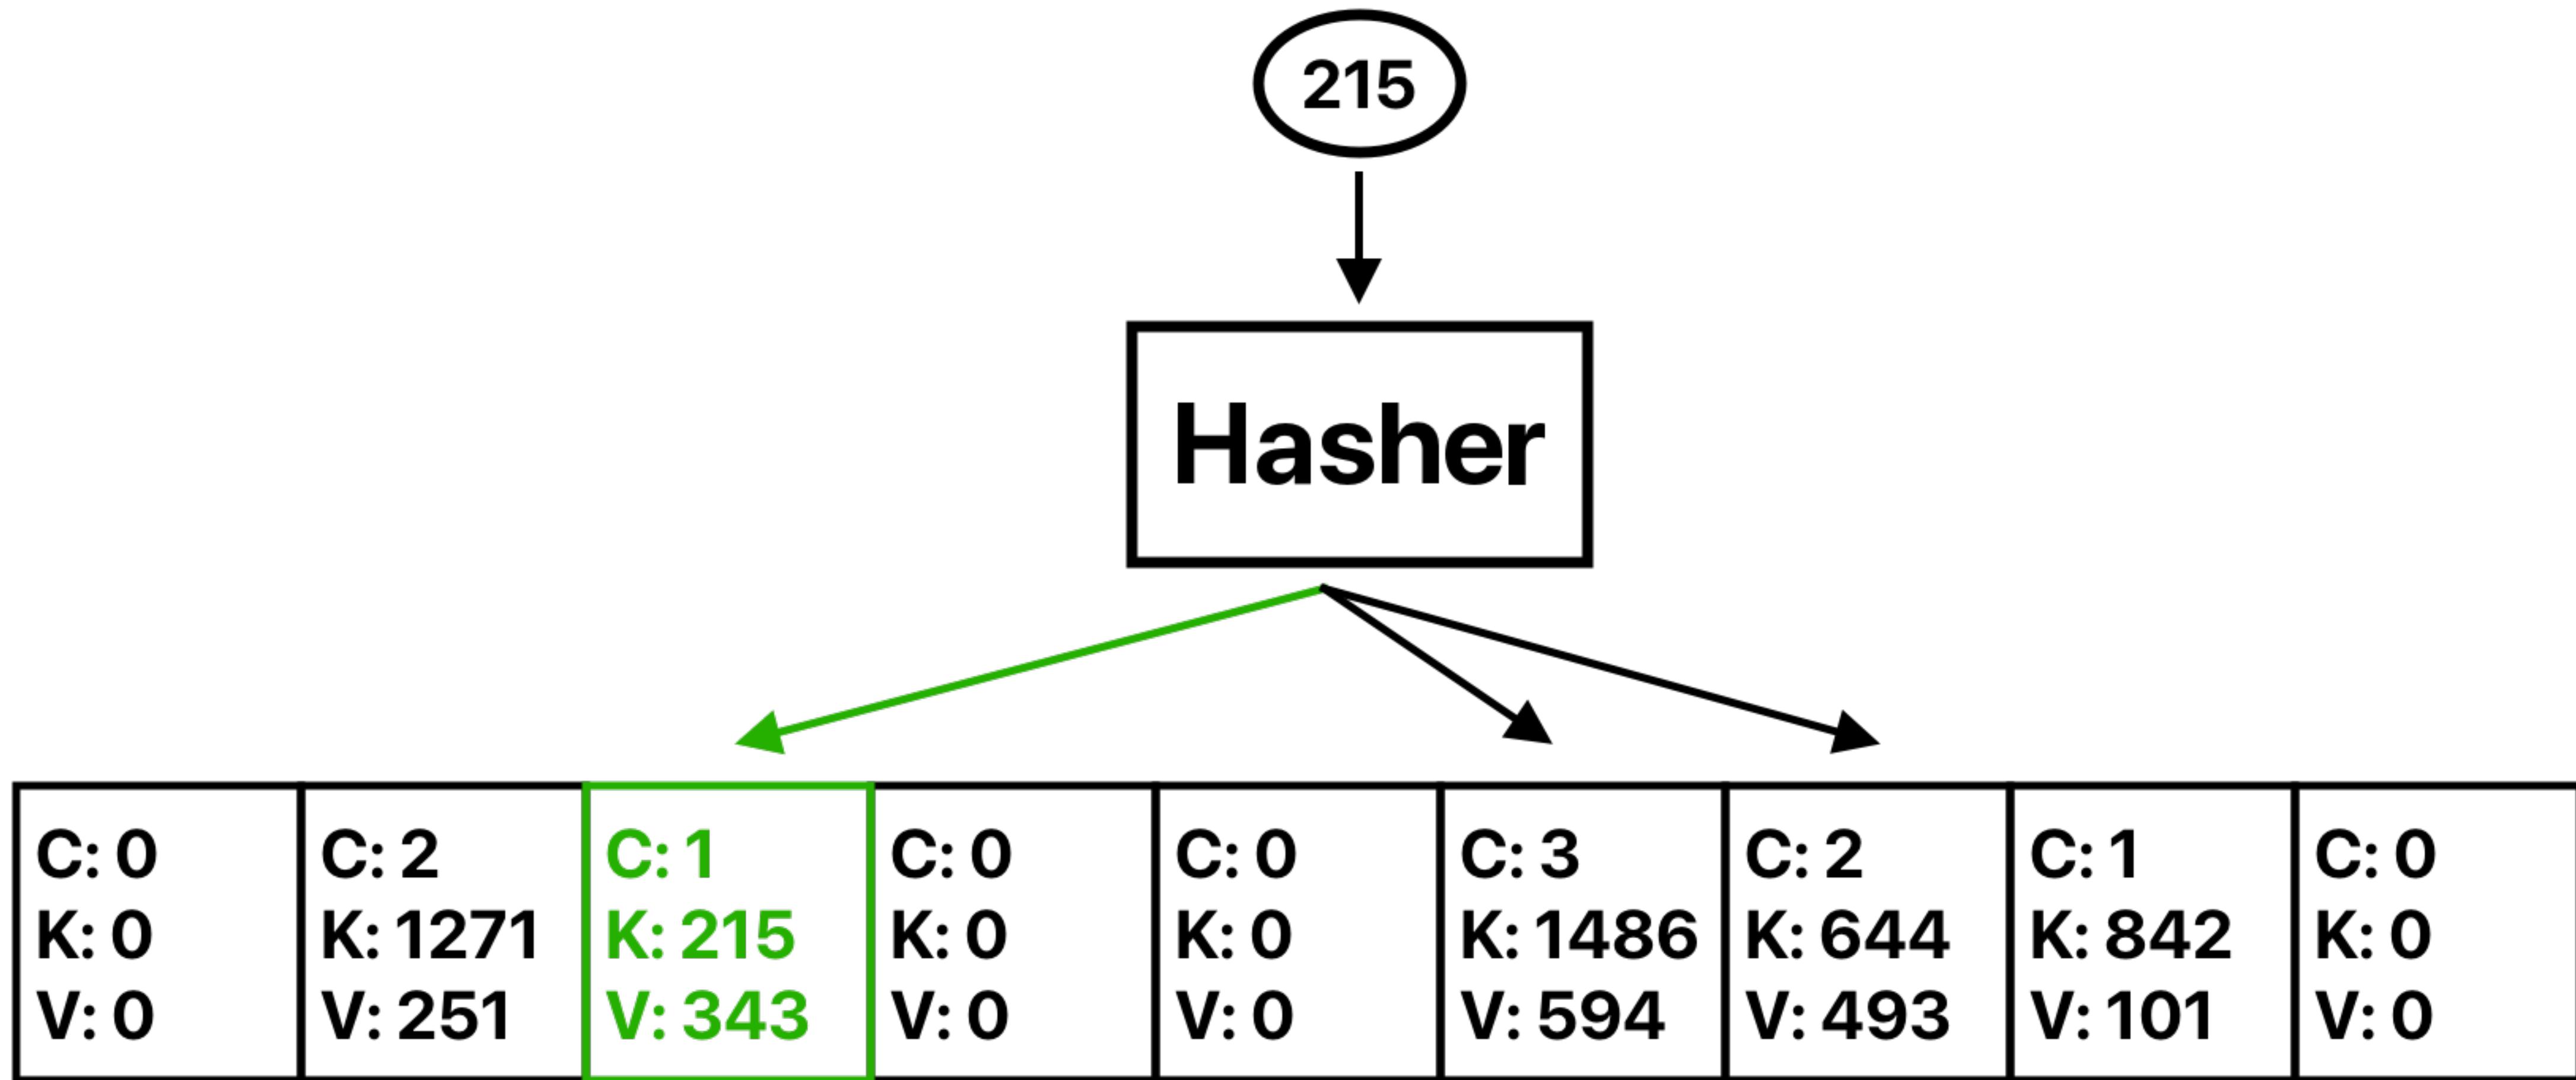
Original IBLT - Delete(x, y)



Original IBLT - Get(x)



Original IBLT - Get(x)



Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

C: 0	C: 2	C: 1	C: 0	C: 0	C: 3	C: 2	C: 1	C: 0
K: 0	K: 1271	K: 215	K: 0	K: 0	K: 1486	K: 644	K: 842	K: 0
V: 0	V: 251	V: 343	V: 0	V: 0	V: 594	V: 493	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

C: 0	C: 2	C: 1	C: 0	C: 0	C: 3	C: 2	C: 1	C: 0
K: 0	K: 1271	K: 215	K: 0	K: 0	K: 1486	K: 644	K: 842	K: 0
V: 0	V: 251	V: 343	V: 0	V: 0	V: 594	V: 493	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

C: 0	C: 2	C: 1	C: 0	C: 0	C: 3	C: 2	C: 1	C: 0
K: 0	K: 1271	K: 215	K: 0	K: 0	K: 1486	K: 644	K: 842	K: 0
V: 0	V: 251	V: 343	V: 0	V: 0	V: 594	V: 493	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

C: 0	C: 2	C: 1	C: 0	C: 0	C: 3	C: 2	C: 1	C: 0
K: 0	K: 1271	K: 215	K: 0	K: 0	K: 1486	K: 644	K: 842	K: 0
V: 0	V: 251	V: 343	V: 0	V: 0	V: 594	V: 493	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

C: 0	C: 2	C: 0	C: 0	C: 0	C: 2	C: 1	C: 1	C: 0
K: 0	K: 1271	K: 0	K: 0	K: 0	K: 1271	K: 429	K: 842	K: 0
V: 0	V: 251	V: 0	V: 0	V: 0	V: 251	V: 150	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

C: 0	C: 2	C: 0	C: 0	C: 0	C: 2	C: 1	C: 1	C: 0
K: 0	K: 1271	K: 0	K: 0	K: 0	K: 1271	K: 429	K: 842	K: 0
V: 0	V: 251	V: 0	V: 0	V: 0	V: 251	V: 150	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

(842, 101)

C: 0	C: 2	C: 0	C: 0	C: 0	C: 2	C: 1	C: 1	C: 0
K: 0	K: 1271	K: 0	K: 0	K: 0	K: 1271	K: 429	K: 842	K: 0
V: 0	V: 251	V: 0	V: 0	V: 0	V: 251	V: 150	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)
(842, 101)

C: 0	C: 2	C: 0	C: 0	C: 0	C: 2	C: 1	C: 1	C: 0
K: 0	K: 1271	K: 0	K: 0	K: 0	K: 1271	K: 429	K: 842	K: 0
V: 0	V: 251	V: 0	V: 0	V: 0	V: 251	V: 150	V: 101	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

(842, 101)

C: 0	C: 1	C: 0	C: 0	C: 0	C: 1	C: 1	C: 0	C: 0
K: 0	K: 429	K: 0	K: 0	K: 0	K: 429	K: 429	K: 0	K: 0
V: 0	V: 150	V: 0	V: 0	V: 0	V: 150	V: 150	V: 0	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)

(842, 101)

C: 0	C: 1	C: 0	C: 0	C: 0	C: 1	C: 1	C: 0	C: 0
K: 0	K: 429	K: 0	K: 0	K: 0	K: 429	K: 429	K: 0	K: 0
V: 0	V: 150	V: 0	V: 0	V: 0	V: 150	V: 150	V: 0	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)
(842, 101)
(429, 150)

C: 0	C: 1	C: 0	C: 0	C: 0	C: 1	C: 1	C: 0	C: 0
K: 0	K: 429	K: 0	K: 0	K: 0	K: 429	K: 429	K: 0	K: 0
V: 0	V: 150	V: 0	V: 0	V: 0	V: 150	V: 150	V: 0	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)
(842, 101)
(429, 150)

C: 0	C: 1	C: 0	C: 0	C: 0	C: 1	C: 1	C: 0	C: 0
K: 0	K: 429	K: 0	K: 0	K: 0	K: 429	K: 429	K: 0	K: 0
V: 0	V: 150	V: 0	V: 0	V: 0	V: 150	V: 150	V: 0	V: 0

Original IBLT - ListEntries()

```
while exists cell w/ C=1:  
    add (key, val) to results  
    DELETE(key, val)
```

Result:

(215, 343)
(842, 101)
(429, 150)

C: 0								
K: 0								
V: 0								



Success!

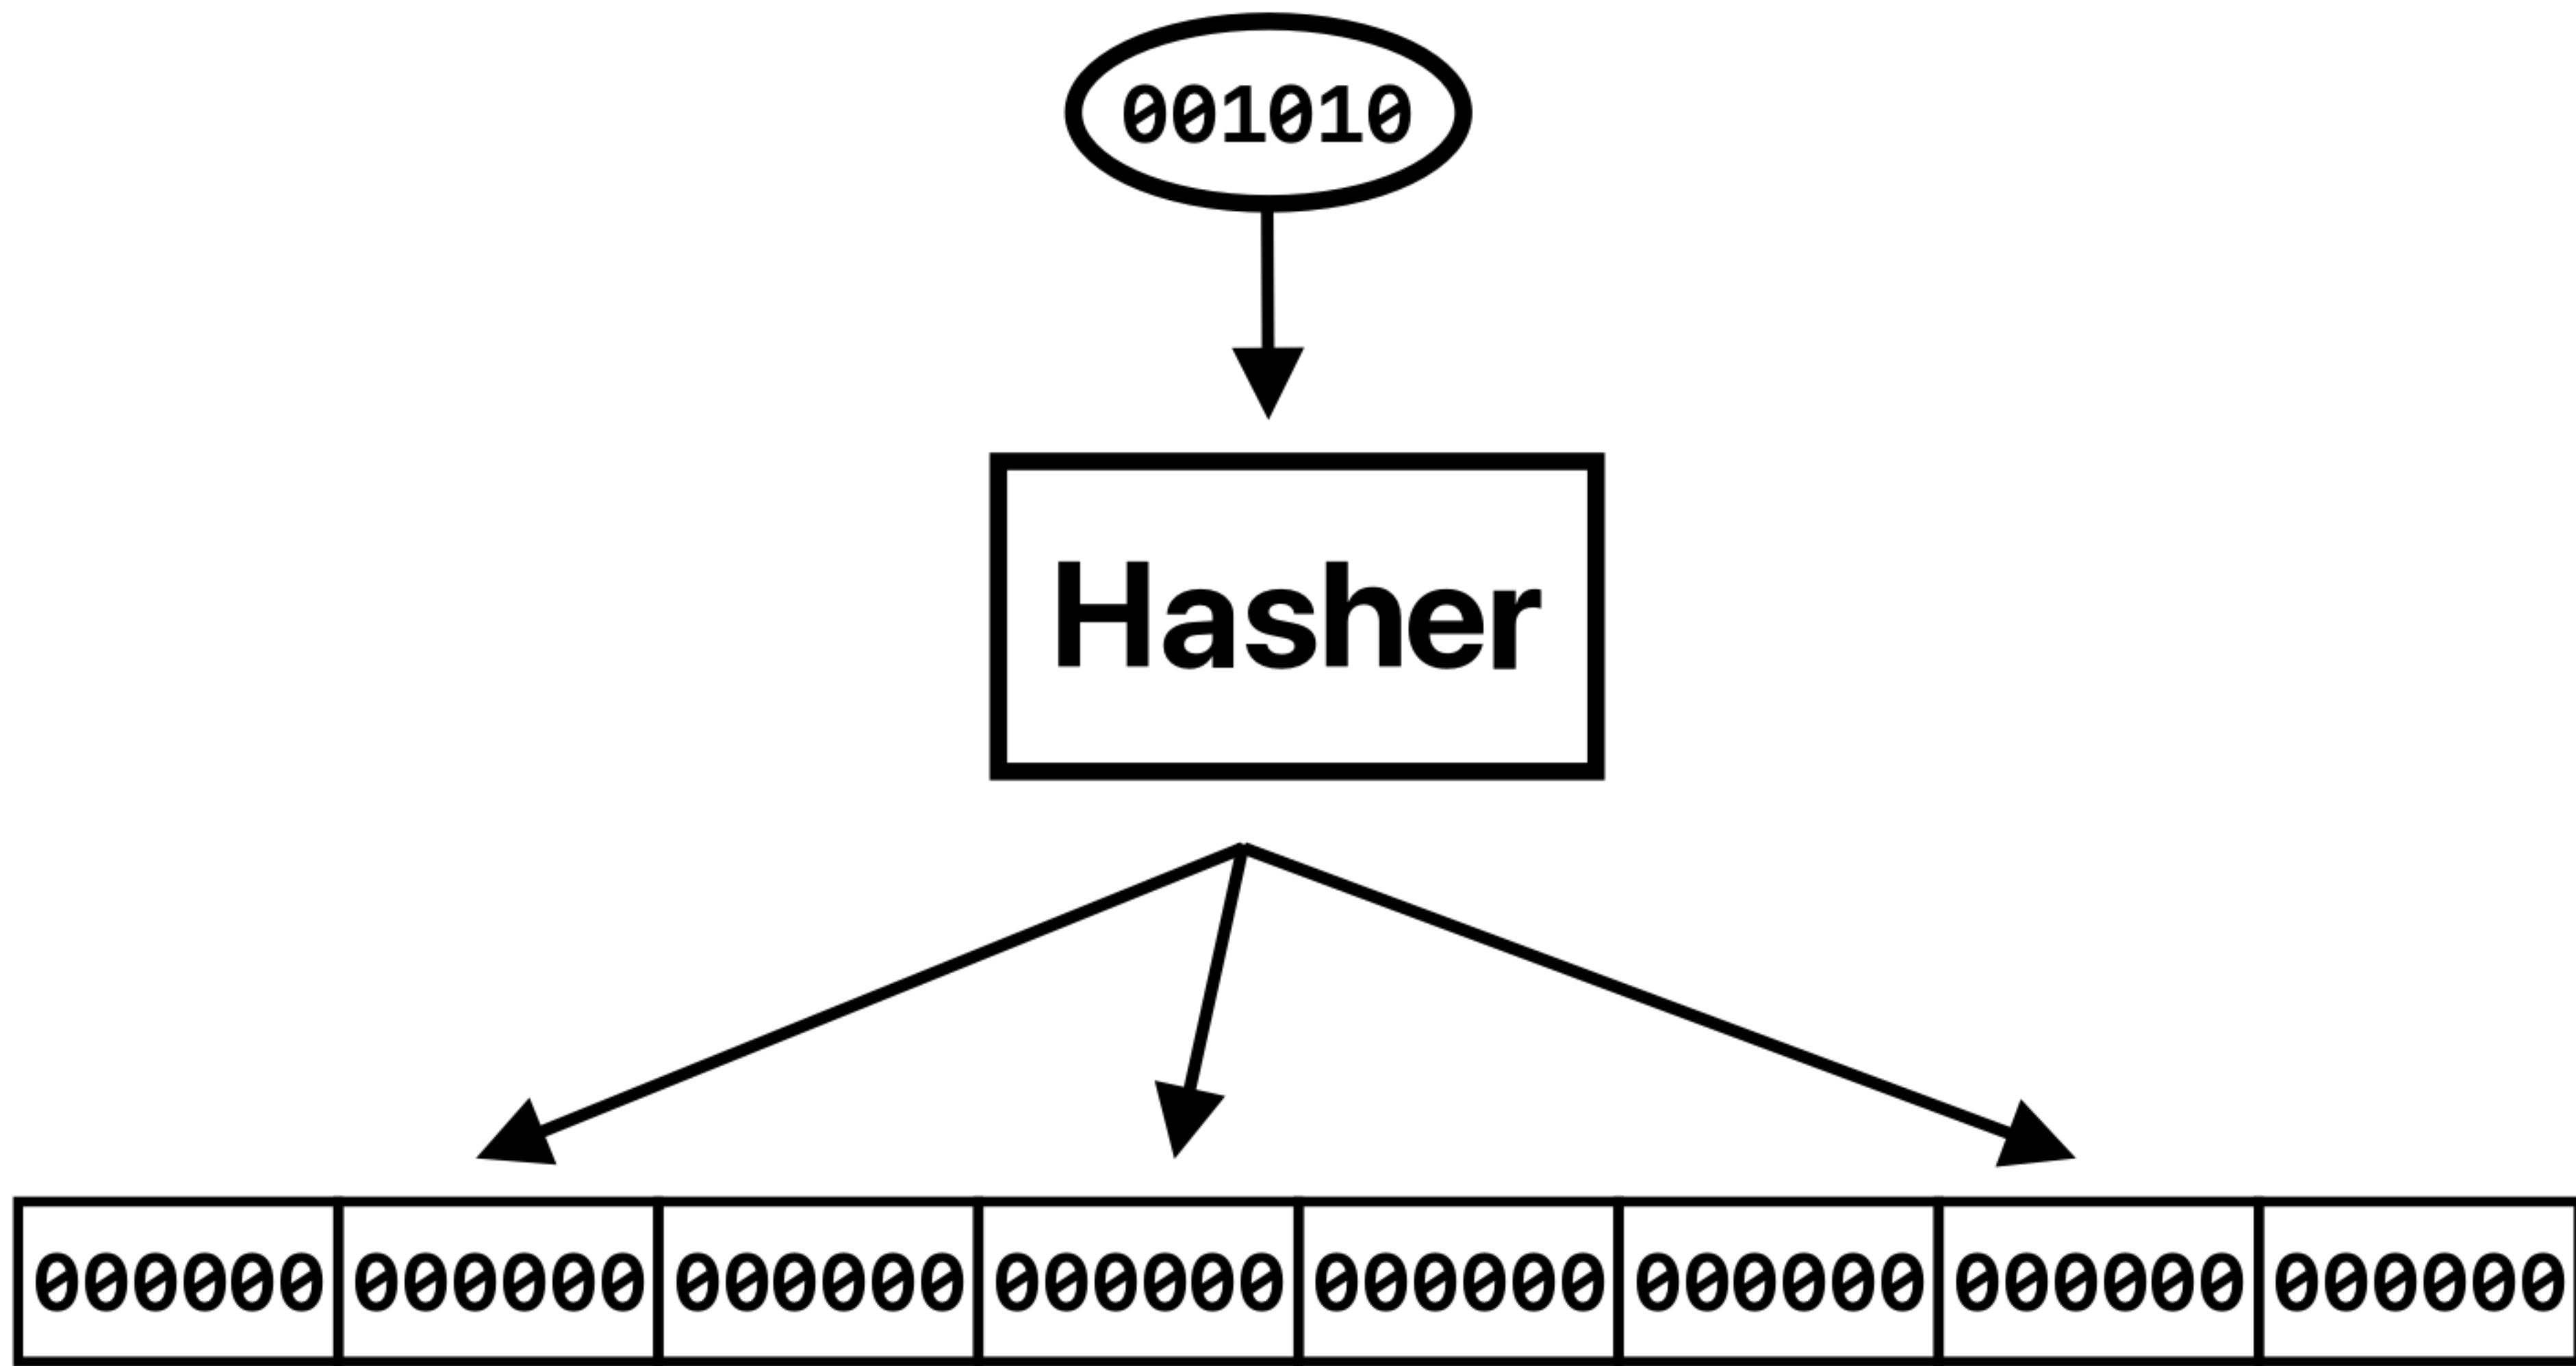
Original IBLT - Limitations

- Extraneous Deletions
- Same Key, Same Values
- Same Key, Different Values

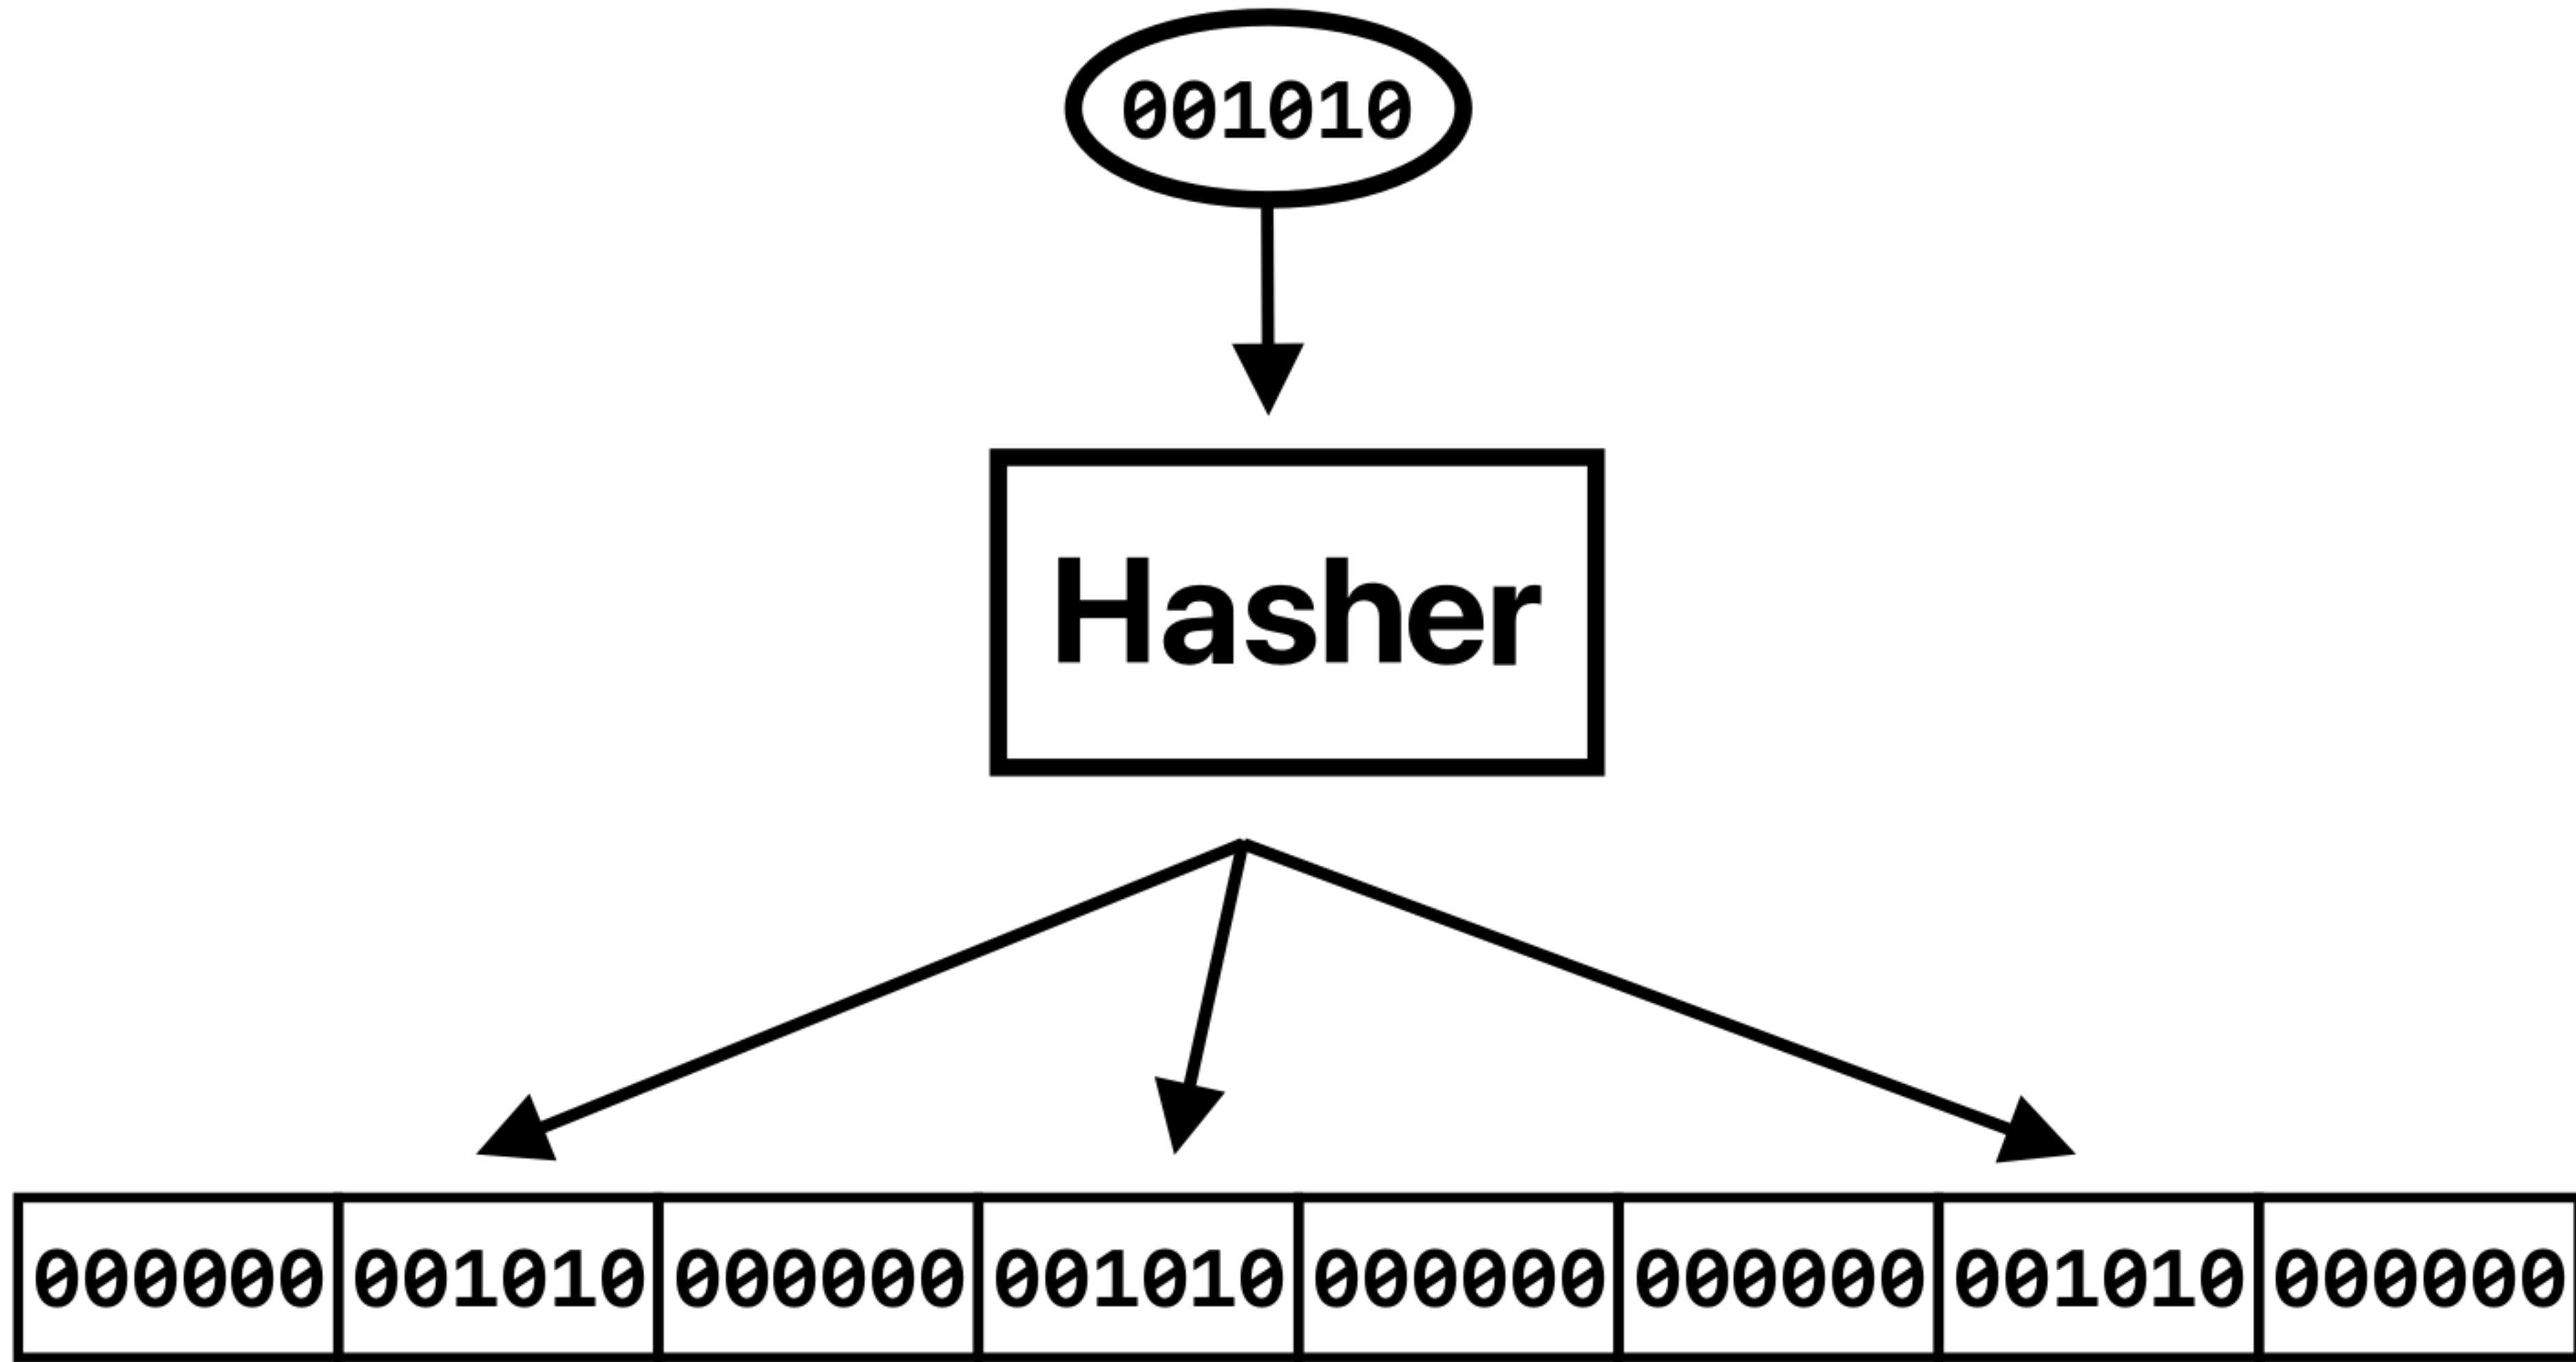
HPW IBLT

- Only stores keys
- Operations:
 - Toggle(x)
 - Merge(other_sketch)
 - ListEntries()

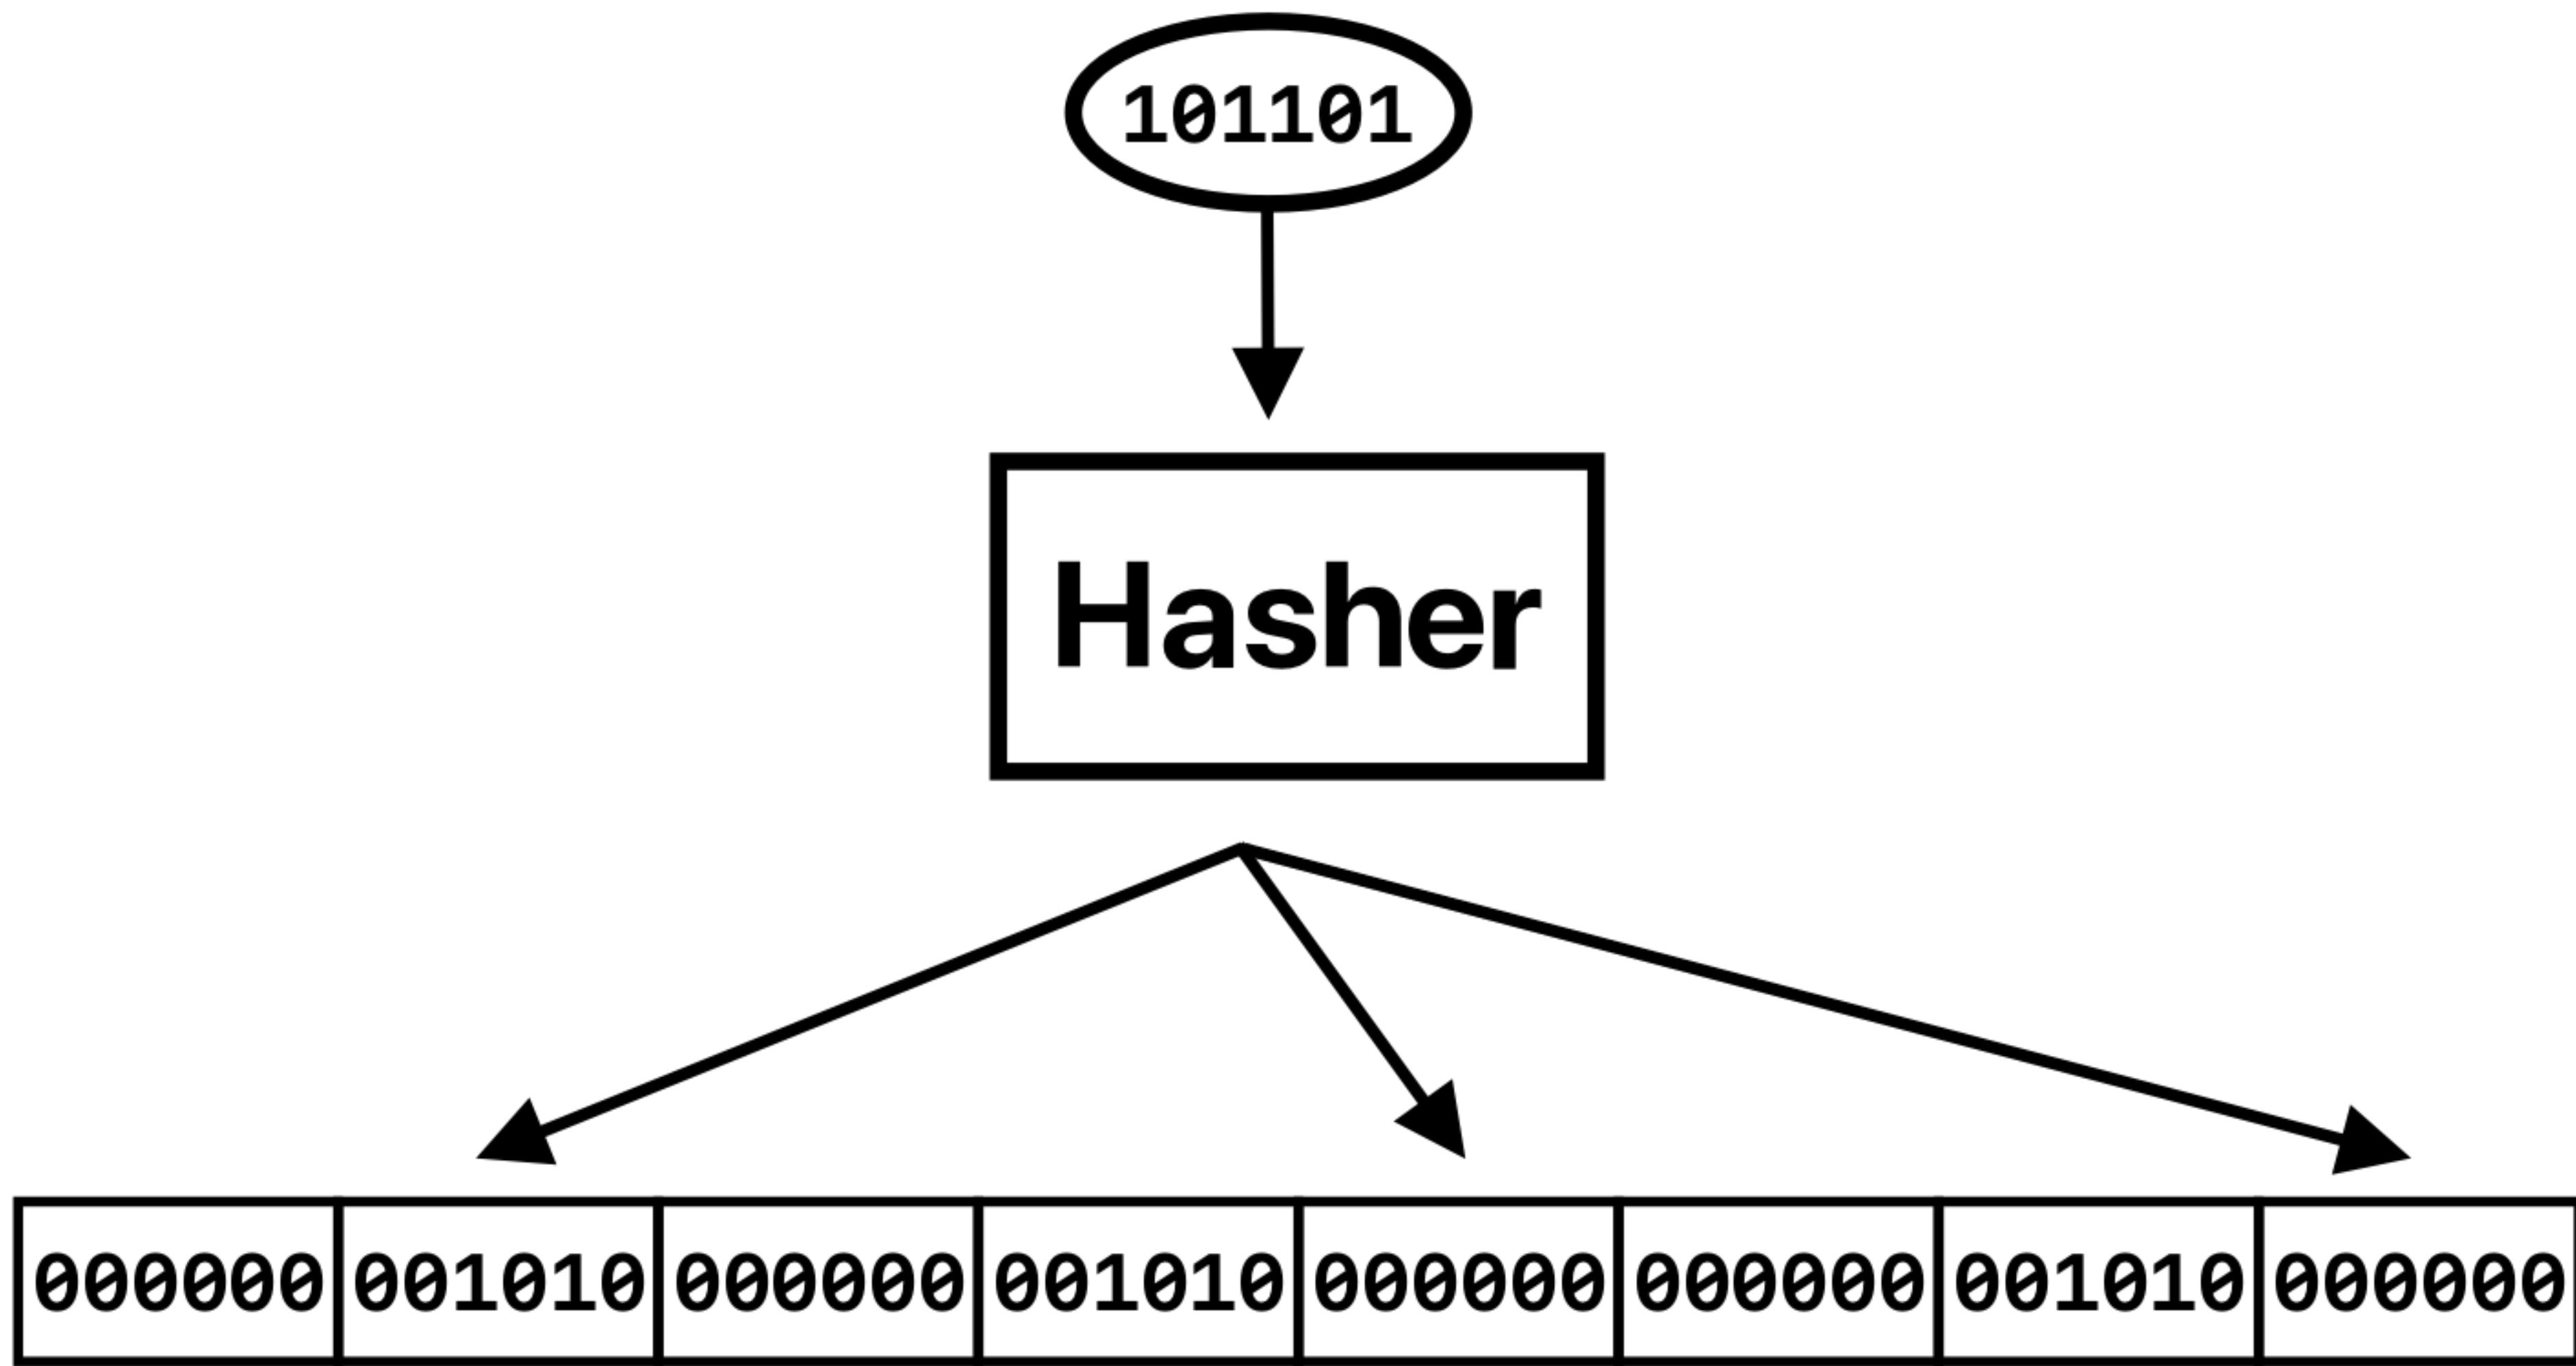
HPW IBLT - Toggle(x)



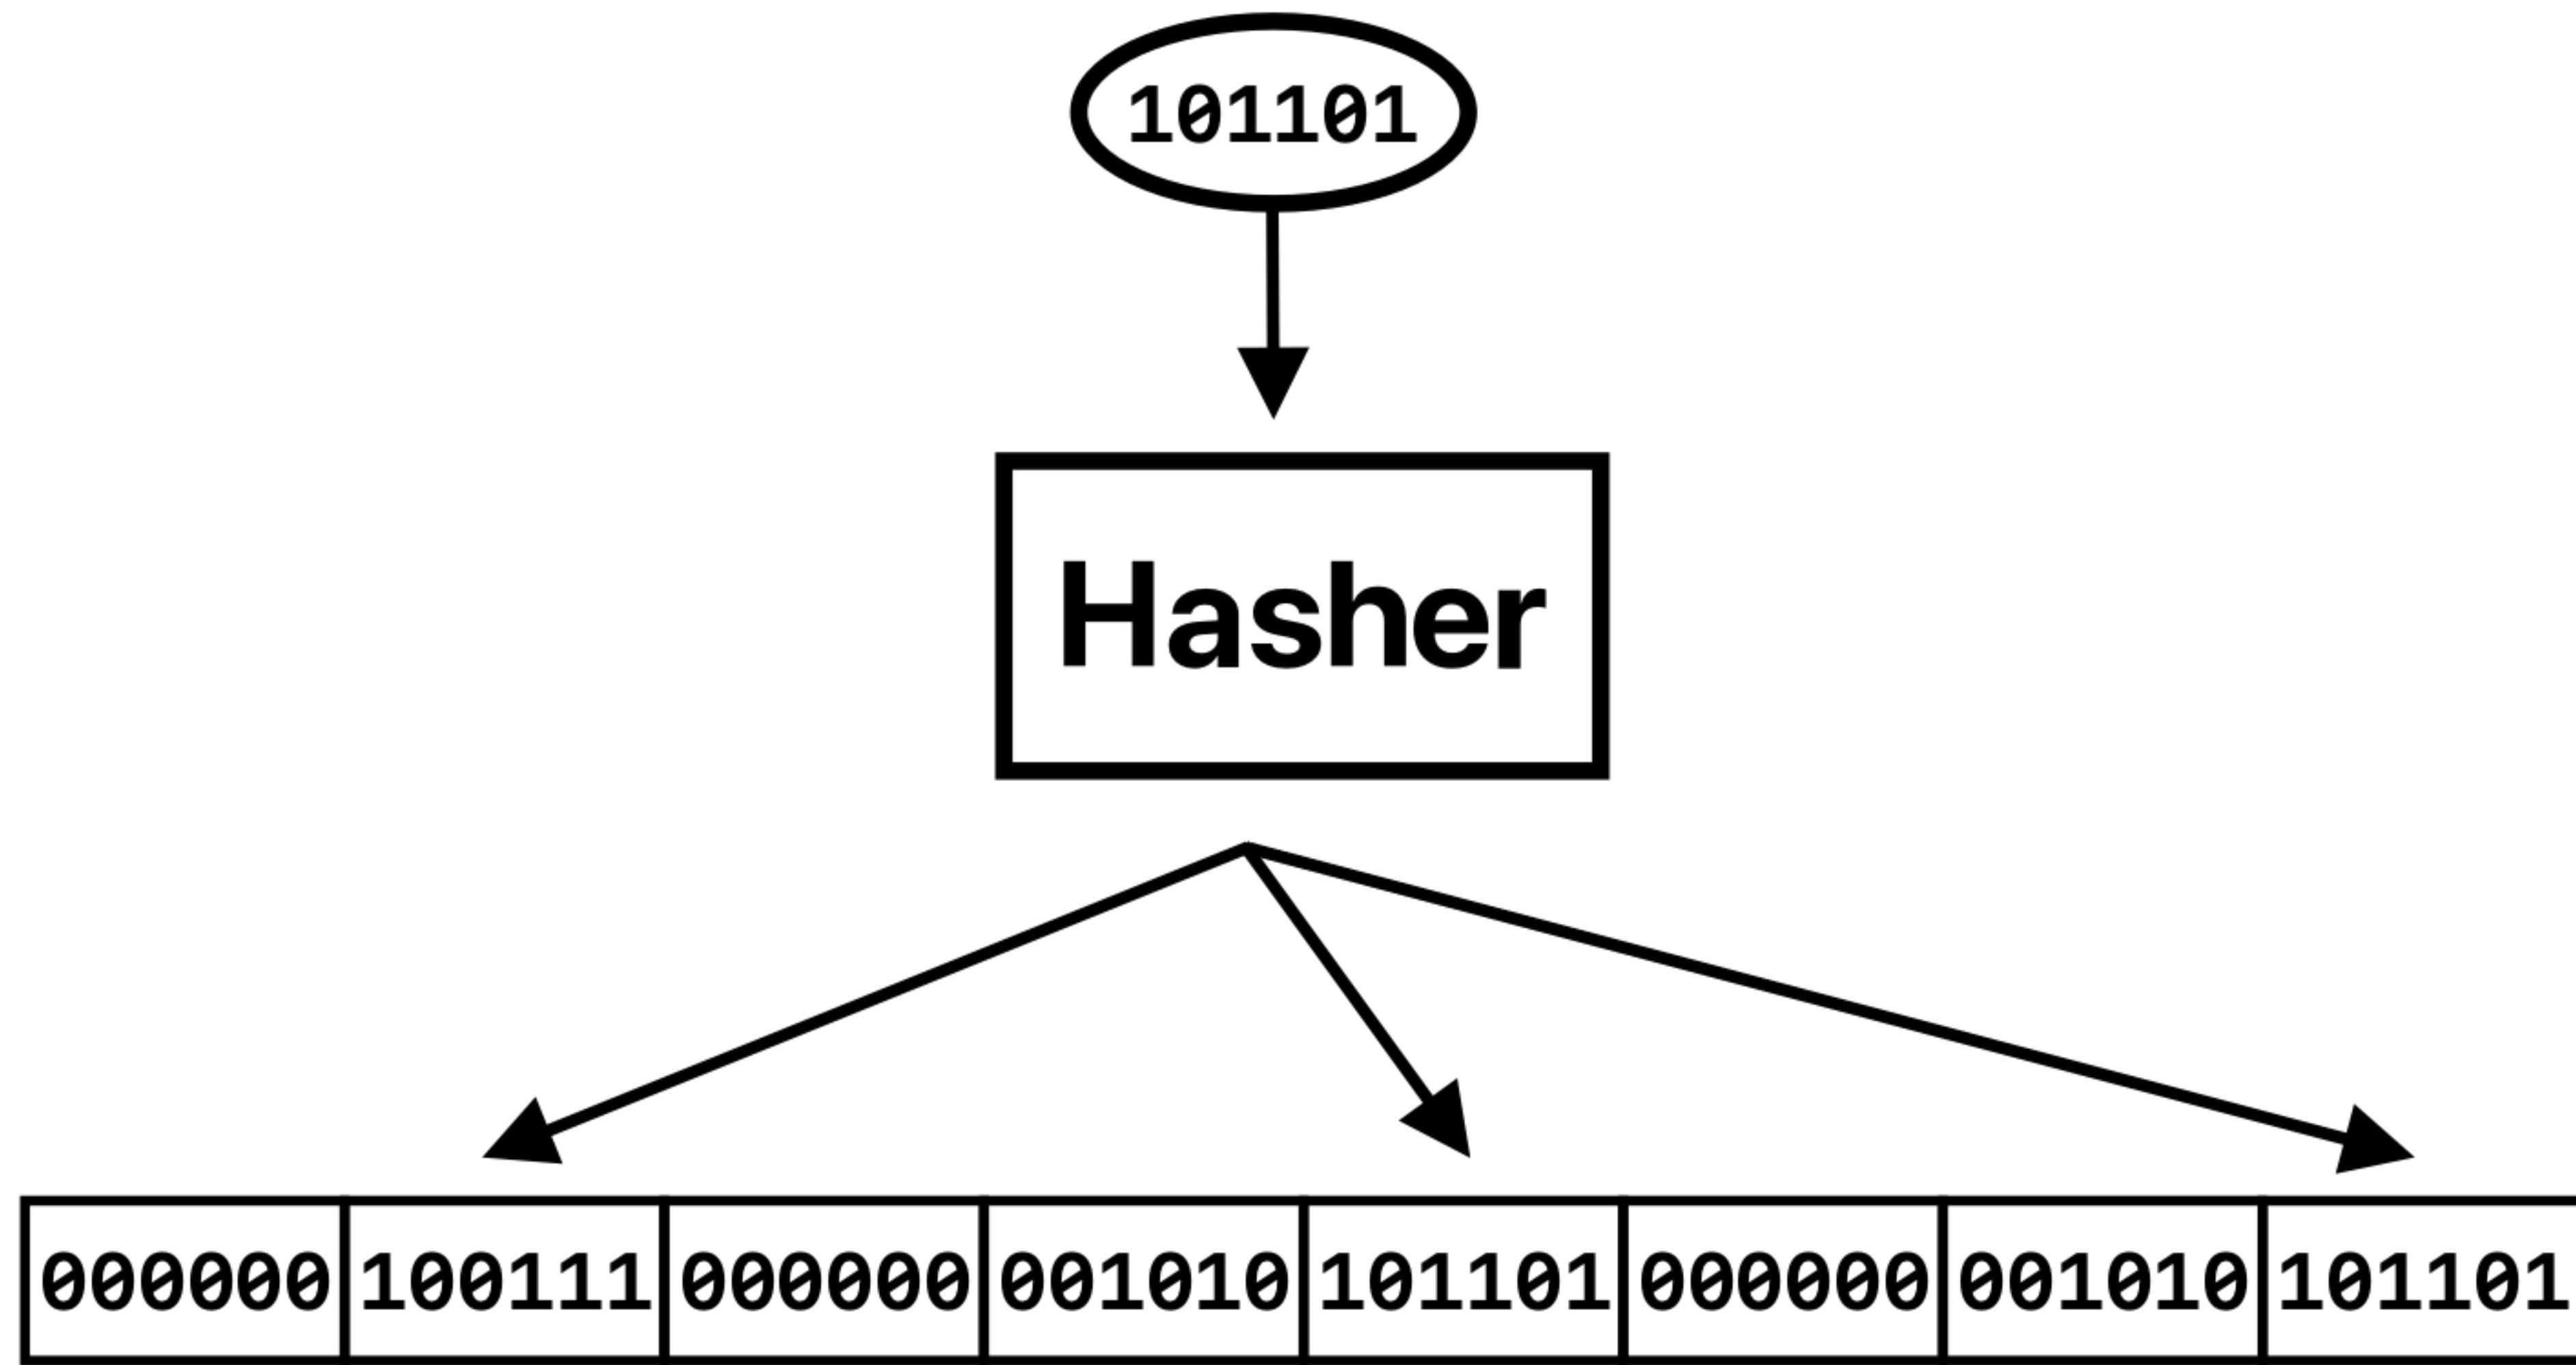
HPW IBLT - Toggle(x)



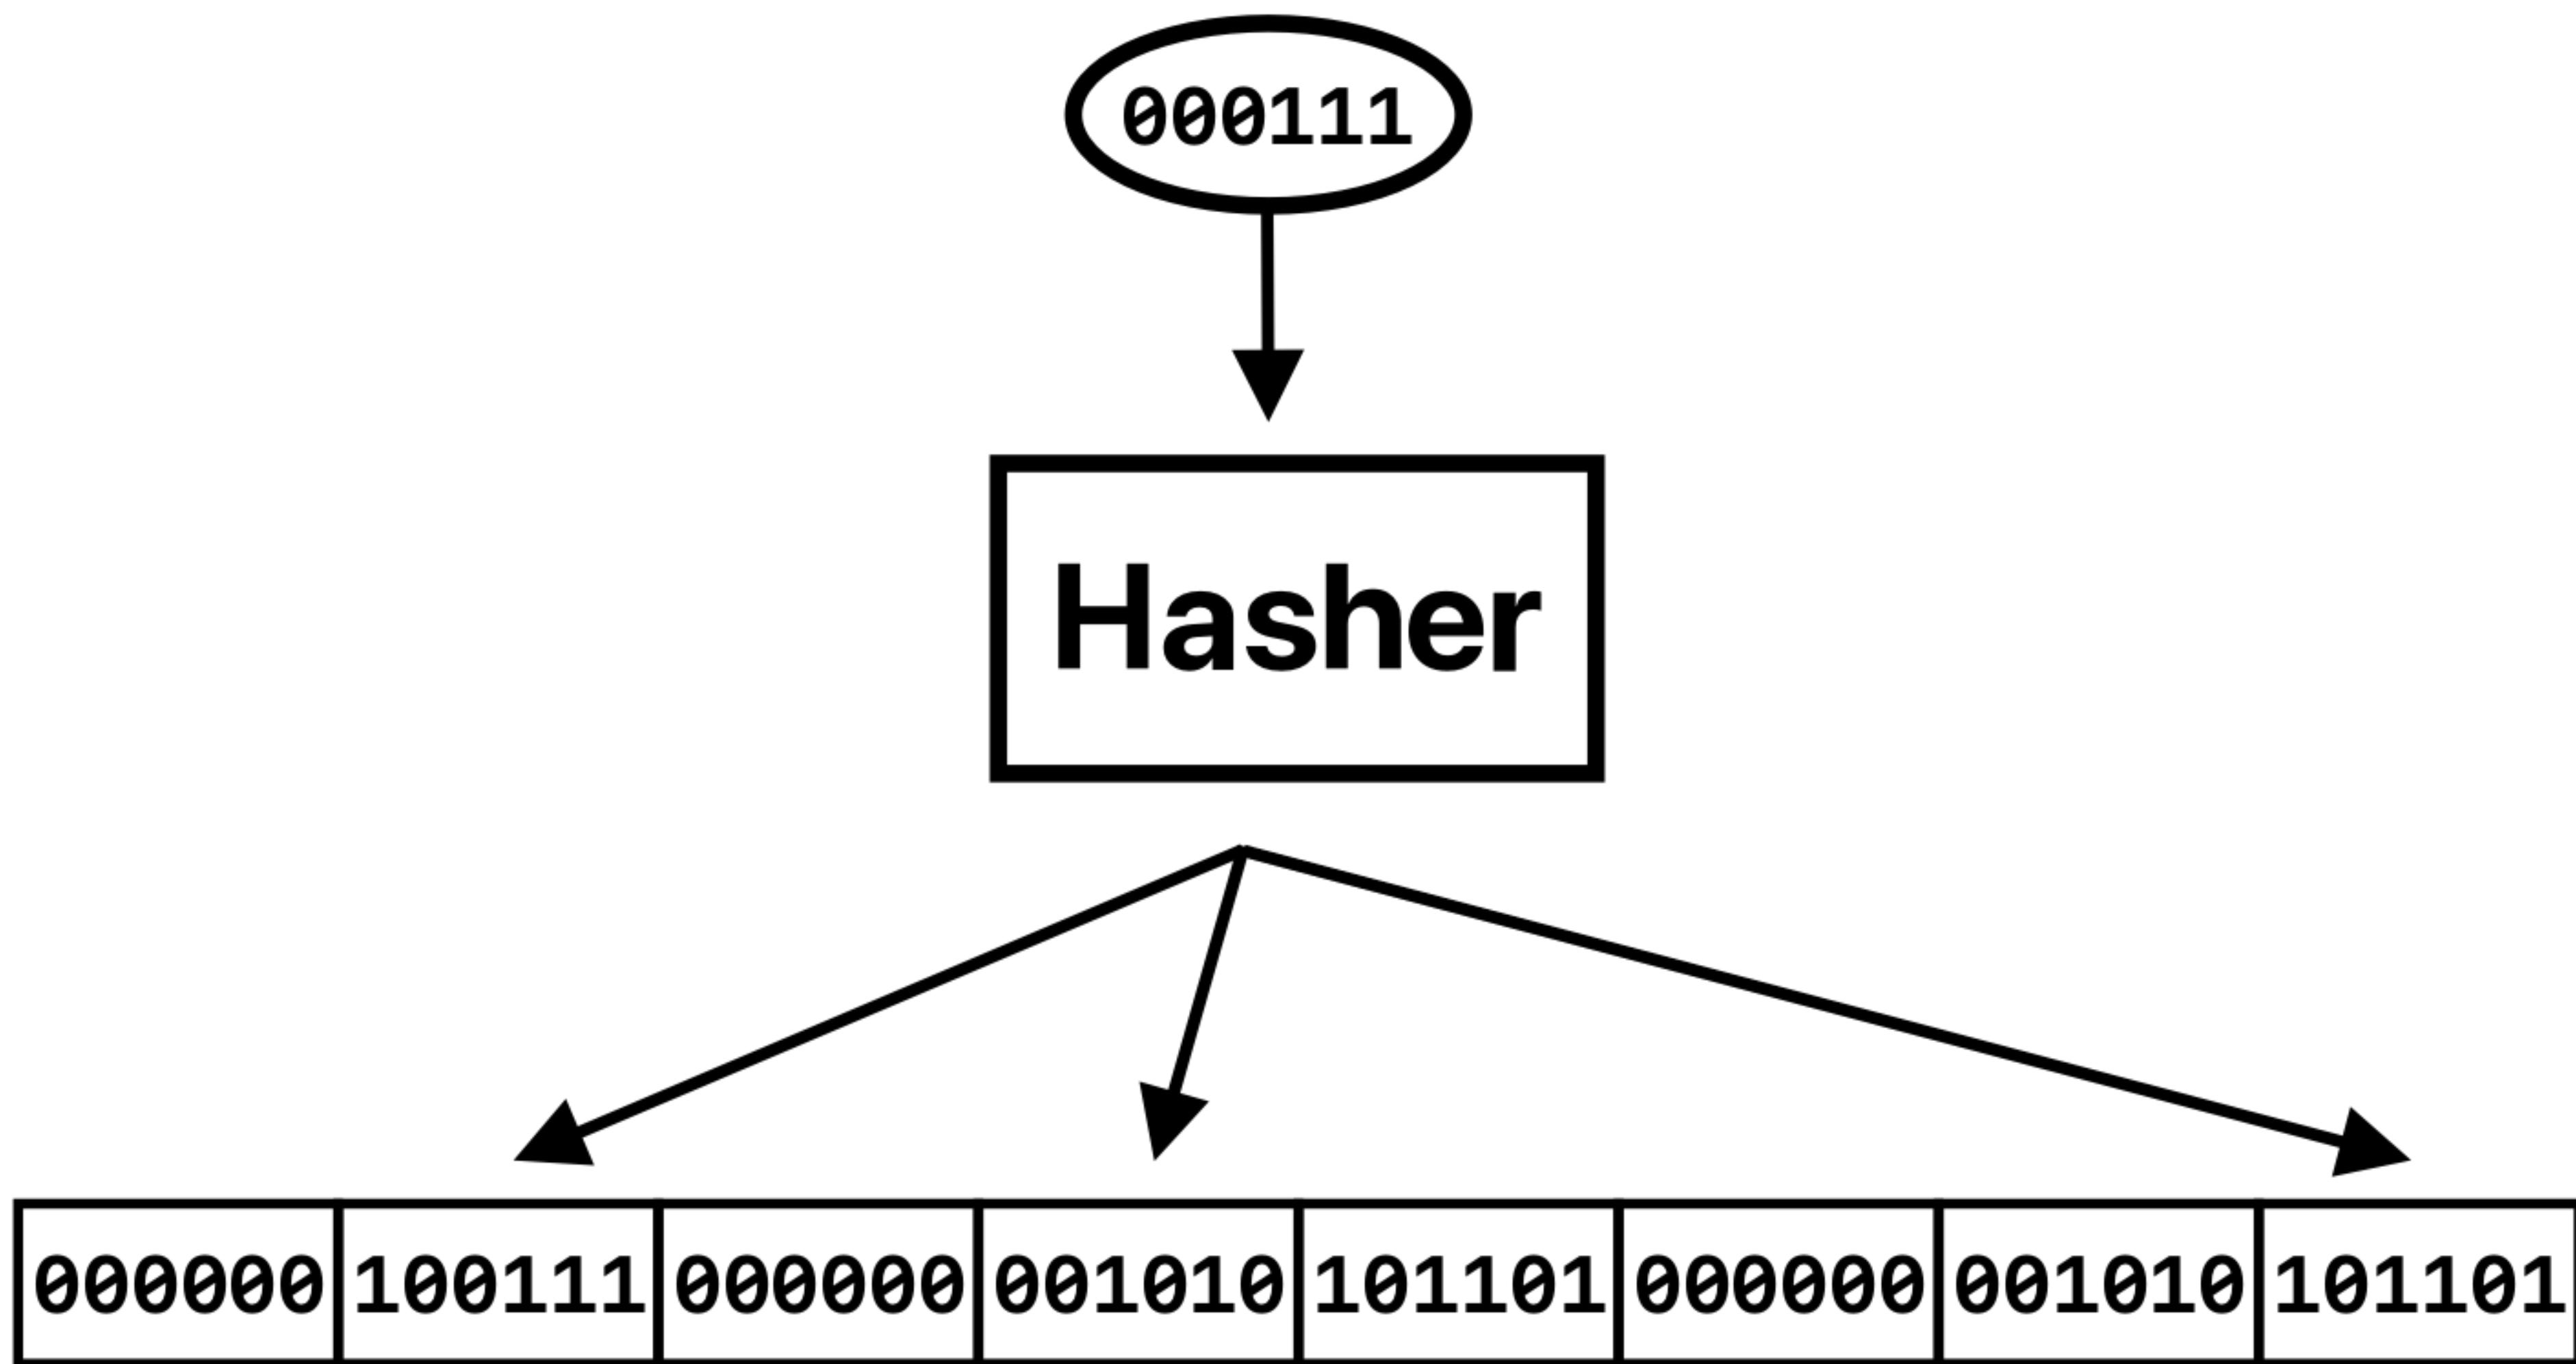
HPW IBLT - Toggle(x)



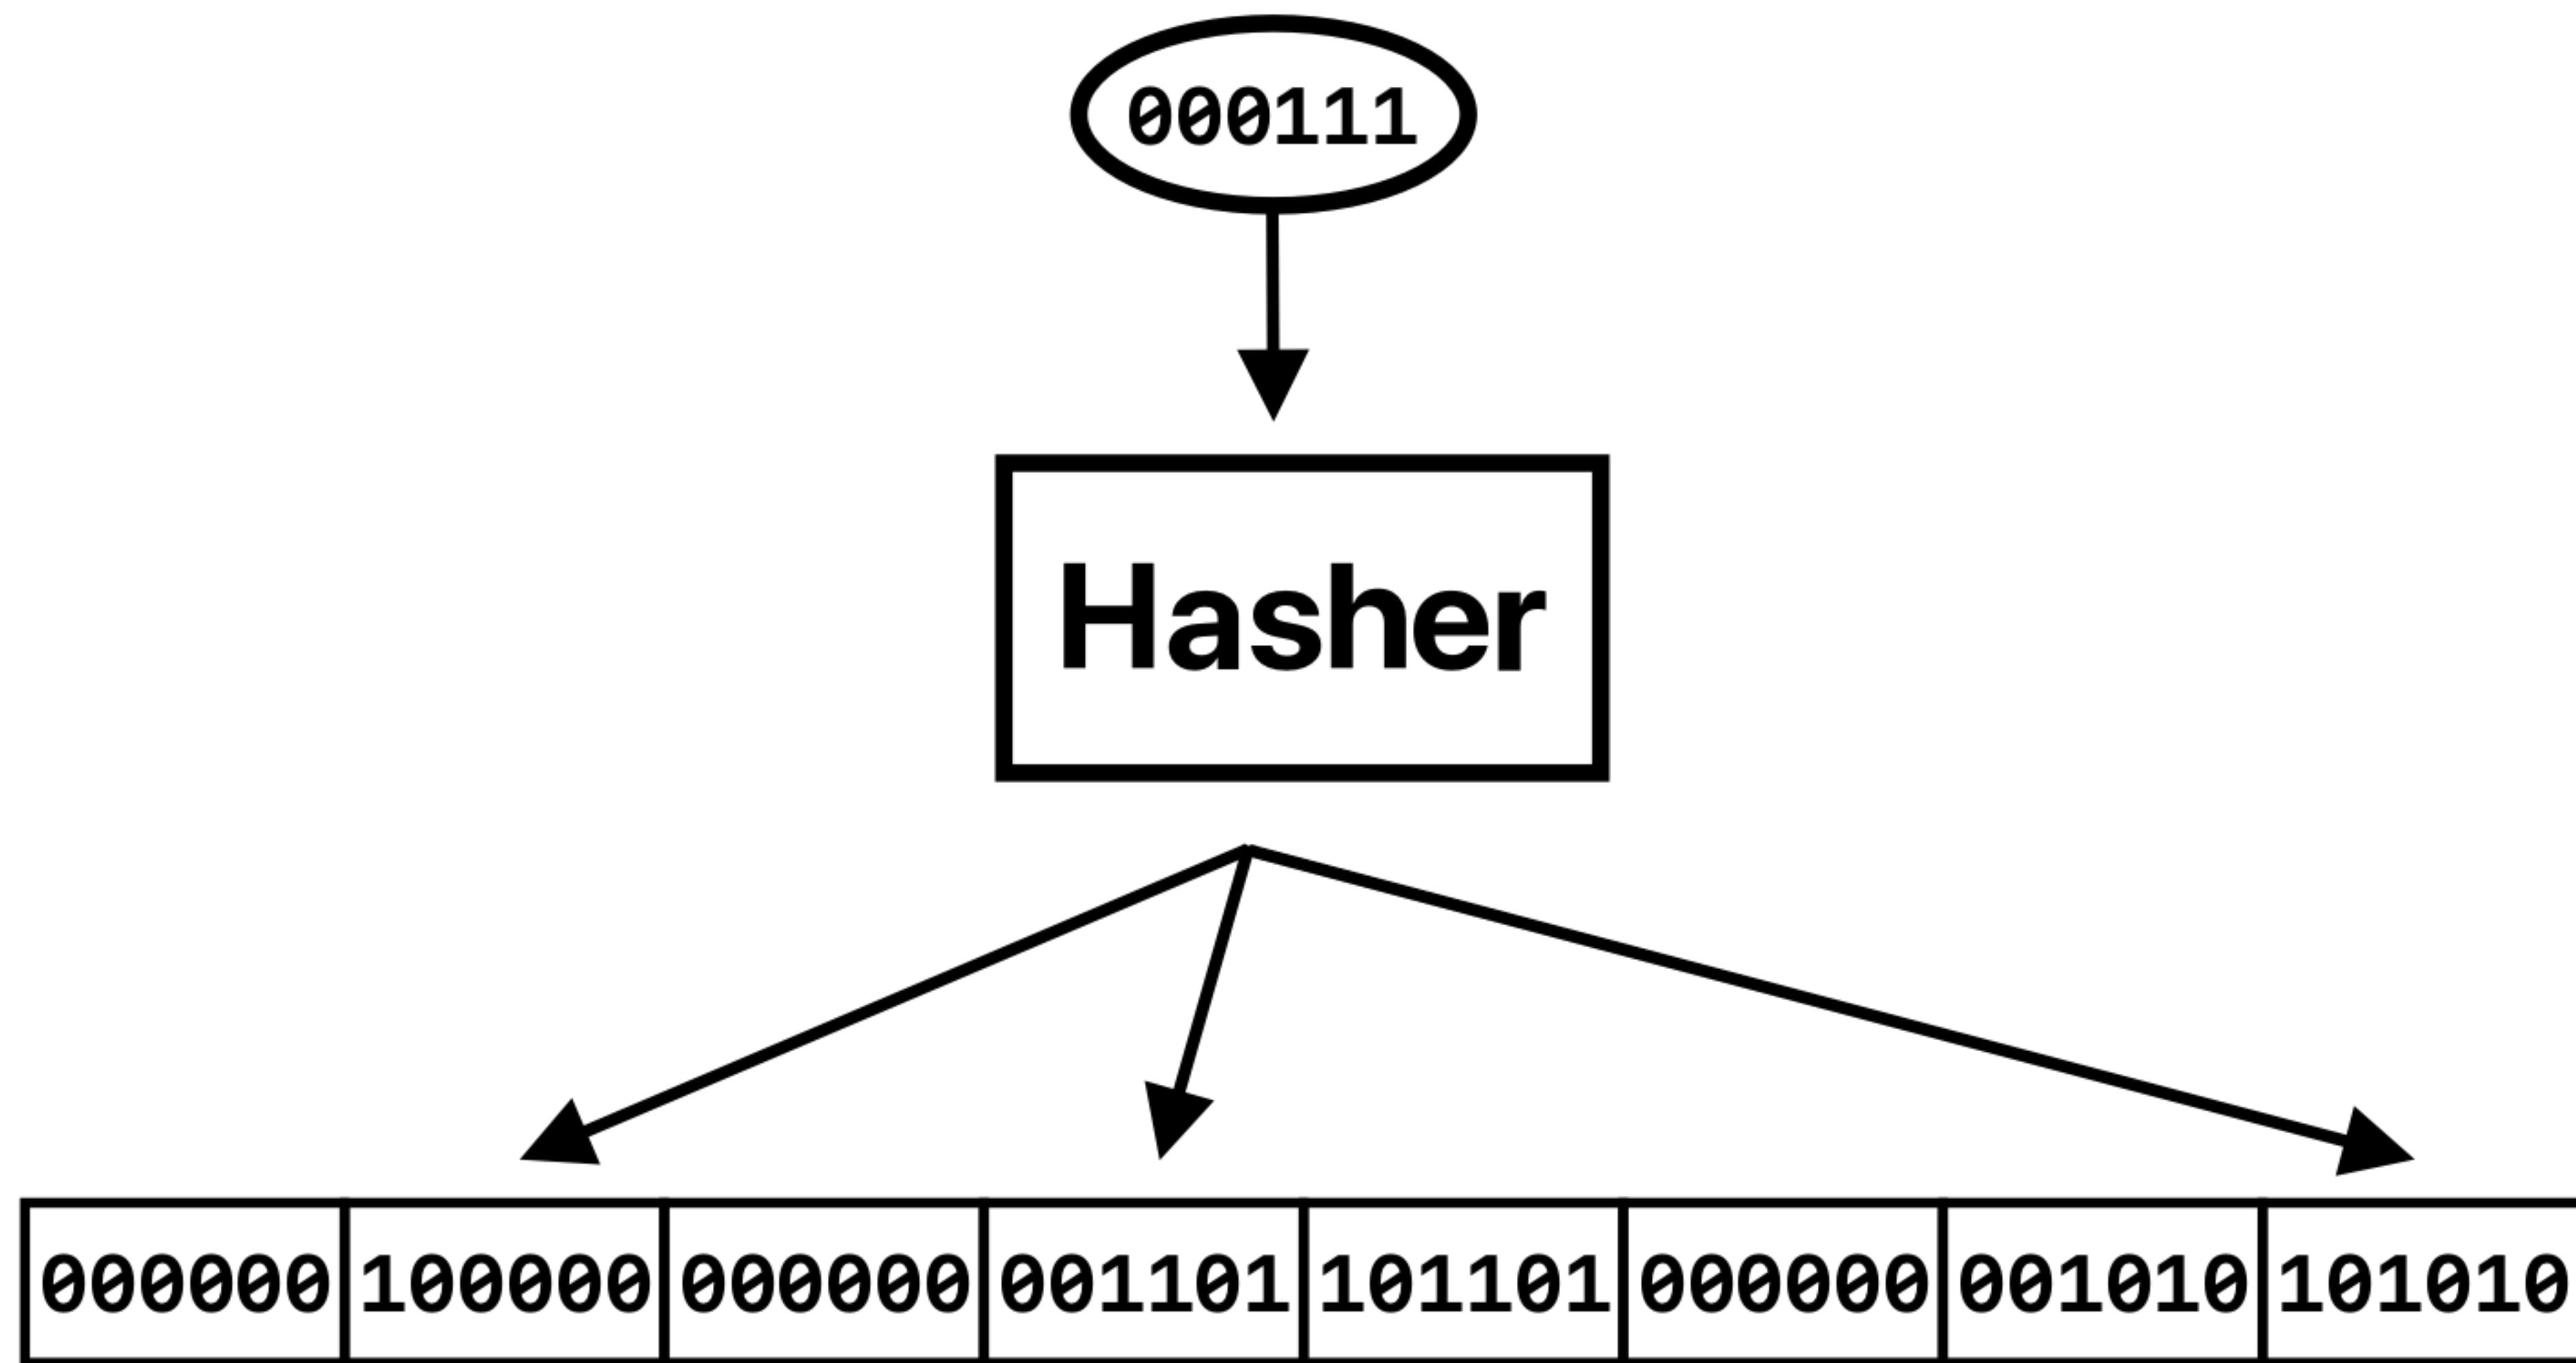
HPW IBLT - Toggle(x)



HPW IBLT - Toggle(x)



HPW IBLT - Toggle(x)



HPW IBLT - Merge(other_sketch)

101010	010000	000100	111110	101000	011100	101011	111110
--------	--------	--------	--------	--------	--------	--------	--------



011011	101000	001111	000000	100100	011011	011000	001000
--------	--------	--------	--------	--------	--------	--------	--------

HPW IBLT - Merge(other_sketch)

1	1	0	0	0	1	1	1	0	0	0	1	0	1	1	1	0	0	0	1	1	1	1	0	0	1	1	1	0	1	1	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

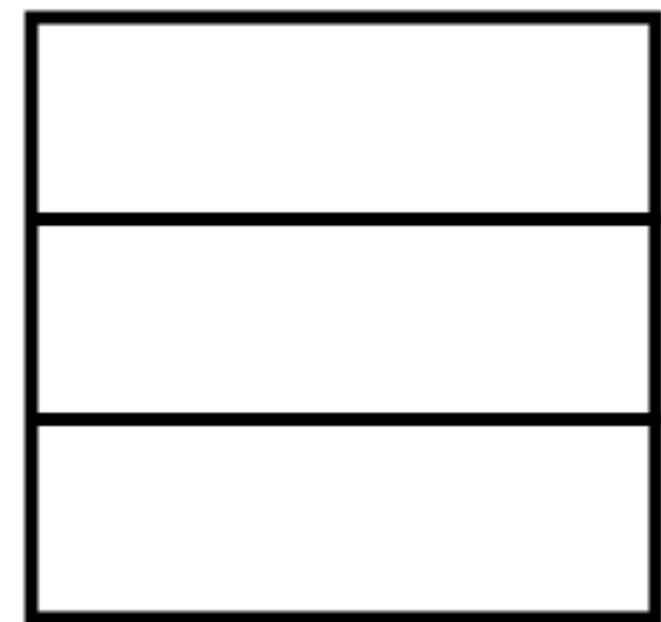
HPW IBLT - ListEntries()

while exists pure cell:

add/remove x to results

TOGGLE(x)

Result:

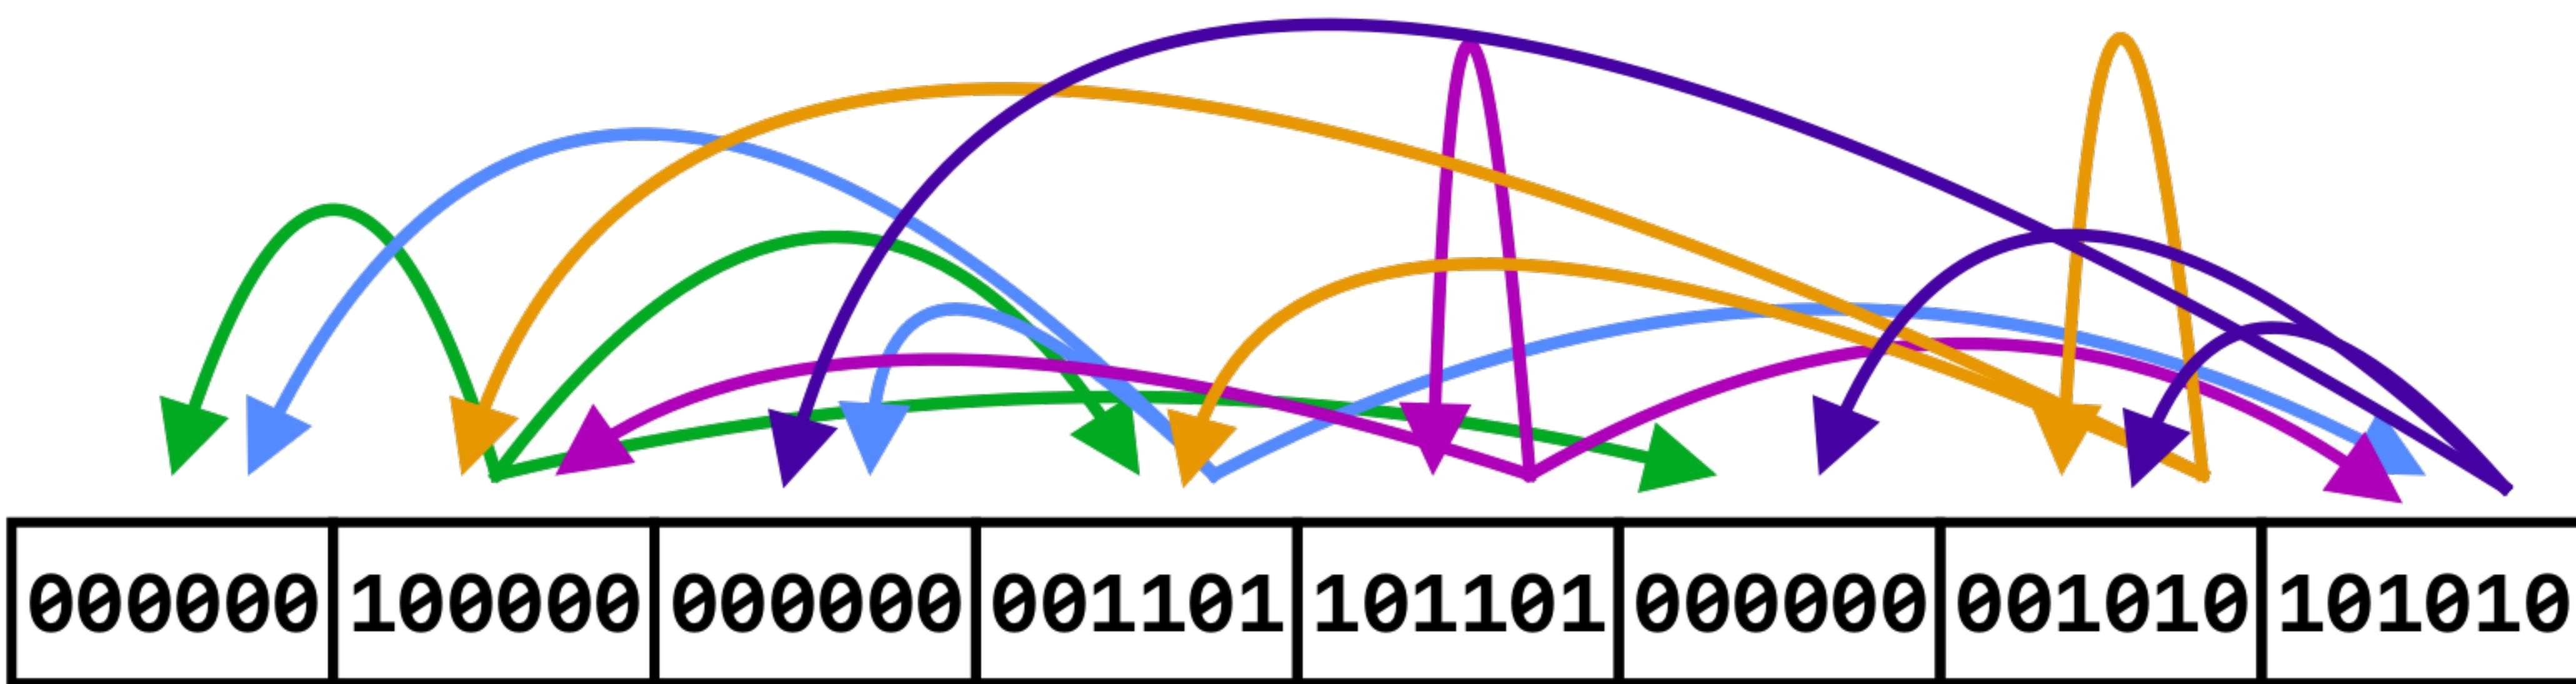


000000	100000	000000	001101	101101	000000	001010	101010
--------	--------	--------	--------	--------	--------	--------	--------

HPWIBLT - ListEntries()

Result:

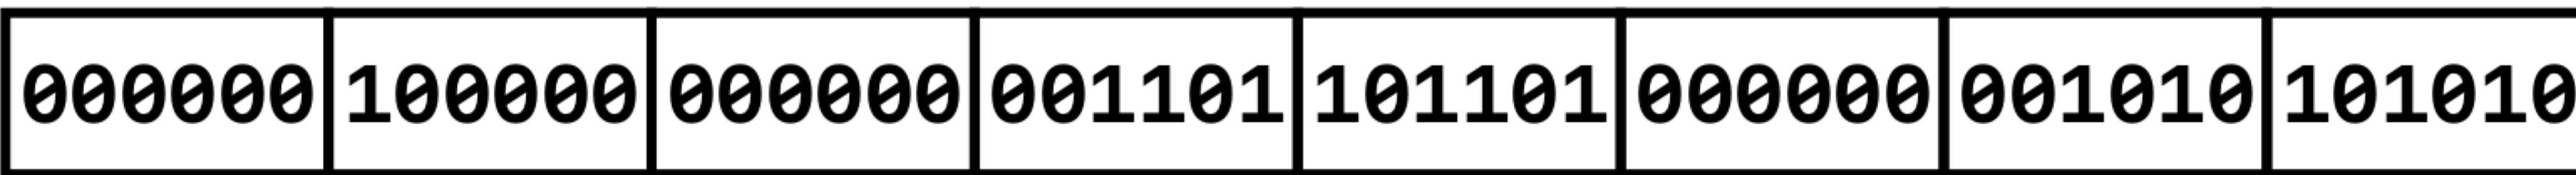
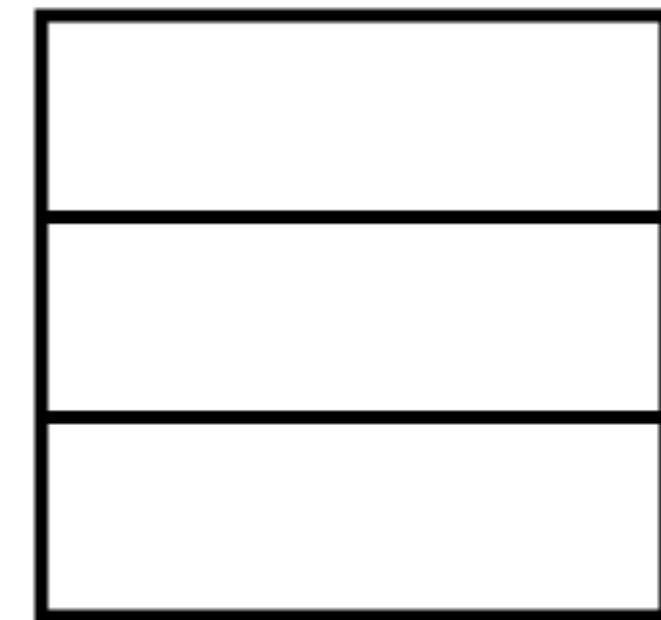
```
while exists pure cell:  
    add/remove x to results  
  
TOGGLE(x)
```



HPW IBLT - ListEntries()

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

Result:

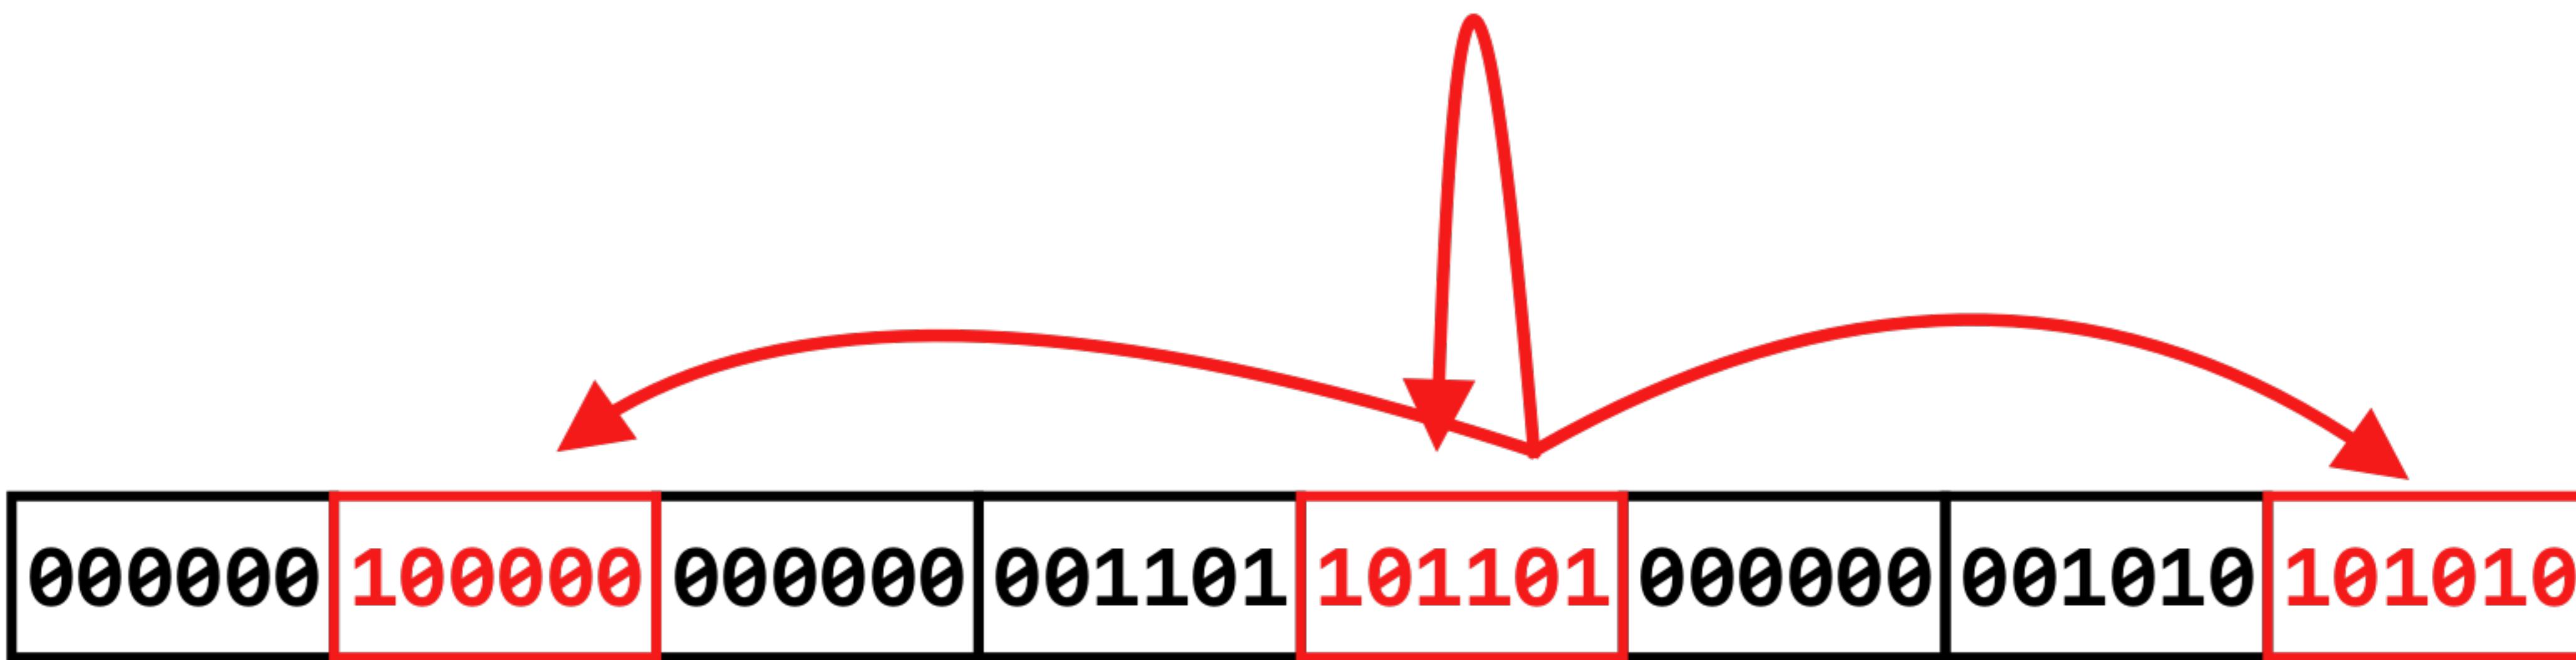


HPW IBLT - ListEntries()

Result:

101101

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

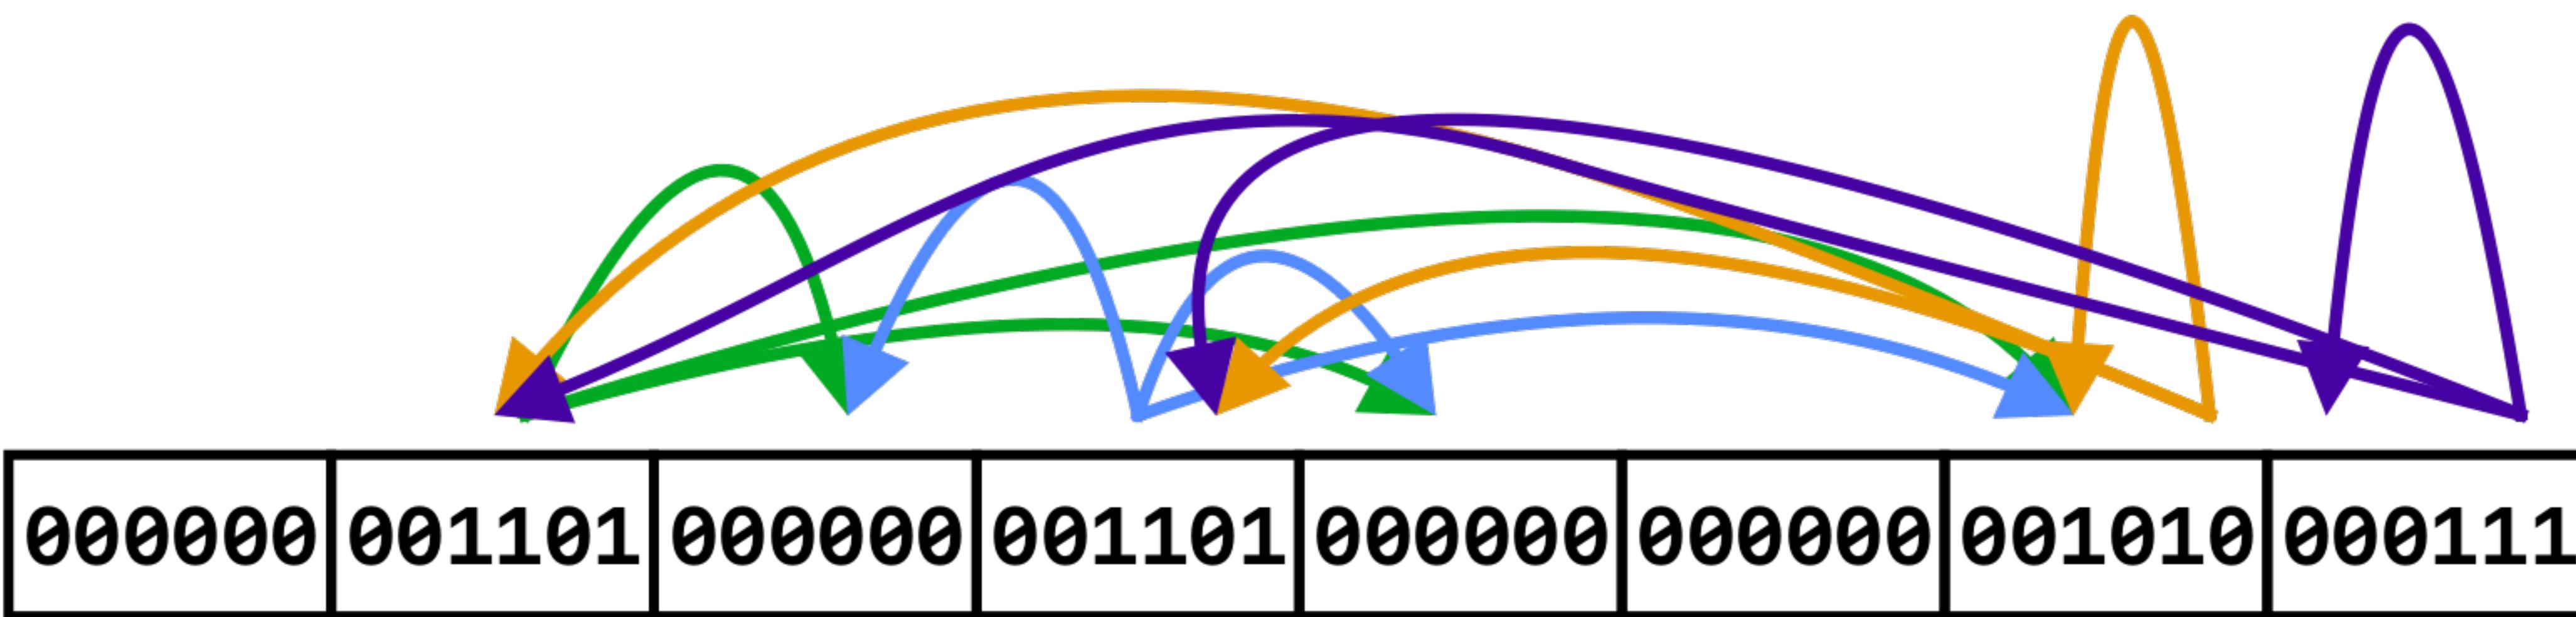


HPW IBLT - ListEntries()

Result:

101101

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

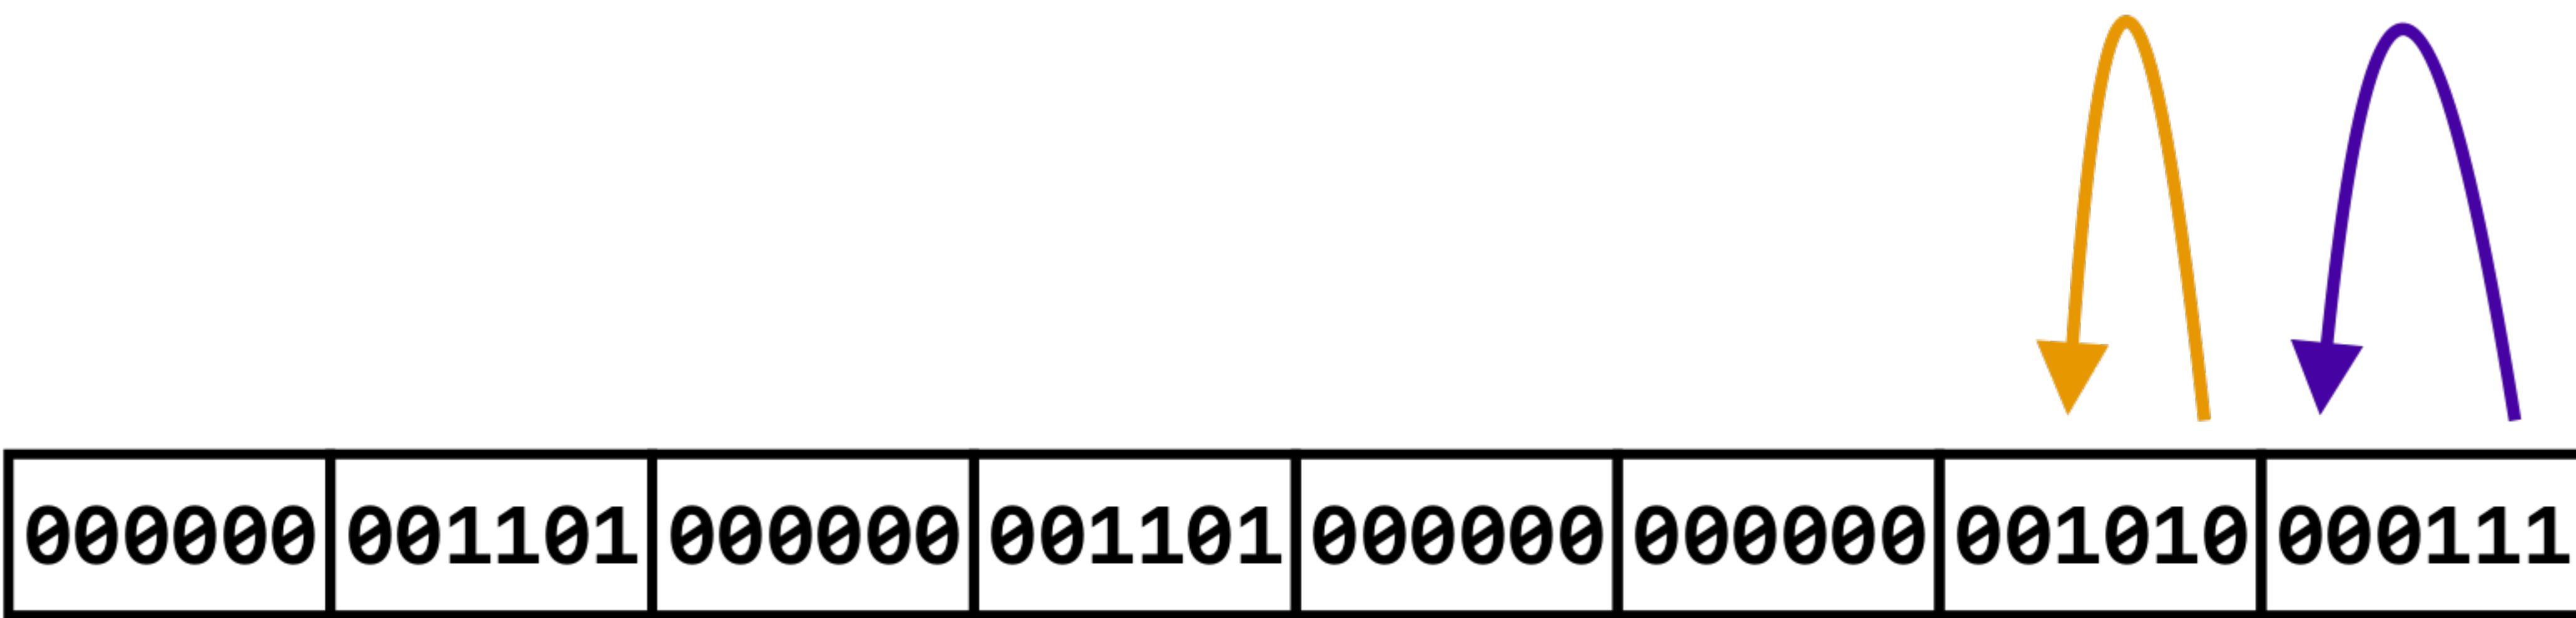


HPW IBLT - ListEntries()

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

Result:

101101

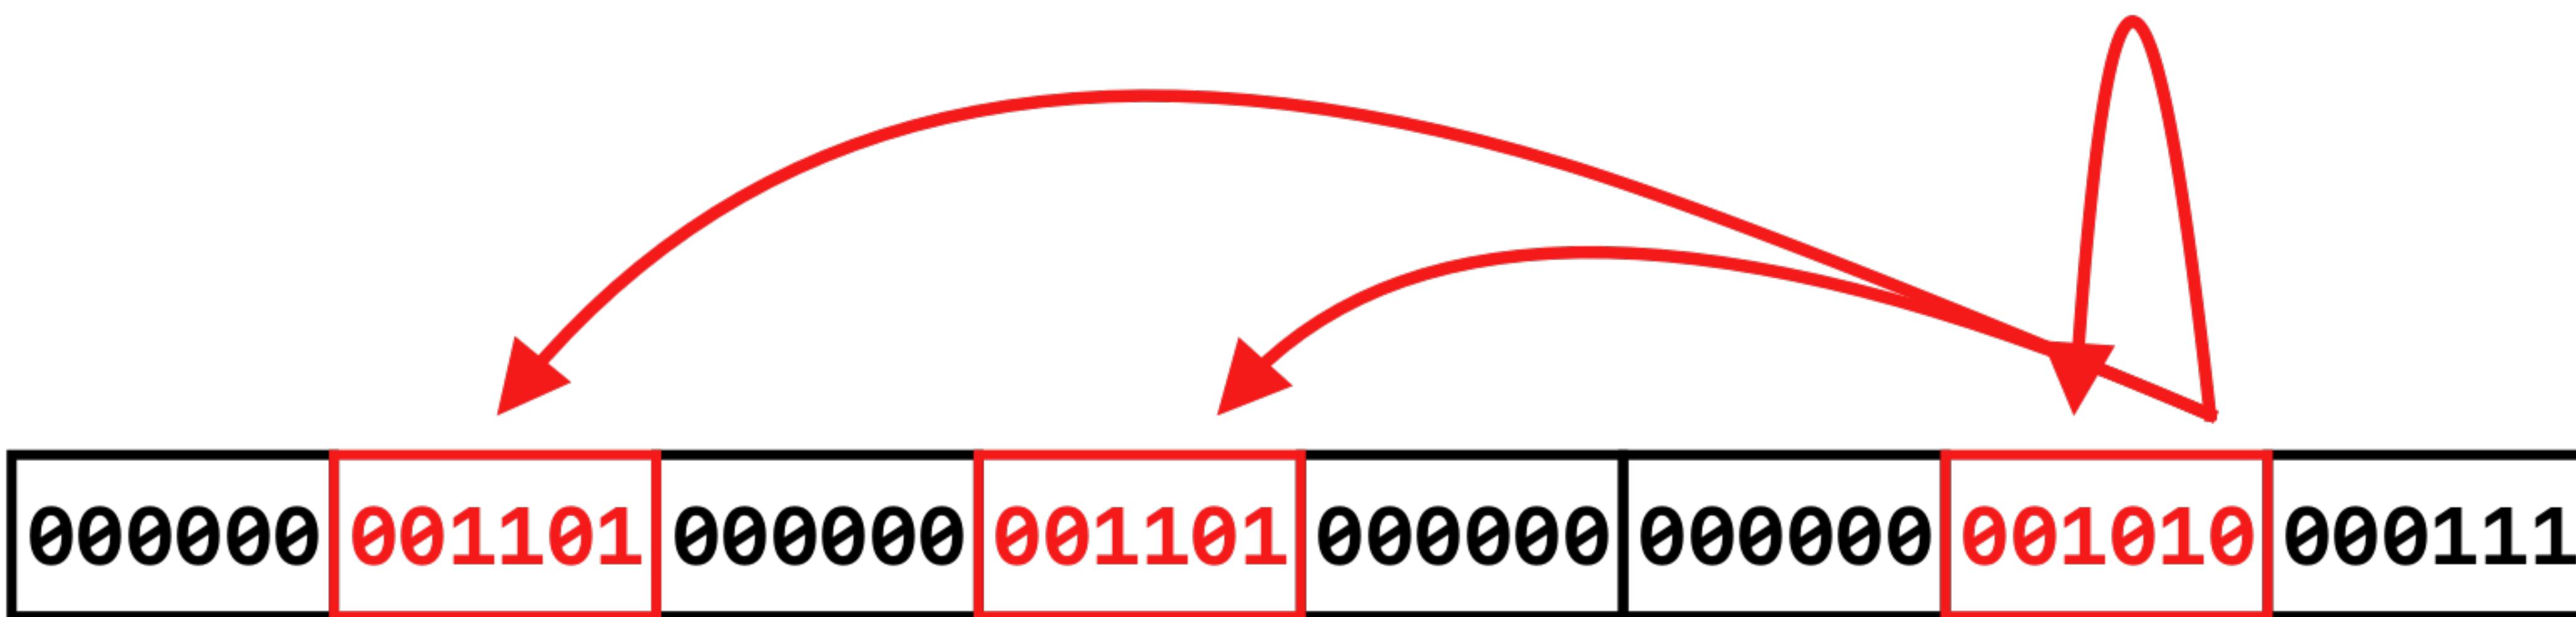


HPW IBLT - ListEntries()

Result:

101101
001010

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

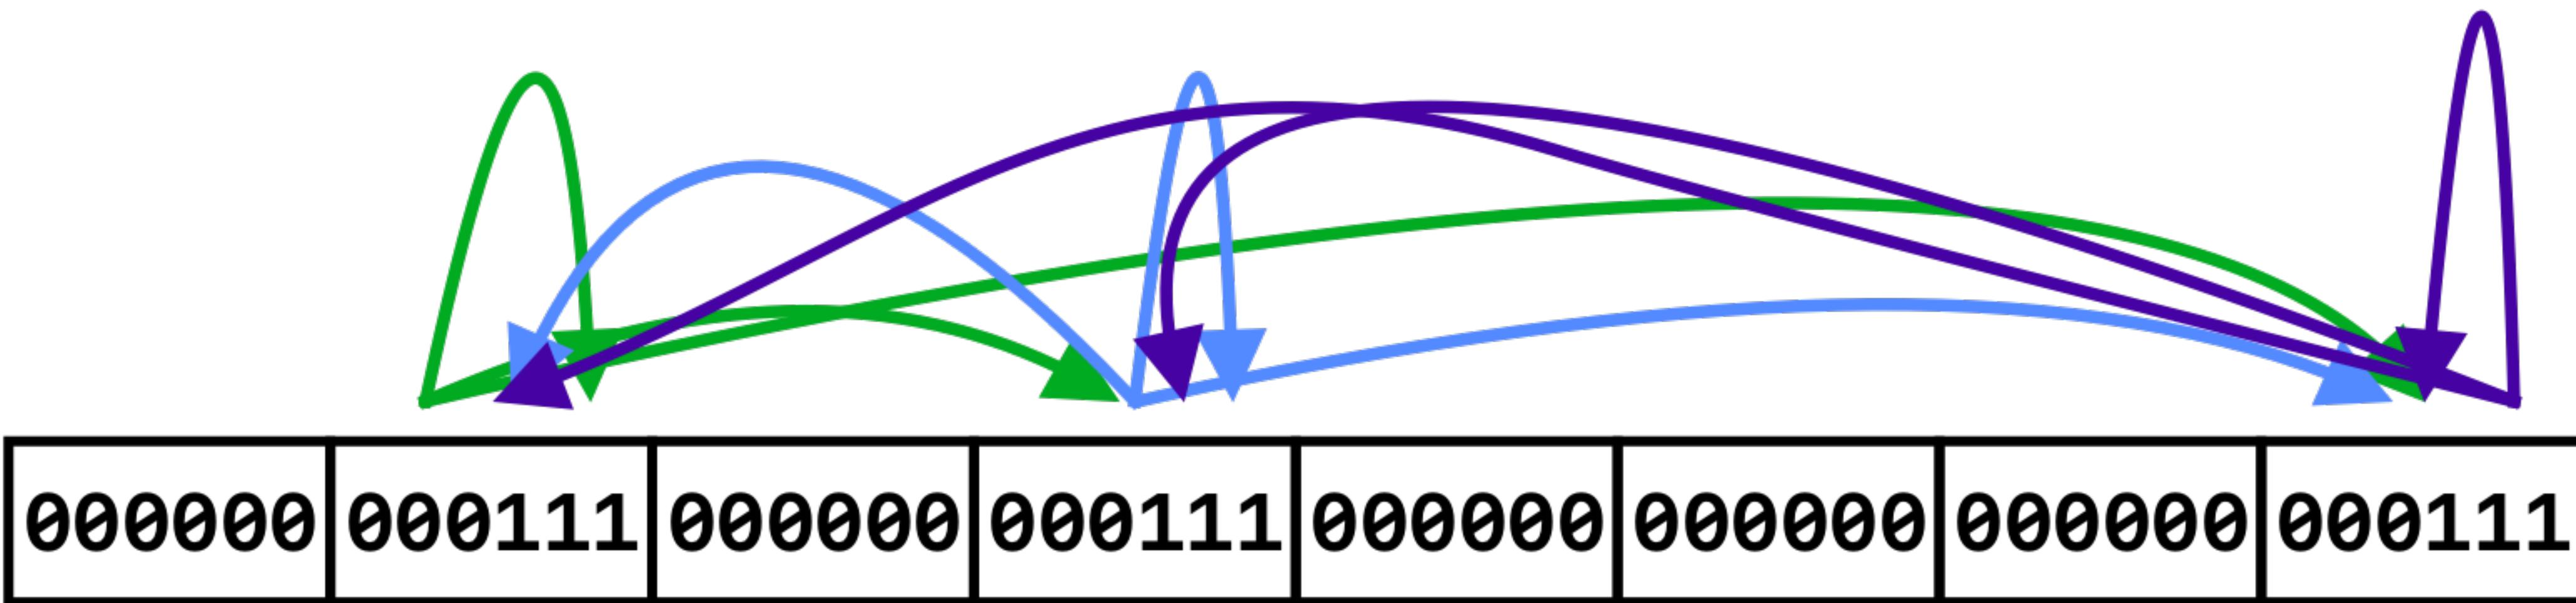


HPW IBLT - ListEntries()

Result:

101101
001010

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

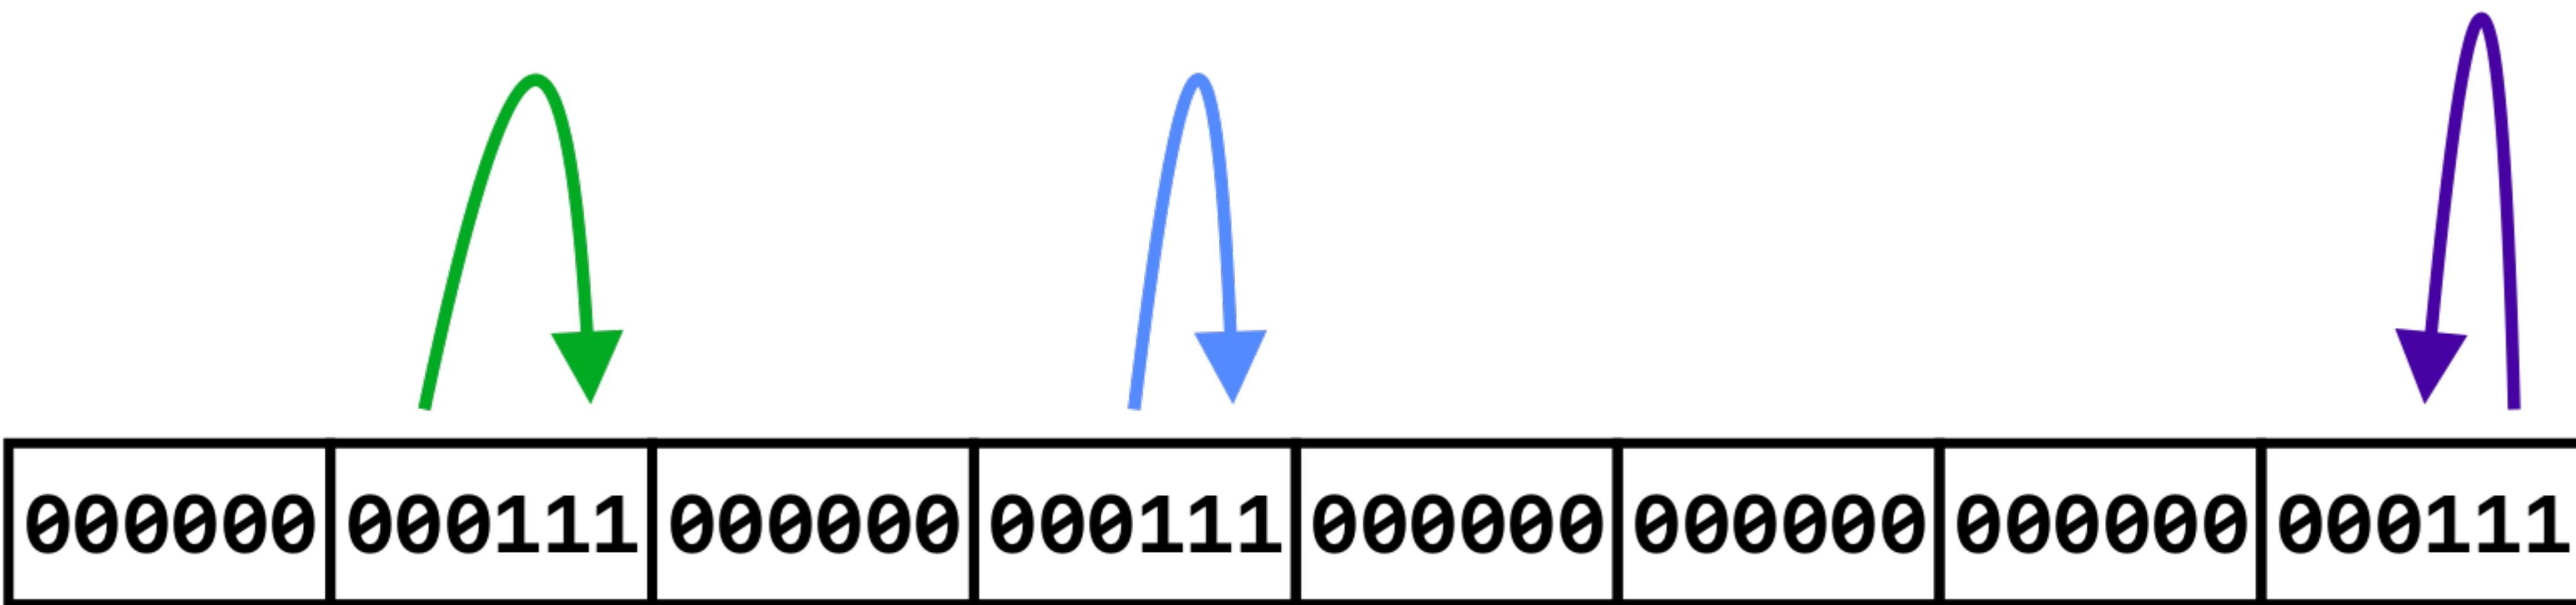


HPW IBLT - ListEntries()

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

Result:

101101
001010

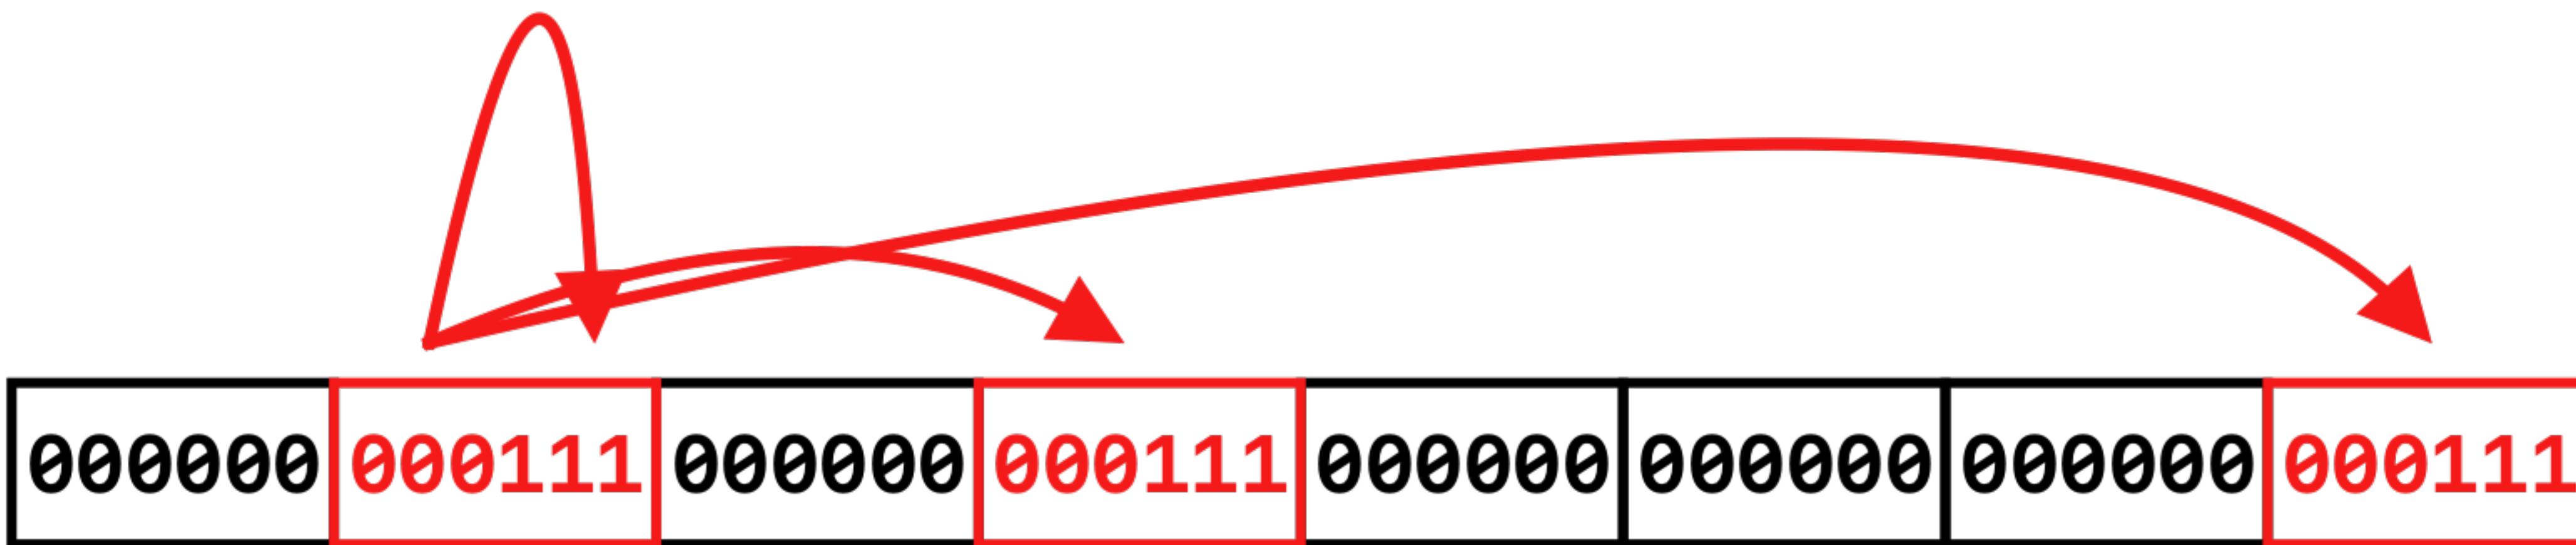


HPW IBLT - ListEntries()

Result:

101101
001010
000111

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

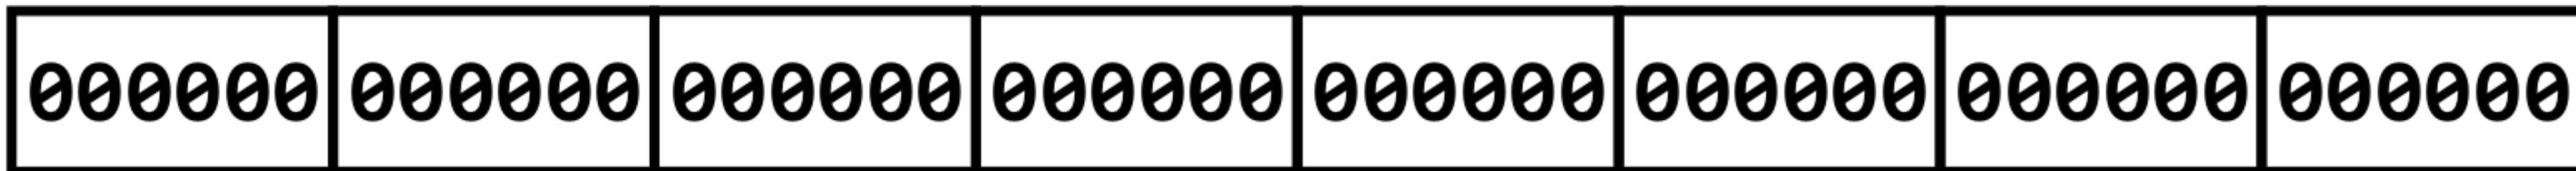


HPW IBLT - ListEntries()

```
while exists pure cell:  
    add/remove x to results  
    TOGGLE(x)
```

Result:

101101
001010
000111



Success!

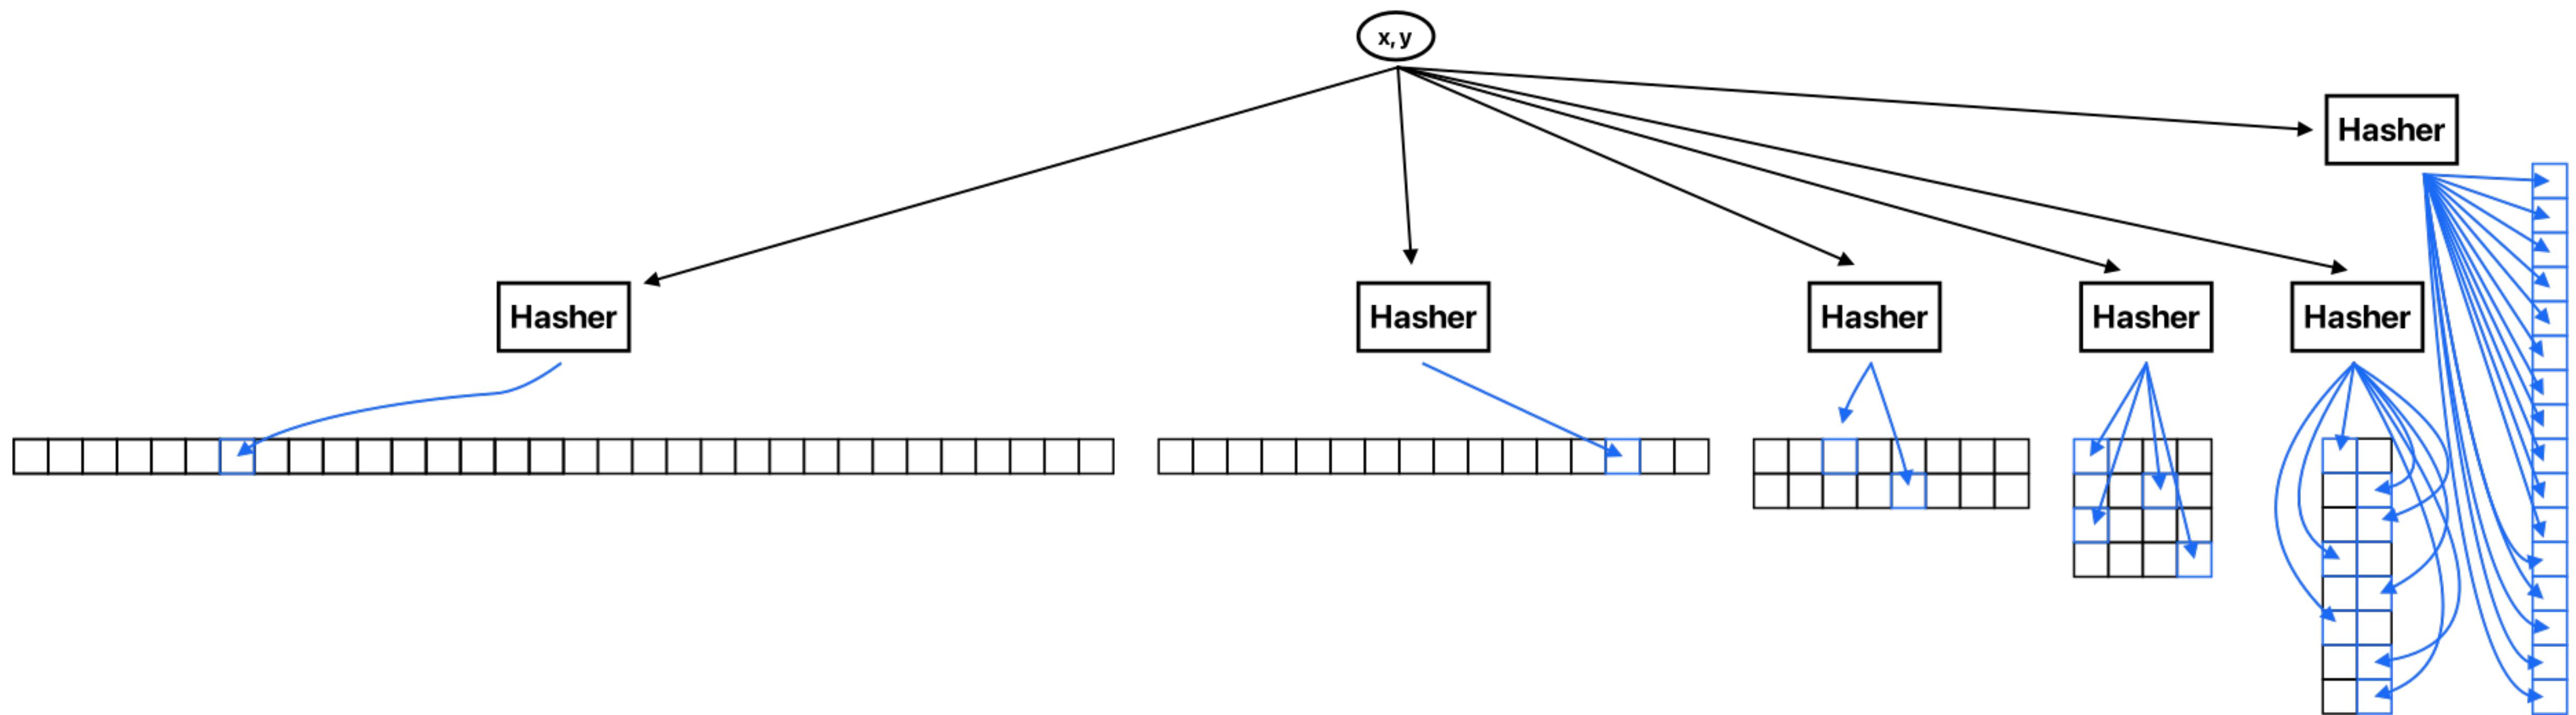
HPW IBLT Quirks

- Only stores keys
- Anomalies are mostly self-correcting
- No partial failure, all or nothing.

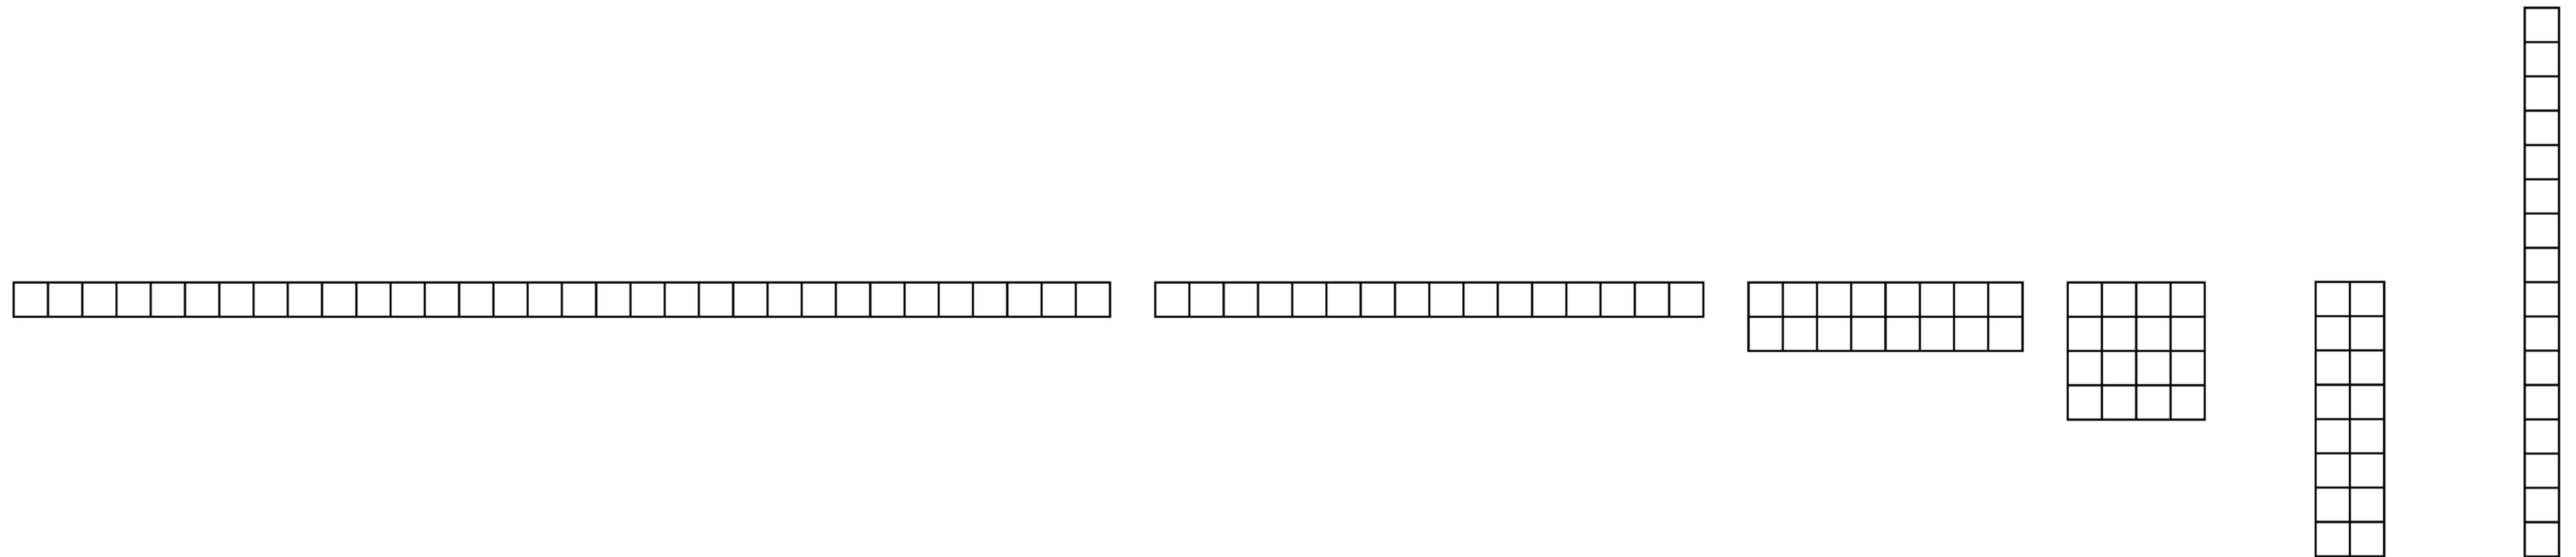
Stacked IBLT

- Key/Value Store
- Uses multiple "Basic IBLTs"
- Operations:
 - Insert(key, value)
 - Remove(key, value)
 - ListEntries()

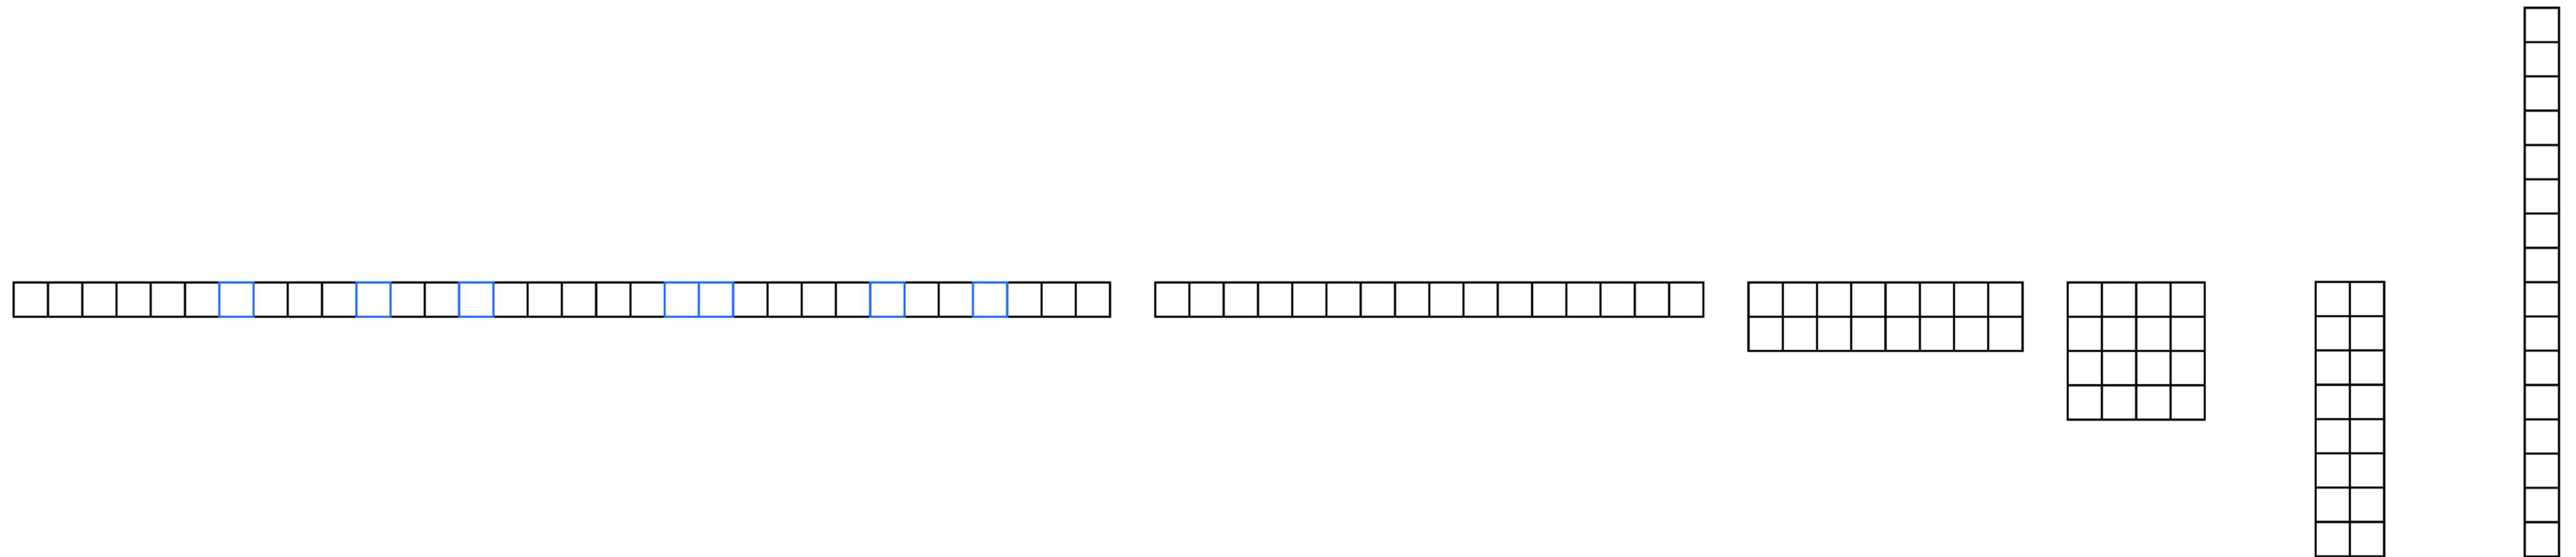
Stacked IBLT



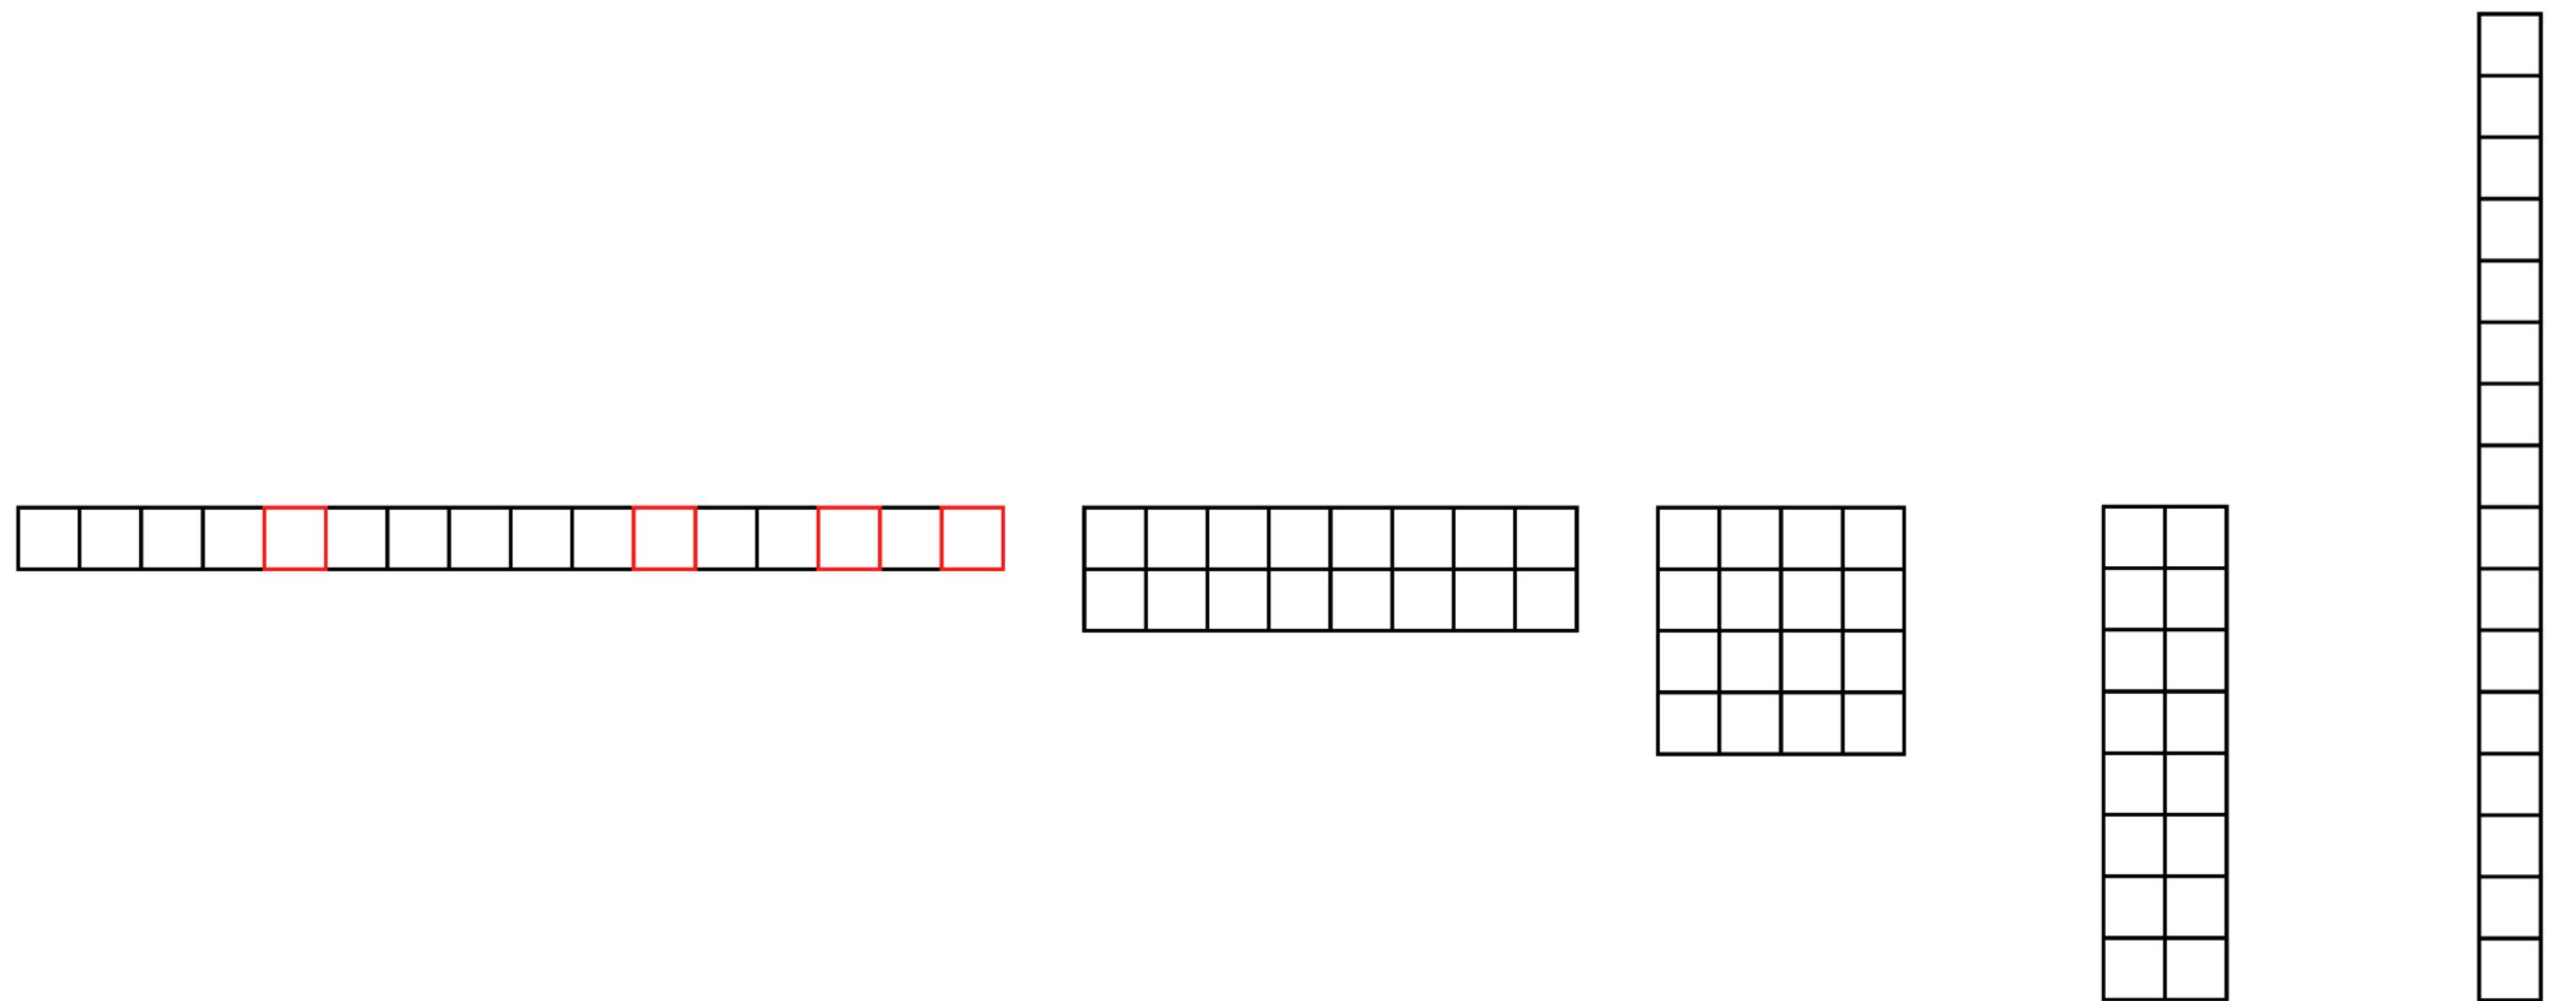
Stacked IBLT - ListEntries()



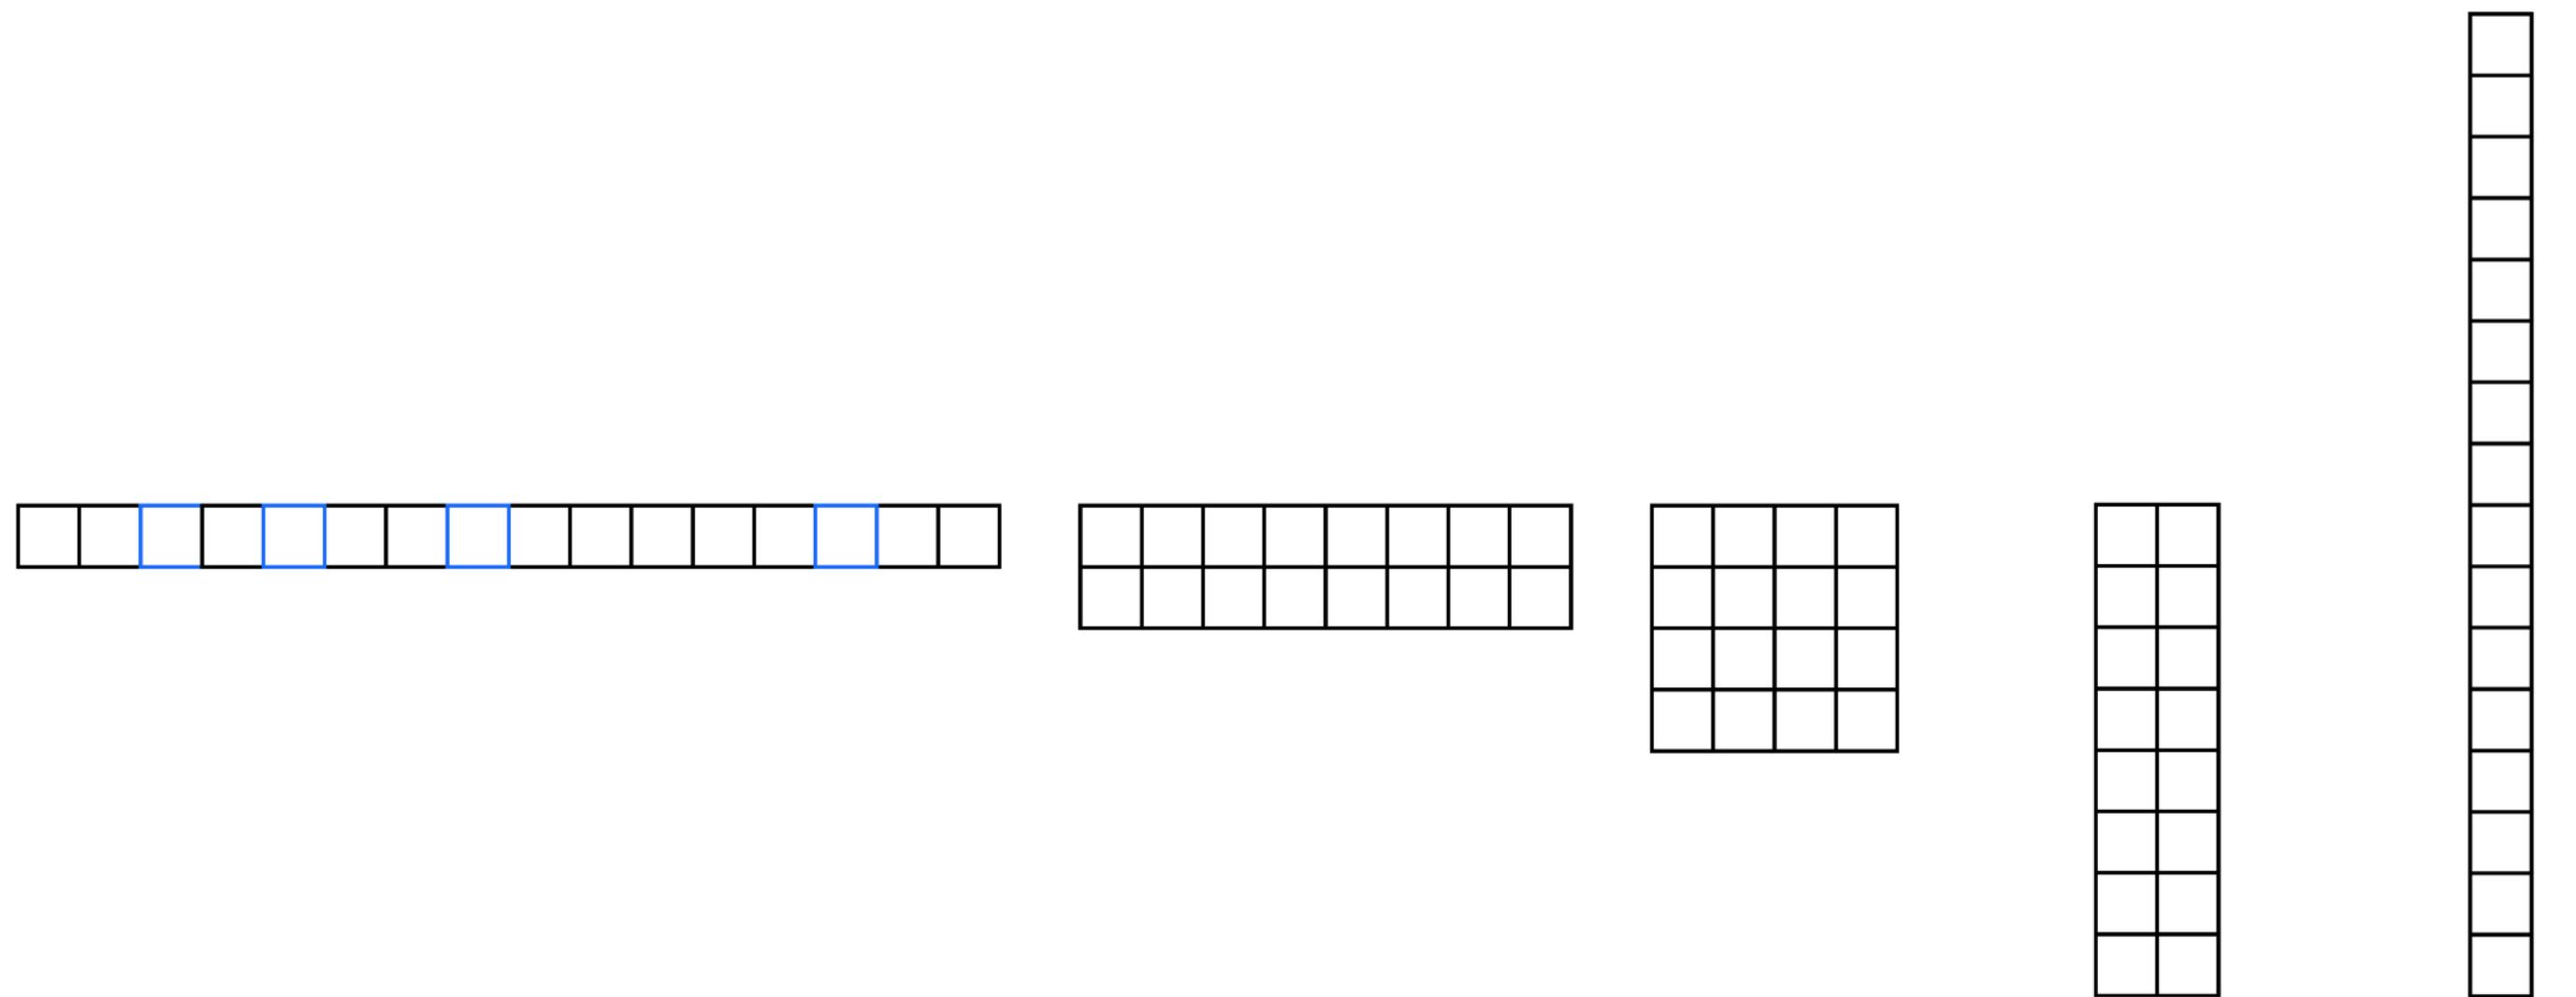
Stacked IBLT - ListEntries()



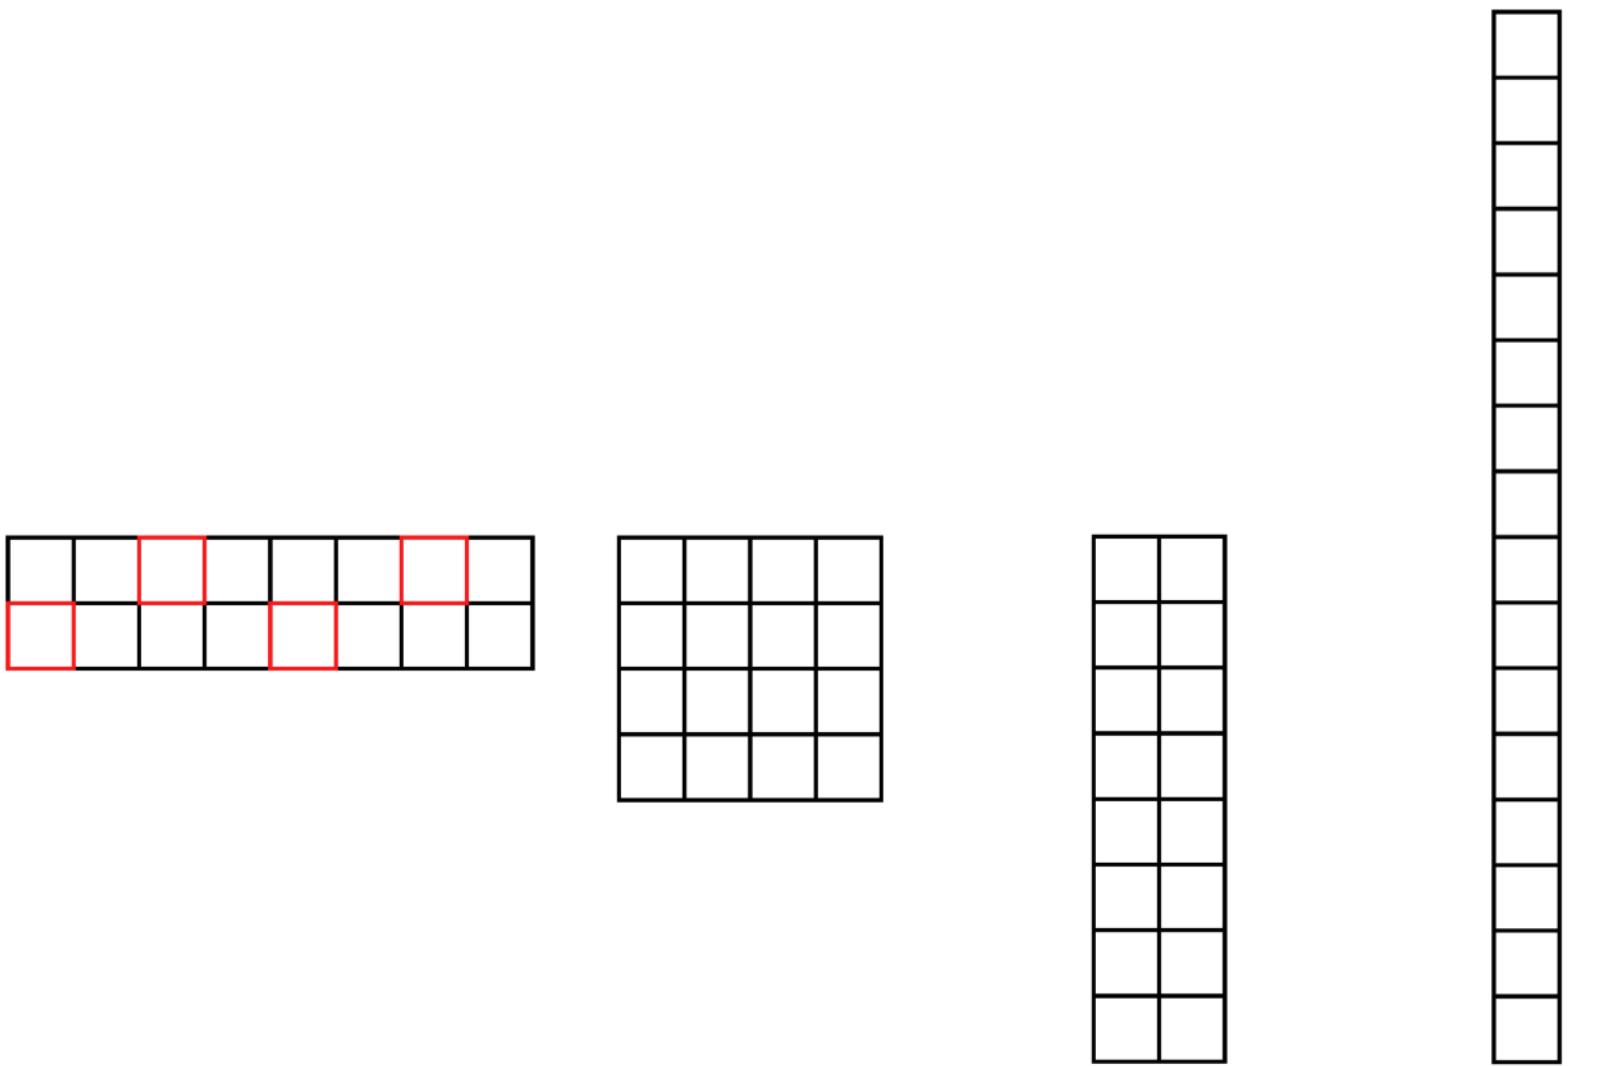
Stacked IBLT - ListEntries()



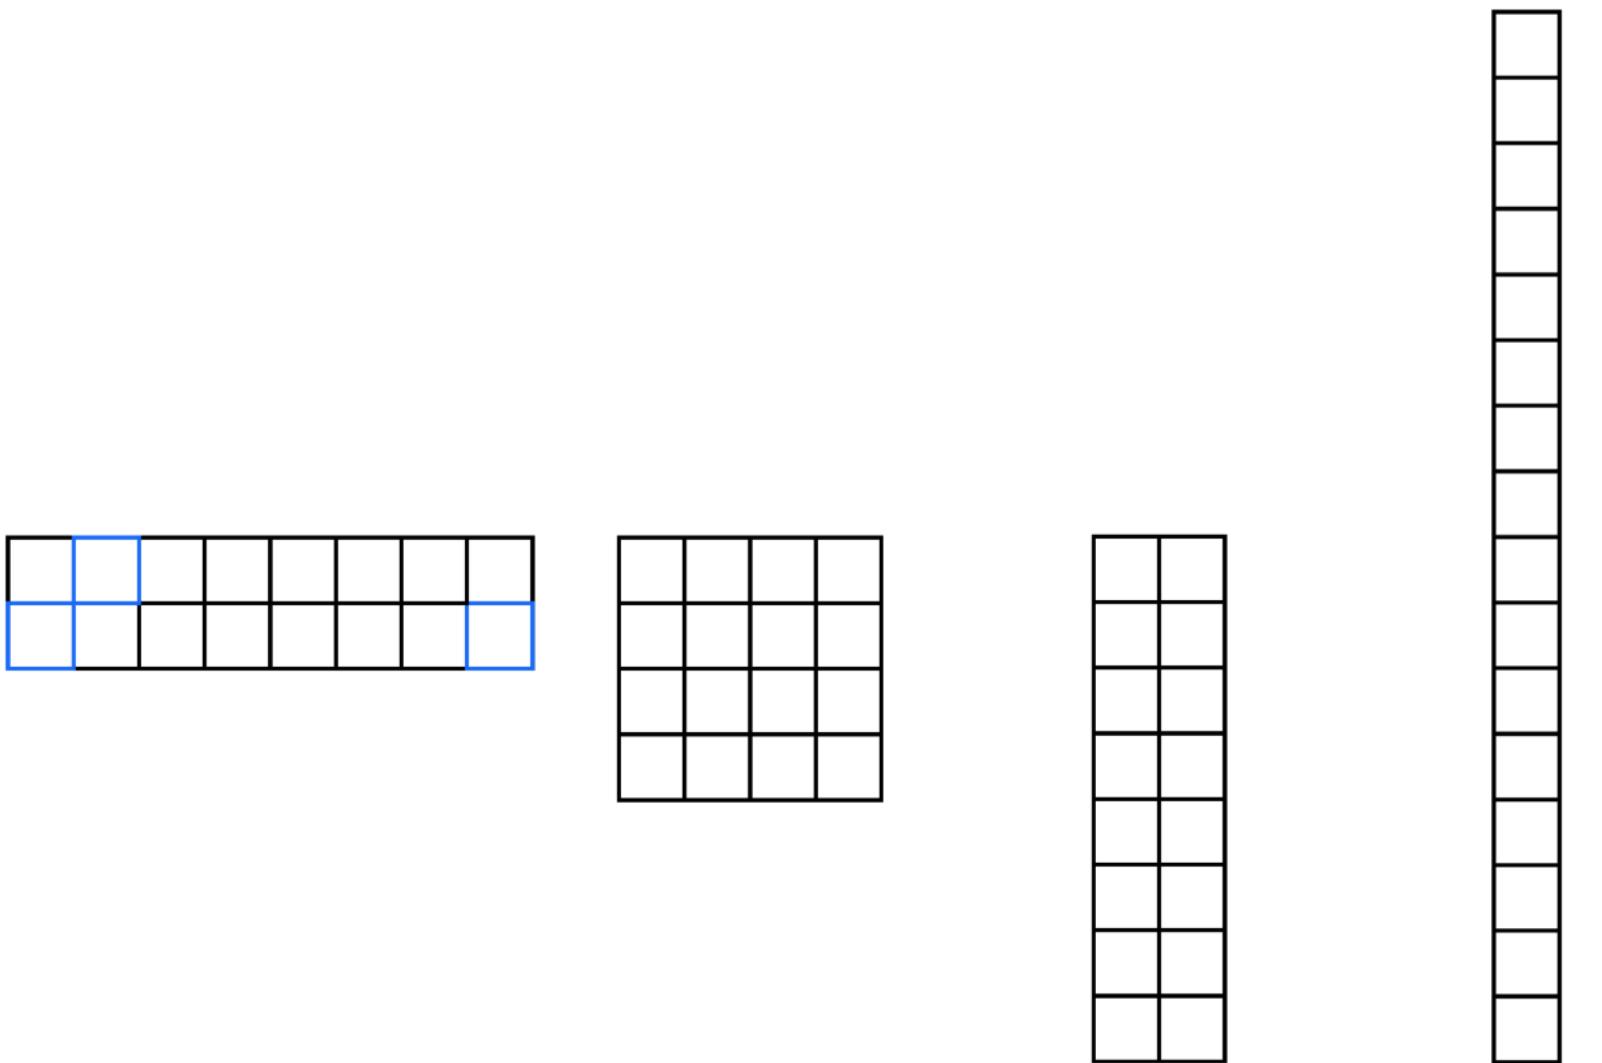
Stacked IBLT - ListEntries()



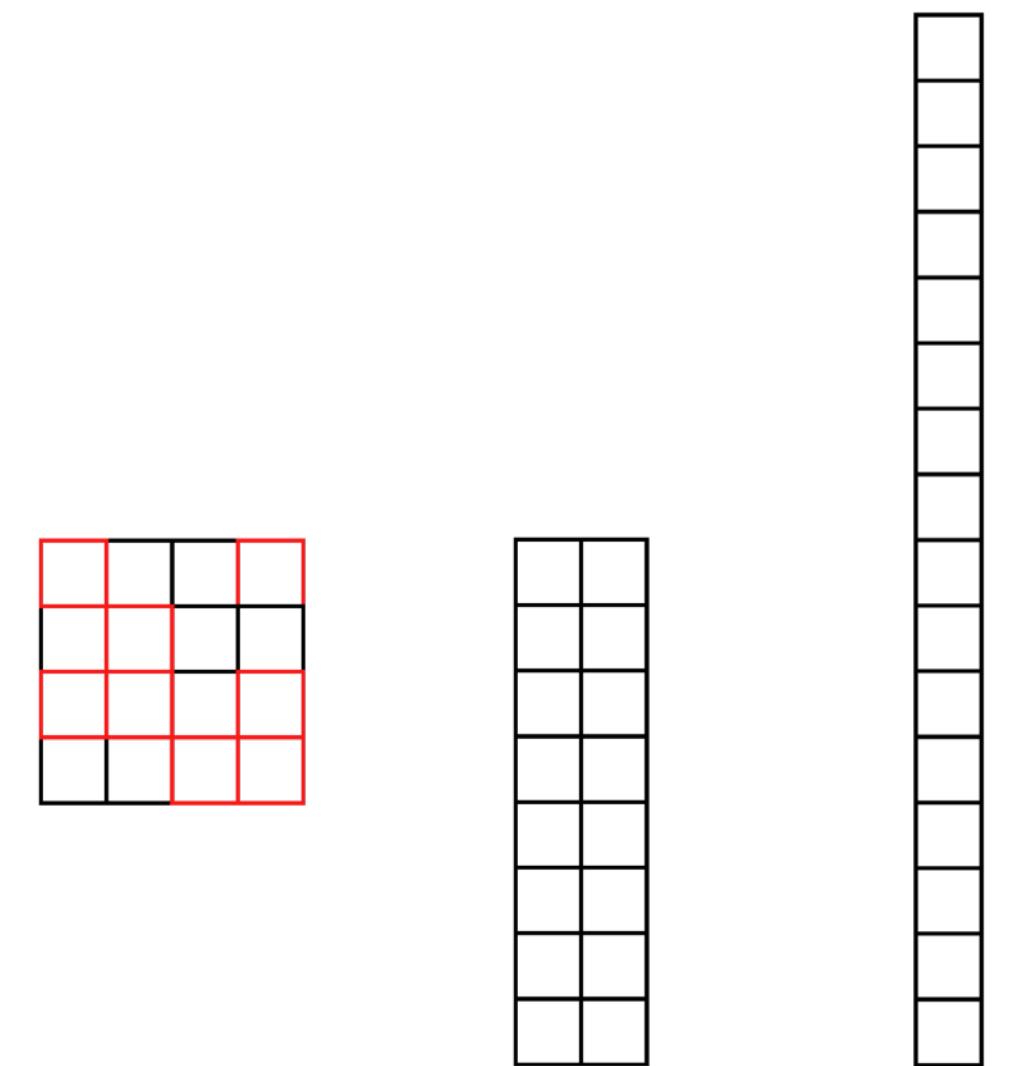
Stacked IBLT - ListEntries()



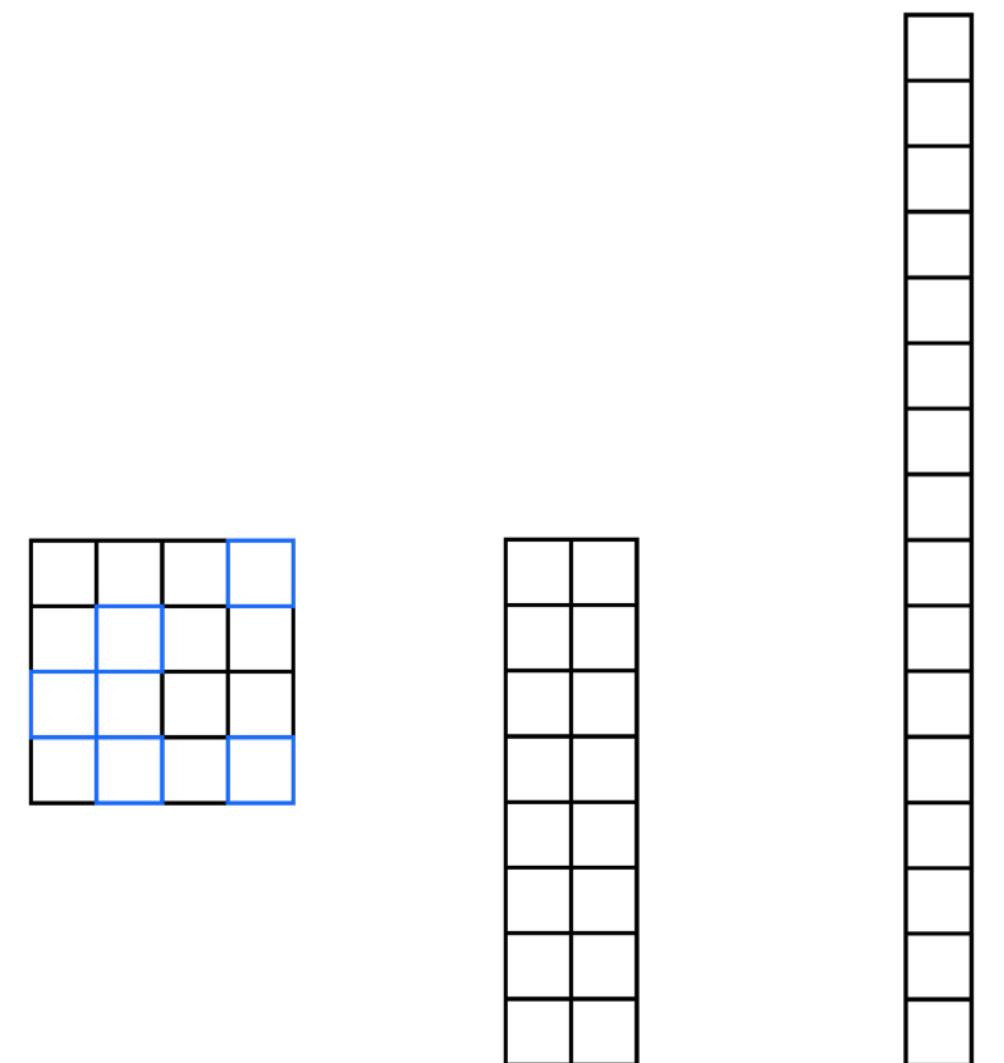
Stacked IBLT - ListEntries()



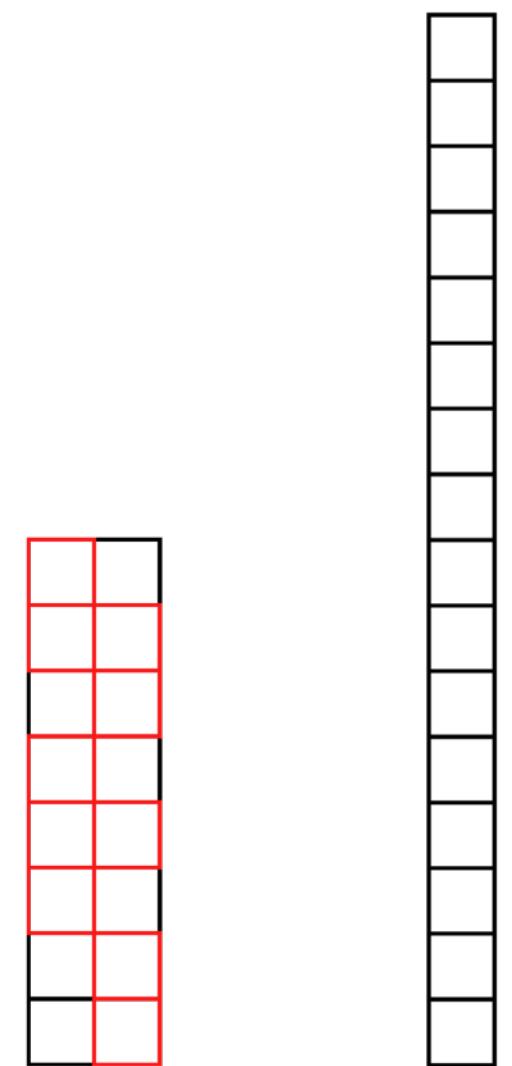
Stacked IBLT - ListEntries()



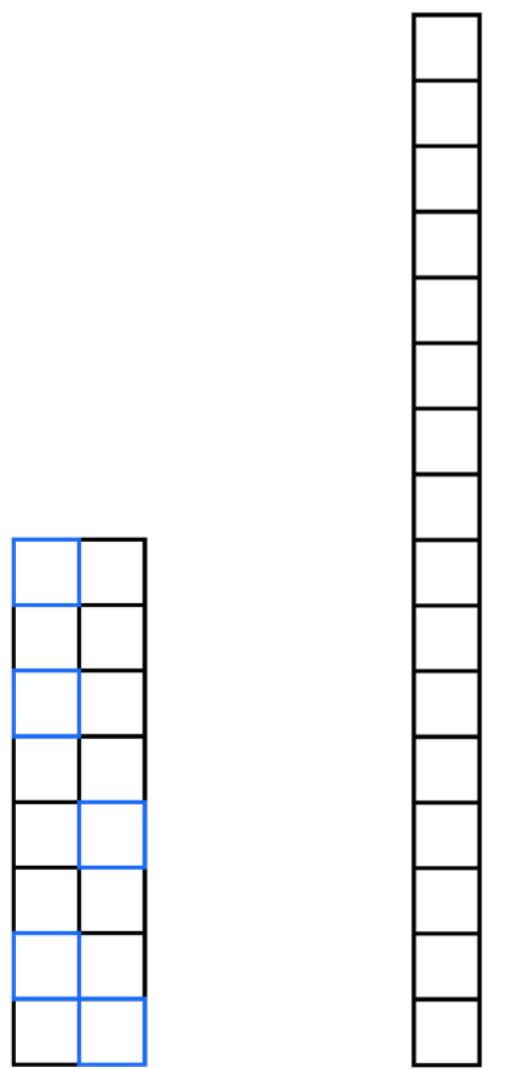
Stacked IBLT - ListEntries()



Stacked IBLT - ListEntries()



Stacked IBLT - ListEntries()



Stacked IBLT - ListEntries()



Stacked IBLT - ListEntries()

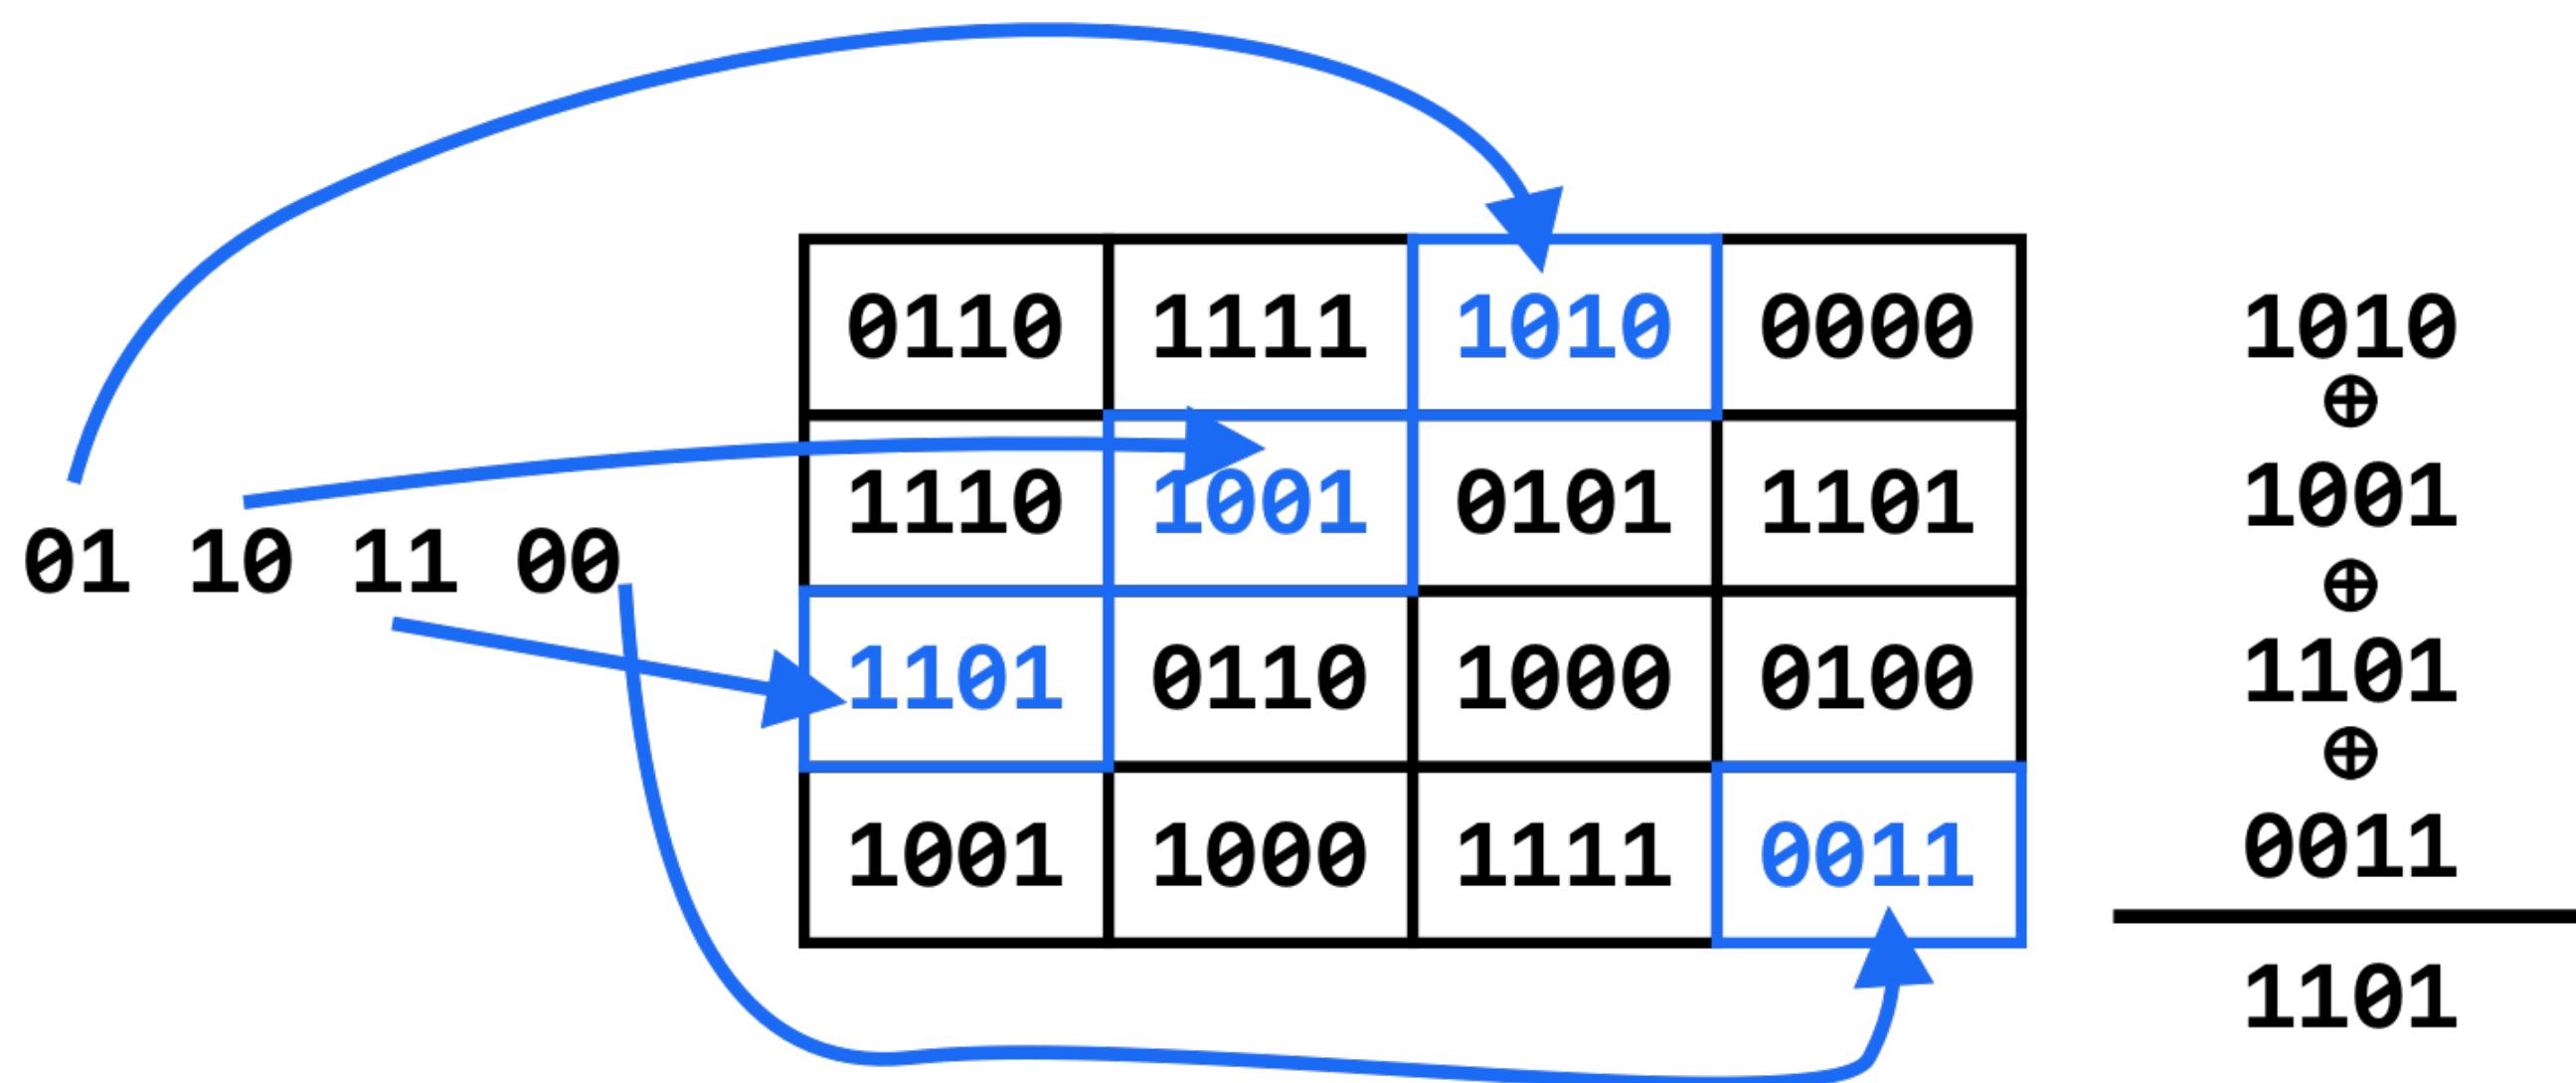


Stacked IBLT Quirks

- No individual key->value access
- Supports partial decodes, like Original
- Comes with more configuration options

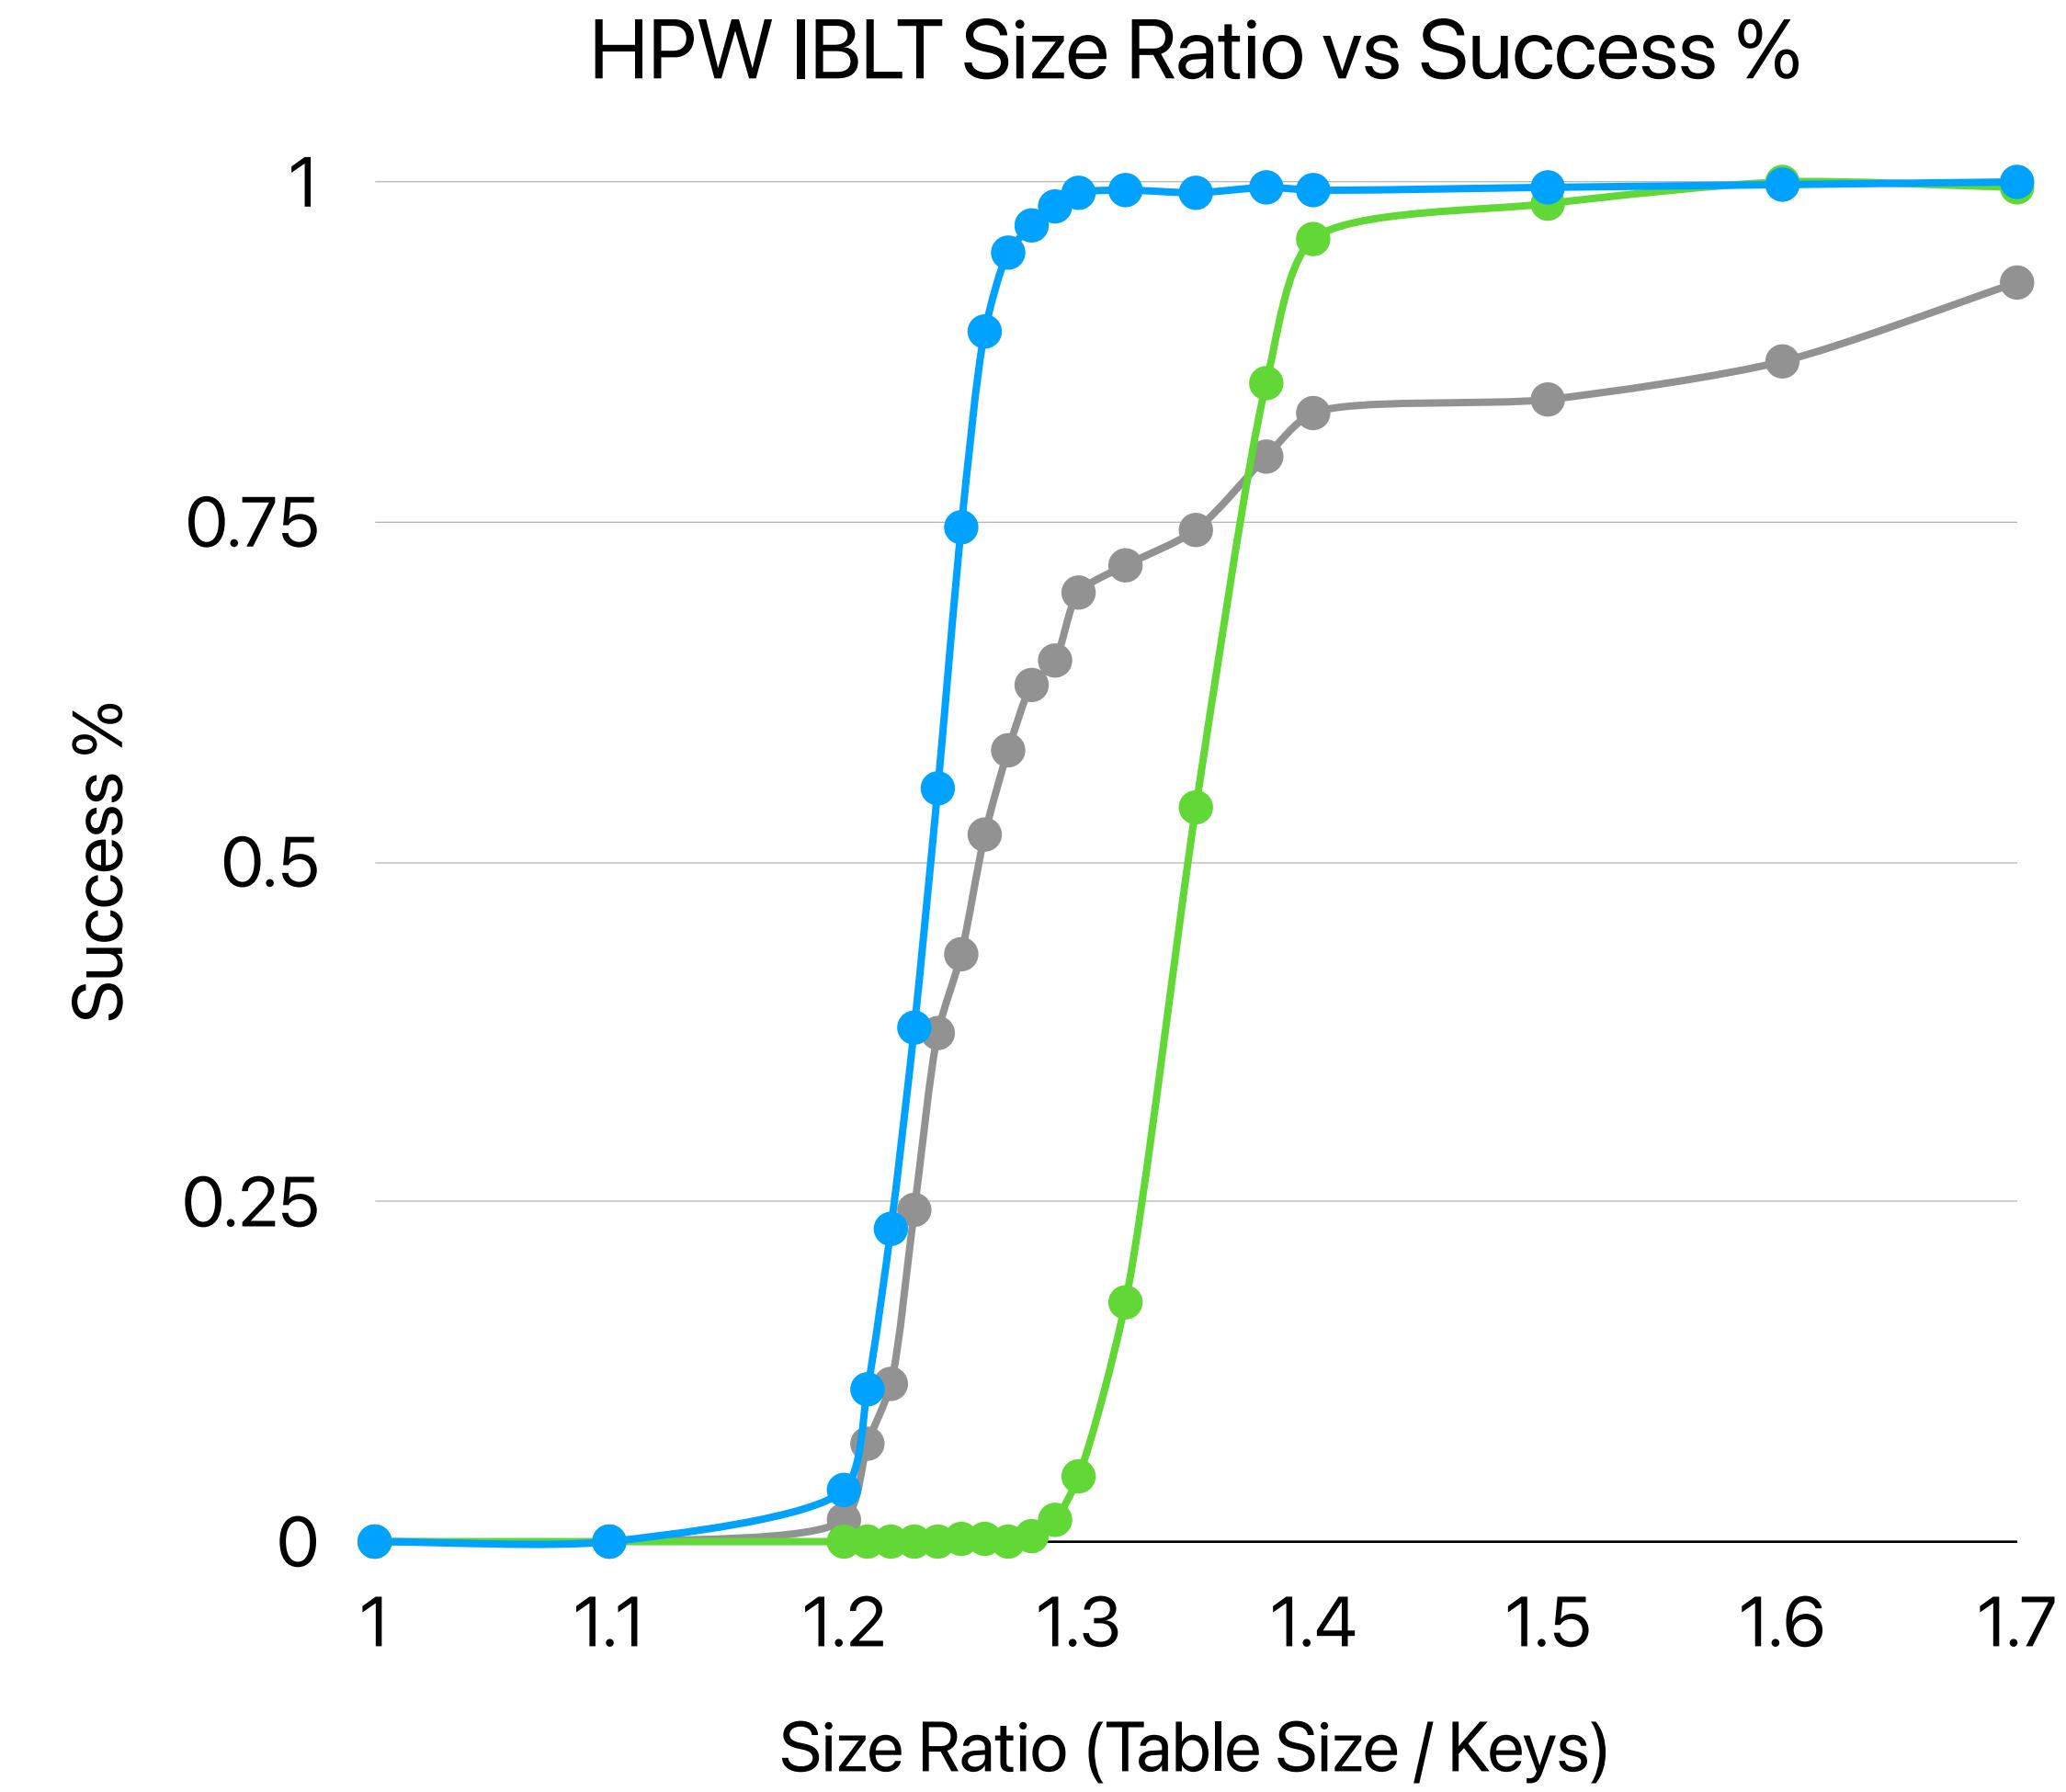
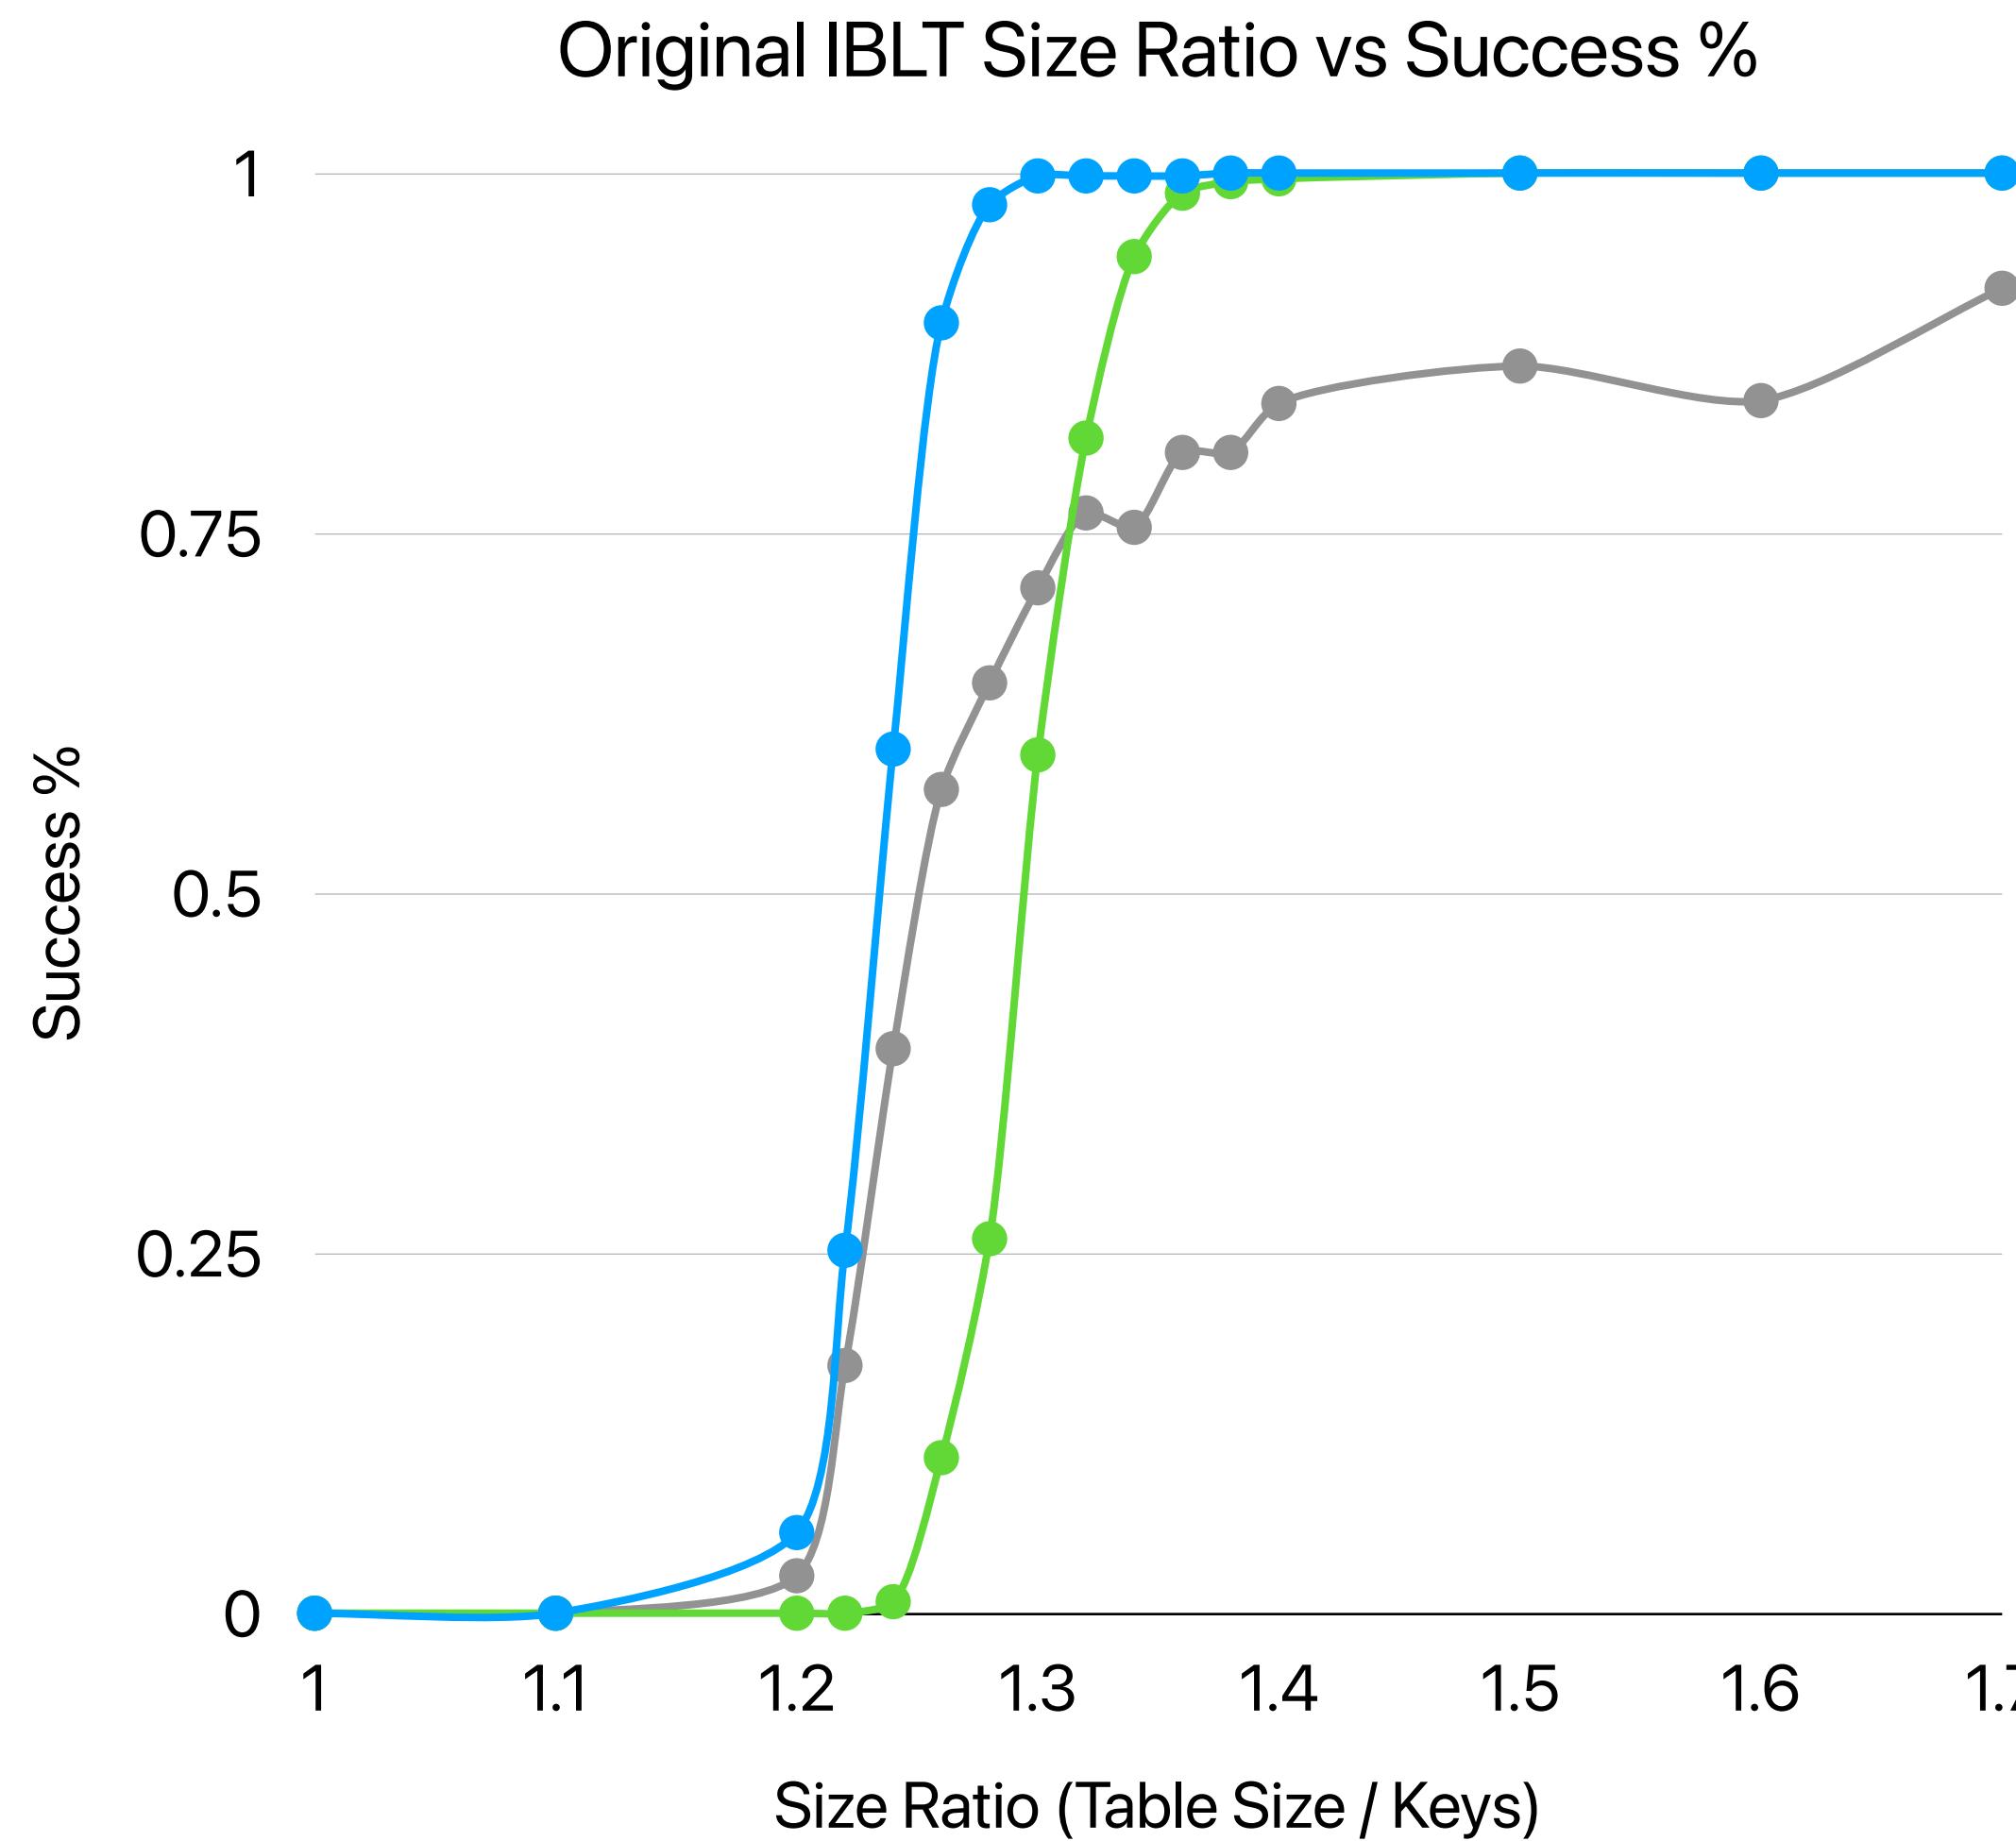
Hash Functions

- Linear Hash:
 - $ax+b \% p \% \text{array_size}$
- Knuth's Multiplicative Hash:
 - $ax \gg (w-b) \% \text{array_size}$
- Tabular Hash



Original & HPW IBLT

● Knuth Hash ● Tabular Hash ● Linear Hash



500 trials per dot - k=3 hash functions

Stacked IBLT Parameters

- C : relates to column size in Basic IBLTs
- C_0 and δ : controls # of single row Basic IBLTs & column count of multirow Basic IBLTs
 - C_0 is hardcoded
 - δ is set by user
- $\tau = C * \log_2(1/\delta)$

```
Init( $\mathbf{h}$ )
_____
for  $0 \leq i < \lg(n) - \lg(\tau)$ 
     $T_i := \text{BasicInit}(1, \lceil Cn2^{-i} \rceil, \mathbf{h}_i)$ 
for  $0 \leq i < \lg(\tau)$ 
     $i' := \lfloor \lg(n) - \lg(\tau) \rfloor + i$ 
     $T_{i'} := \text{BasicInit}(2^i, \lceil C\tau2^{-i} \rceil, \mathbf{h}_{i'})$ 
return  $(T_0, \dots, T_{\lceil \lg n \rceil - 1})$ 
```

```
double tau = C_0 * log2(1/error_rate);
int logtau = std::ceil(log2(tau));
int singlerow_count = sketchCount - logtau;
```

```
// Single Row Columns
for (int index = 0; index < singlerow_count; index++) {
    int columns = std::ceil(C * threshold * std::pow(2, -index));
    this->addNextSketch(1, columns);
}

// Multirow Columns
for (int index = 0; index < logtau; index++) {
    int columns = std::ceil(C * tau * std::pow(2, -index));
    int rows = std::pow(2, index);

    this->addNextSketch(rows, columns);
}
```

Stacked Results

- C is the most responsive
 - Any $C > 5/e$ is robust
- $\tau (C_0 * \log_2 (1/\delta))$ is less responsive and less correlated at times

Stacked Success: C vs C_0 ($\delta=0.1$)

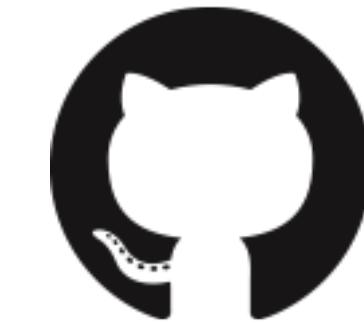
C	5	10	20
8/e		1	1
6/e		1	
5/e	1	1	
4.5/e		1	
4.25/e	0.99	0.98	
4.1/e	0.89	0.91	
4.05/e	0.85	0.948	
4/e		0.81	0.65

Summary

- Stacked IBLT is complex to use
 - Practical benefits over Original aren't substantial
- KnuthHash performs best
 - Original and HPW perform well with KnuthHash & Space Ratio ≥ 1.3
- Future areas to explore:
 - Store data types other than ints
 - Handle more anomalies
 - Performance / Memory Usage

Sources

- Fleischhacker, N., Larsen, K. G., Obremski, M., & Simkin, M. (2024). Invertible bloom lookup tables with less memory and randomness. 32nd Annual European Symposium on Algorithms (ESA 2024), 54:1-54:17. <https://doi.org/10.4230/LIPIcs.ESA.2024.54>
- Goodrich, M. T., & Mitzenmacher, M. (2011). Invertible bloom lookup tables. 2011 49th Annual Allerton Conference on Communication, Control, and Computing (Allerton), 792–799. <https://doi.org/10.1109/Allerton.2011.6120248>
- Houen, J. B. T., Pagh, R., & Walzer, S. (2023). Simple set sketching. Symposium on Simplicity in Algorithms. <https://doi.org/10.1137/1.9781611977585>



Ikellar/iblt